

Copyright by Pierian Data Inc.
For more information, visit us at www.pieriandata.com

Series

The first main data type we will learn about for pandas is the Series data type. Let's import Pandas and explore the Series object.

A Series is very similar to a NumPy array (in fact it is built on top of the NumPy array object). What differentiates the NumPy array from a Series, is that a Series can have axis labels, meaning it can be indexed by a label, instead of just a number location. It also doesn't need to hold numeric data, it can hold any arbitrary Python Object.

Let's explore this concept through some examples:

Imports

In [12]:

```
import numpy as np
import pandas as pd
```

Creating a Series from Python Objects

In [13]:

```
help(pd.Series)
```

Help on class Series in module pandas.core.series:

```
class Series(pandas.core.base.IndexOpsMixin, pandas.core.generic.NDFrame)
|   One-dimensional ndarray with axis labels (including time series).
|
|   Labels need not be unique but must be a hashable type. The object
|   supports both integer- and label-based indexing and provides a host of
|   methods for performing operations involving the index. Statistical
|   methods from ndarray have been overridden to automatically exclude
|   missing data (currently represented as NaN).
|
|   Operations between Series (+, -, /, *, **) align values based on their
|   associated index values-- they need not be the same length. The result
|   index will be the sorted union of the two indexes.
|
|   Parameters
|   -----
|   data : array-like, Iterable, dict, or scalar value
|       Contains data stored in Series.
|
|       .. versionchanged :: 0.23.0
|           If data is a dict, argument order is maintained for Python 3.6
|           and later.
|
|   index : array-like or Index (1d)
|       Values must be hashable and have the same length as `data`.
|       Non-unique index values are allowed. Will default to
|       RangeIndex (0, 1, 2, ..., n) if not provided. If both a dict and index
|       sequence are used, the index will override the keys found in the
```

```
dict.
dtype : str, numpy.dtype, or ExtensionDtype, optional
dtype for the output Series. If not specified, this will be
inferred from `data`.
See the :ref:`user guide <basics.dtypes>` for more usages.
copy : bool, default False
Copy input data.
```

Method resolution order:

```
Series
pandas.core.base.IndexOpsMixin
pandas.core.generic.NDFrame
pandas.core.base.PandasObject
pandas.core.base.StringMixin
pandas.core.accessor.DirNamesMixin
pandas.core.base.SelectionMixin
builtins.object
```

Methods defined here:

```
__add__(left, right)
```

```
__and__(self, other)
```

```
__array__(self, dtype=None)
Return the values as a NumPy array.
```

Users should not call this directly. Rather, it is invoked by
:func:`numpy.array` and :func:`numpy.asarray`.

Parameters

dtype : str or numpy.dtype, optional
The dtype to use for the resulting NumPy array. By default,
the dtype is inferred from the data.

Returns

numpy.ndarray
The values in the series converted to a :class:`numpy.ndarray`
with the specified `dtype`.

See Also

pandas.array : Create a new array from data.
Series.array : Zero-copy view to the array backing the Series.
Series.to_numpy : Series method for similar behavior.

Examples

```
>>> ser = pd.Series([1, 2, 3])
>>> np.asarray(ser)
array([1, 2, 3])
```

For timezone-aware data, the timezones may be retained with
``dtype='object'``

```
>>> tzser = pd.Series(pd.date_range('2000', periods=2, tz="CET"))
>>> np.asarray(tzser, dtype="object")
array([Timestamp('2000-01-01 00:00:00+0100', tz='CET', freq='D'),
      Timestamp('2000-01-02 00:00:00+0100', tz='CET', freq='D')],
      dtype=object)
```

Or the values may be localized to UTC and the tzinfo discarded with
``dtype='datetime64[ns]'``

```
>>> np.asarray(tzser, dtype="datetime64[ns]") # doctest: +ELLIPSIS
array(['1999-12-31T23:00:00.000000000', ...],
      dtype='datetime64[ns]')
```

```
__arrayprepare__(self, result, context=None)
Gets called prior to a ufunc.
```

```

|   __array_wrap__(self, result, context=None)
|       Gets called after a ufunc.
|
|   __div__ = __truediv__(left, right)
|
|   __divmod__(left, right)
|
|   __eq__(self, other, axis=None)
|
|   __float__(self)
|
|   __floordiv__(left, right)
|
|   __ge__(self, other, axis=None)
|
|   __getitem__(self, key)
|
|   __gt__(self, other, axis=None)
|
|   __iadd__(self, other)
|
|   __iand__(self, other)
|
|   __ifloordiv__(self, other)
|
|   __imod__(self, other)
|
|   __imul__(self, other)
|
|   __init__(self, data=None, index=None, dtype=None, name=None, copy=False, fastpath=False)
|       Initialize self. See help(type(self)) for accurate signature.
|
|   __int__(self)
|
|   __ior__(self, other)
|
|   __ipow__(self, other)
|
|   __isub__(self, other)
|
|   __itruediv__(self, other)
|
|   __ixor__(self, other)
|
|   __le__(self, other, axis=None)
|
|   __len__(self)
|       Return the length of the Series.
|
|   __long__ = __int__(self)
|
|   __lt__(self, other, axis=None)
|
|   __matmul__(self, other)
|       Matrix multiplication using binary `@` operator in Python>=3.5.
|
|   __mod__(left, right)
|
|   __mul__(left, right)
|
|   __ne__(self, other, axis=None)
|
|   __or__(self, other)
|
|   __pow__(left, right)
|
|   __radd__(left, right)
|
|   __rand__(self, other)

```

```

__rdiv__ = __rtruediv__(left, right)

__rdivmod__(left, right)

__rfloordiv__(left, right)

__rmatmul__(self, other)
    Matrix multiplication using binary `@` operator in Python>=3.5.

__rmod__(left, right)

__rmul__(left, right)

__ror__(self, other)

__rpow__(left, right)

__rsub__(left, right)

__rtruediv__(left, right)

__rxor__(self, other)

__setitem__(self, key, value)

__sub__(left, right)

__truediv__(left, right)

__unicode__(self)
    Return a string representation for a particular DataFrame.

    Invoked by unicode(df) in py2 only. Yields a Unicode String in both
    py2/py3.

__xor__(self, other)

add(self, other, level=None, fill_value=None, axis=0)
    Addition of series and other, element-wise (binary operator `add`).

    Equivalent to ``series + other``, but with support to substitute a fill_value for
    missing data in one of the inputs.

    Parameters
    -----
    other : Series or scalar value
    fill_value : None or float value, default None (NaN)
        Fill existing missing (NaN) values, and any new element needed for
        successful Series alignment, with this value before computation.
        If data in both corresponding Series locations is missing
        the result will be missing
    level : int or name
        Broadcast across a level, matching Index values on the
        passed MultiIndex level

    Returns
    -----
    result : Series

    See Also
    -----
    Series.radd

    Examples
    -----
    >>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
    >>> a
    a    1.0
    b    1.0
    c    1.0
    d    NaN
    dtype: float64

```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a    1.0
b    NaN
d    1.0
e    NaN
dtype: float64
>>> a.add(b, fill_value=0)
a    2.0
b    1.0
c    1.0
d    1.0
e    NaN
dtype: float64
```

```
agg = aggregate(self, func, axis=0, *args, **kwargs)
```

```
aggregate(self, func, axis=0, *args, **kwargs)
```

Aggregate using one or more operations over the specified axis.

.. versionadded:: 0.20.0

Parameters

func : function, str, list or dict

Function to use for aggregating the data. If a function, must either work when passed a Series or when passed to Series.apply.

Accepted combinations are:

- function
- string function name
- list of functions and/or function names, e.g. ``[np.sum, 'mean']``
- dict of axis labels -> functions, function names or list of such.

axis : {0 or 'index'}

Parameter needed for compatibility with DataFrame.

*args

Positional arguments to pass to `func`.

**kwargs

Keyword arguments to pass to `func`.

Returns

DataFrame, Series or scalar

if DataFrame.agg is called with a single function, returns a Series

if DataFrame.agg is called with several functions, returns a DataFrame

if Series.agg is called with single function, returns a scalar

if Series.agg is called with several functions, returns a Series

See Also

Series.apply : Invoke function on a Series.

Series.transform : Transform function producing a Series with like indexes.

Notes

`agg` is an alias for `aggregate`. Use the alias.

A passed user-defined-function will be passed a Series for evaluation.

Examples

```
>>> s = pd.Series([1, 2, 3, 4])
```

```
>>> s
```

```
0    1
```

```
1    2
```

```
2    3
```

```
3    4
```

```
dtype: int64
```

```
>>> s.agg('min')
1
```

```
>>> s.agg(['min', 'max'])
min    1
max    4
dtype: int64
```

```
align(self, other, join='outer', axis=None, level=None, copy=True, fill_value=None, method=None, limit=None, fill_axis=0, broadcast_axis=None)
```

Align two objects on their axes with the specified join method for each axis Index.

Parameters

other : DataFrame or Series

join : {'outer', 'inner', 'left', 'right'}, default 'outer'

axis : allowed axis of the other object, default None

Align on index (0), columns (1), or both (None)

level : int or level name, default None

Broadcast across a level, matching Index values on the passed MultiIndex level

copy : boolean, default True

Always returns new objects. If copy=False and no reindexing is required then original objects are returned.

fill_value : scalar, default np.NaN

Value to use for missing values. Defaults to NaN, but can be any "compatible" value

method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None

Method to use for filling holes in reindexed Series

pad / ffill: propagate last valid observation forward to next valid

backfill / bfill: use NEXT valid observation to fill gap

limit : int, default None

If method is specified, this is the maximum number of consecutive NaN values to forward/backward fill. In other words, if there is a gap with more than this number of consecutive NaNs, it will only be partially filled. If method is not specified, this is the maximum number of entries along the entire axis where NaNs will be filled. Must be greater than 0 if not None.

fill_axis : {0 or 'index'}, default 0

Filling axis, method and limit

broadcast_axis : {0 or 'index'}, default None

Broadcast values along this axis, if aligning two objects of different dimensions

Returns

(left, right) : (Series, type of other)

Aligned objects

```
all(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)
```

Return whether all elements are True, potentially over an axis.

Returns True unless there at least one element within a series or along a Dataframe axis that is False or equivalent (e.g. zero or empty).

Parameters

axis : {0 or 'index', 1 or 'columns', None}, default 0

Indicate which axis or axes should be reduced.

* 0 / 'index' : reduce the index, return a Series whose index is the original column labels.

* 1 / 'columns' : reduce the columns, return a Series whose index is the original index.

* None : reduce all axes, return a scalar.

bool_only : bool, default None

Include only boolean columns. If None, will attempt to use everything, then use only boolean data. Not implemented for Series.

`skipna : bool, default True`
 Exclude NA/null values. If the entire row/column is NA and `skipna` is True, then the result will be True, as for an empty row/column.
 If `skipna` is False, then NA are treated as True, because these are not equal to zero.
`level : int or level name, default None`
 If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.
`**kwargs : any, default None`
 Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

scalar or Series

If level is specified, then, Series is returned; otherwise, scalar is returned.

See Also

`Series.all` : Return True if all elements are True.

`DataFrame.any` : Return True if one (or more) elements are True.

Examples

****Series****

```

>>> pd.Series([True, True]).all()
True
>>> pd.Series([True, False]).all()
False
>>> pd.Series([]).all()
True
>>> pd.Series([np.nan]).all()
True
>>> pd.Series([np.nan]).all(skipna=False)
True

```

****DataFrames****

Create a dataframe from a dictionary.

```

>>> df = pd.DataFrame({'col1': [True, True], 'col2': [True, False]})
>>> df
   col1  col2
0  True  True
1  True False

```

Default behaviour checks if column-wise values all return True.

```

>>> df.all()
col1    True
col2   False
dtype: bool

```

Specify ``axis='columns'`` to check if row-wise values all return True.

```

>>> df.all(axis='columns')
0    True
1   False
dtype: bool

```

Or ``axis=None`` for whether every value is True.

```

>>> df.all(axis=None)
False

```

`any(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)`
 Return whether any element is True, potentially over an axis.

Returns False unless there at least one element within a series or along a Dataframe axis that is True or equivalent (e.g. non-zero or

```
non-empty).
```

Parameters

axis : {0 or 'index', 1 or 'columns', None}, default 0

Indicate which axis or axes should be reduced.

- * 0 / 'index' : reduce the index, return a Series whose index is the original column labels.

- * 1 / 'columns' : reduce the columns, return a Series whose index is the original index.

- * None : reduce all axes, return a scalar.

bool_only : bool, default None

Include only boolean columns. If None, will attempt to use everything, then use only boolean data. Not implemented for Series.

skipna : bool, default True

Exclude NA/null values. If the entire row/column is NA and skipna is True, then the result will be False, as for an empty row/column.

If skipna is False, then NA are treated as True, because these are not equal to zero.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

****kwargs** : any, default None

Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

scalar or Series

If level is specified, then, Series is returned; otherwise, scalar is returned.

See Also

numpy.any : Numpy version of this method.

Series.any : Return whether any element is True.

Series.all : Return whether all elements are True.

DataFrame.any : Return whether any element is True over requested axis.

DataFrame.all : Return whether all elements are True over requested axis.

Examples

****Series****

For Series input, the output is a scalar indicating whether any element is True.

```
>>> pd.Series([False, False]).any()
```

```
False
```

```
>>> pd.Series([True, False]).any()
```

```
True
```

```
>>> pd.Series([]).any()
```

```
False
```

```
>>> pd.Series([np.nan]).any()
```

```
False
```

```
>>> pd.Series([np.nan]).any(skipna=False)
```

```
True
```

****DataFrame****

Whether each column contains at least one True element (the default).

```
>>> df = pd.DataFrame({"A": [1, 2], "B": [0, 2], "C": [0, 0]})
```

```
>>> df
```

```
   A  B  C
```

```
0  1  0  0
```

```
1  2  2  0
```

```
>>> df.any()
```

```
A      True
```



```
B    True
C    False
dtype: bool
```

Aggregating over the columns.

```
>>> df = pd.DataFrame({"A": [True, False], "B": [1, 2]})
>>> df
```

```
   A  B
0  True  1
1 False  2
```

```
>>> df.any(axis='columns')
0    True
1    True
dtype: bool
```

```
>>> df = pd.DataFrame({"A": [True, False], "B": [1, 0]})
>>> df
```

```
   A  B
0  True  1
1 False  0
```

```
>>> df.any(axis='columns')
0    True
1   False
dtype: bool
```

Aggregating over the entire DataFrame with ``axis=None``.

```
>>> df.any(axis=None)
True
```

``any`` for an empty DataFrame is an empty Series.

```
>>> pd.DataFrame([]).any()
Series([], dtype: bool)
```

```
append(self, to_append, ignore_index=False, verify_integrity=False)
Concatenate two or more Series.
```

Parameters

to_append : Series or list/tuple of Series

ignore_index : boolean, default False

If True, do not use the index labels.

.. versionadded:: 0.19.0

verify_integrity : boolean, default False

If True, raise Exception on creating index with duplicates

Returns

appended : Series

See Also

concat : General function to concatenate DataFrame, Series
or Panel objects.

Notes

Iteratively appending to a Series can be more computationally intensive than a single concatenate. A better solution is to append values to a list and then concatenate the list with the original Series all at once.

Examples

```
>>> s1 = pd.Series([1, 2, 3])
```

```
>>> s2 = pd.Series([4, 5, 6])
```

```
>>> s3 = pd.Series([4, 5, 6], index=[3,4,5])
>>> s1.append(s2)
0    1
1    2
2    3
0    4
1    5
2    6
dtype: int64
```

```
>>> s1.append(s3)
0    1
1    2
2    3
3    4
4    5
5    6
dtype: int64
```

With `ignore_index` set to True:

```
>>> s1.append(s2, ignore_index=True)
0    1
1    2
2    3
3    4
4    5
5    6
dtype: int64
```

With `verify_integrity` set to True:

```
>>> s1.append(s2, verify_integrity=True)
Traceback (most recent call last):
...
ValueError: Indexes have overlapping values: [0, 1, 2]
```

```
apply(self, func, convert_dtype=True, args=(), **kwds)
Invoke function on values of Series.
```

Can be ufunc (a NumPy function that applies to the entire Series) or a Python function that only works on single values.

Parameters

```
func : function
    Python function or NumPy ufunc to apply.
convert_dtype : bool, default True
    Try to find better dtype for elementwise function results. If
    False, leave as dtype=object.
args : tuple
    Positional arguments passed to func after the series value.
**kwds
    Additional keyword arguments passed to func.
```

Returns

```
Series or DataFrame
    If func returns a Series object the result will be a DataFrame.
```

See Also

```
Series.map: For element-wise operations.
Series.agg: Only perform aggregating type operations.
Series.transform: Only perform transforming type operations.
```

Examples

Create a series with typical summer temperatures for each city.

```
>>> s = pd.Series([20, 21, 12],
...                 index=['London', 'New York', 'Helsinki'])
```

```
>>> s
London      20
New York    21
Helsinki    12
dtype: int64
```

Square the values by defining a function and passing it as an argument to ``apply``.

```
>>> def square(x):
...     return x ** 2
>>> s.apply(square)
London      400
New York    441
Helsinki    144
dtype: int64
```

Square the values by passing an anonymous function as an argument to ``apply``.

```
>>> s.apply(lambda x: x ** 2)
London      400
New York    441
Helsinki    144
dtype: int64
```

Define a custom function that needs additional positional arguments and pass these additional arguments using the ``args`` keyword.

```
>>> def subtract_custom_value(x, custom_value):
...     return x - custom_value

>>> s.apply(subtract_custom_value, args=(5,))
London      15
New York    16
Helsinki     7
dtype: int64
```

Define a custom function that takes keyword arguments and pass these arguments to ``apply``.

```
>>> def add_custom_values(x, **kwargs):
...     for month in kwargs:
...         x += kwargs[month]
...     return x

>>> s.apply(add_custom_values, june=30, july=20, august=25)
London      95
New York    96
Helsinki    87
dtype: int64
```

Use a function from the Numpy library.

```
>>> s.apply(np.log)
London      2.995732
New York    3.044522
Helsinki    2.484907
dtype: float64
```

```
argmax = idxmax(self, axis=0, skipna=True, *args, **kwargs)
Return the row label of the maximum value.
```

```
.. deprecated:: 0.21.0
```

The current behaviour of 'Series.argmax' is deprecated, use 'idxmax' instead.

The behavior of 'argmax' will be corrected to return the positional maximum in the future. For now, use 'series.values.argmax' or 'np.argmax(np.array(values))' to get the position of the maximum row.

If multiple values equal the maximum, the first row label with that value is returned.

Parameters

`skipna` : boolean, default True
Exclude NA/null values. If the entire Series is NA, the result will be NA.
`axis` : int, default 0
For compatibility with `DataFrame.idxmax`. Redundant for application on Series.
`*args, **kwargs`
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

`idxmax` : Index of maximum of values.

Raises

`ValueError`
If the Series is empty.

See Also

`numpy.argmax` : Return indices of the maximum values along the given axis.
`DataFrame.idxmax` : Return index of first occurrence of maximum over requested axis.
`Series.idxmin` : Return index *label* of the first occurrence of minimum of values.

Notes

This method is the Series version of ``ndarray.argmax``. This method returns the label of the maximum, while ``ndarray.argmax`` returns the position. To get the position, use ``series.values.argmax()``.

Examples

```
>>> s = pd.Series(data=[1, None, 4, 3, 4],
...                 index=['A', 'B', 'C', 'D', 'E'])
>>> s
A    1.0
B    NaN
C    4.0
D    3.0
E    4.0
dtype: float64
```

```
>>> s.idxmax()
'C'
```

If ``skipna`` is False and there is an NA value in the data, the function returns ``nan``.

```
>>> s.idxmax(skipna=False)
nan
```

`argmin = idxmin(self, axis=0, skipna=True, *args, **kwargs)`
Return the row label of the minimum value.

.. deprecated:: 0.21.0

The current behaviour of `'Series.argmin'` is deprecated, use `'idxmin'` instead.

The behavior of `'argmin'` will be corrected to return the positional minimum in the future. For now, use `'series.values.argmin'` or `'np.argmax(np.array(values))'` to get the position of the minimum row.

If multiple values equal the minimum, the first row label with that value is returned.

Parameters

`skipna` : boolean, default True
Exclude NA/null values. If the entire Series is NA, the result will be NA.
`axis` : int, default 0
For compatibility with `DataFrame.idxmin`. Redundant for application on Series.
`*args, **kwargs`
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

`idxmin` : Index of minimum of values.

Raises

`ValueError`
If the Series is empty.

See Also

`numpy.argmin` : Return indices of the minimum values along the given axis.
`DataFrame.idxmin` : Return index of first occurrence of minimum over requested axis.
`Series.idxmax` : Return index *label* of the first occurrence of maximum of values.

Notes

This method is the Series version of ```ndarray.argmin```. This method returns the label of the minimum, while ```ndarray.argmin``` returns the position. To get the position, use ```series.values.argmin()```.

Examples

```
>>> s = pd.Series(data=[1, None, 4, 1],
...                 index=['A', 'B', 'C', 'D'])
>>> s
A    1.0
B    NaN
C    4.0
D    1.0
dtype: float64
```

```
>>> s.idxmin()
'A'
```

If `skipna` is False and there is an NA value in the data, the function returns ```nan```.

```
>>> s.idxmin(skipna=False)
nan
```

`argsort(self, axis=0, kind='quicksort', order=None)`
Overrides `ndarray.argsort`. Arg sorts the value, omitting NA/null values, and places the result in the same locations as the non-NA values.

Parameters

`axis` : int
Has no effect but is accepted for compatibility with numpy.
`kind` : {'mergesort', 'quicksort', 'heapsort'}, default 'quicksort'
Choice of sorting algorithm. See `np.sort` for more information. 'mergesort' is the only stable algorithm
`order` : None

Has no effect but is accepted for compatibility with numpy.

Returns

argsorted : Series, with -1 indicated where nan values are present

See Also

numpy.ndarray.argsort

autocorr(self, lag=1)

Compute the lag-N autocorrelation.

This method computes the Pearson correlation between the Series and its shifted self.

Parameters

lag : int, default 1

Number of lags to apply before performing autocorrelation.

Returns

float

The Pearson correlation between self and self.shift(lag).

See Also

Series.corr : Compute the correlation between two Series.

Series.shift : Shift index by desired number of periods.

DataFrame.corr : Compute pairwise correlation of columns.

DataFrame.corrwith : Compute pairwise correlation between rows or columns of two DataFrame objects.

Notes

If the Pearson correlation is not well defined return 'NaN'.

Examples

```
>>> s = pd.Series([0.25, 0.5, 0.2, -0.05])
```

```
>>> s.autocorr() # doctest: +ELLIPSIS
```

```
0.10355...
```

```
>>> s.autocorr(lag=2) # doctest: +ELLIPSIS
```

```
-0.99999...
```

If the Pearson correlation is not well defined, then 'NaN' is returned.

```
>>> s = pd.Series([1, 0, 0, 0])
```

```
>>> s.autocorr()
```

```
nan
```

between(self, left, right, inclusive=True)

Return boolean Series equivalent to left <= series <= right.

This function returns a boolean vector containing `True` wherever the corresponding Series element is between the boundary values `left` and `right`. NA values are treated as `False`.

Parameters

left : scalar

Left boundary.

right : scalar

Right boundary.

inclusive : bool, default True

Include boundaries.

Returns

Series

Each element will be a boolean.

See Also

Series.gt : Greater than of series and other.

Series.lt : Less than of series and other.

Notes

This function is equivalent to `((left <= ser) & (ser <= right))`

Examples

```
>>> s = pd.Series([2, 0, 4, 8, np.nan])
```

Boundary values are included by default:

```
>>> s.between(1, 4)
```

```
0      True
1     False
2      True
3     False
4     False
dtype: bool
```

With `'inclusive'` set to `'False'` boundary values are excluded:

```
>>> s.between(1, 4, inclusive=False)
```

```
0      True
1     False
2     False
3     False
4     False
dtype: bool
```

`'left'` and `'right'` can be any scalar value:

```
>>> s = pd.Series(['Alice', 'Bob', 'Carol', 'Eve'])
```

```
>>> s.between('Anna', 'Daniel')
```

```
0     False
1      True
2      True
3     False
dtype: bool
```

`combine(self, other, func, fill_value=None)`

Combine the Series with a Series or scalar according to `'func'`.

Combine the Series and `'other'` using `'func'` to perform elementwise selection for combined Series.

`'fill_value'` is assumed when value is missing at some index from one of the two objects being combined.

Parameters

`other` : Series or scalar

The value(s) to be combined with the `'Series'`.

`func` : function

Function that takes two scalars as inputs and returns an element.

`fill_value` : scalar, optional

The value to assume when an index is missing from one Series or the other. The default specifies to use the appropriate NaN value for the underlying dtype of the Series.

Returns

Series

The result of combining the Series with the other object.

See Also

`Series.combine_first` : Combine Series values, choosing the calling Series' values first.

Examples

Consider 2 Datasets ``s1`` and ``s2`` containing highest clocked speeds of different birds.

```
>>> s1 = pd.Series({'falcon': 330.0, 'eagle': 160.0})
>>> s1
falcon    330.0
eagle     160.0
dtype: float64
>>> s2 = pd.Series({'falcon': 345.0, 'eagle': 200.0, 'duck': 30.0})
>>> s2
falcon    345.0
eagle     200.0
duck       30.0
dtype: float64
```

Now, to combine the two datasets and view the highest speeds of the birds across the two datasets

```
>>> s1.combine(s2, max)
duck      NaN
eagle     200.0
falcon    345.0
dtype: float64
```

In the previous example, the resulting value for duck is missing, because the maximum of a NaN and a float is a NaN. So, in the example, we set ``fill_value=0``, so the maximum value returned will be the value from some dataset.

```
>>> s1.combine(s2, max, fill_value=0)
duck       30.0
eagle     200.0
falcon    345.0
dtype: float64
```

`combine_first(self, other)`

Combine Series values, choosing the calling Series's values first.

Parameters

`other` : Series

The value(s) to be combined with the `Series`.

Returns

Series

The result of combining the Series with the other object.

See Also

`Series.combine` : Perform elementwise operation on two Series using a given function.

Notes

Result index will be the union of the two indexes.

Examples

```
>>> s1 = pd.Series([1, np.nan])
>>> s2 = pd.Series([3, 4])
>>> s1.combine_first(s2)
0    1.0
1    4.0
dtype: float64
```

`compound(self, axis=None, skipna=None, level=None)`

Return the compound percentage of the values for the requested axis.

Parameters

axis : {index (0)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

**kwargs

Additional keyword arguments to be passed to the function.

Returns

compounded : scalar or Series (if level specified)

compress(self, condition, *args, **kwargs)

Return selected slices of an array along given axis as a Series.

.. deprecated:: 0.24.0

See Also

numpy.ndarray.compress

corr(self, other, method='pearson', min_periods=None)

Compute correlation with `other` Series, excluding missing values.

Parameters

other : Series

method : {'pearson', 'kendall', 'spearman'} or callable

* pearson : standard correlation coefficient

* kendall : Kendall Tau correlation coefficient

* spearman : Spearman rank correlation

* callable: callable with input two 1d ndarray and returning a float

.. versionadded:: 0.24.0

min_periods : int, optional

Minimum number of observations needed to have a valid result

Returns

correlation : float

Examples

```
>>> histogram_intersection = lambda a, b: np.minimum(a, b  
... ).sum().round(decimals=1)
```

```
>>> s1 = pd.Series([.2, .0, .6, .2])
```

```
>>> s2 = pd.Series([.3, .6, .0, .1])
```

```
>>> s1.corr(s2, method=histogram_intersection)
```

```
0.3
```

count(self, level=None)

Return number of non-NA/null observations in the Series.

Parameters

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a smaller Series

Returns

nobs : int or Series (if level specified)

cov(self, other, min_periods=None)

Compute covariance with Series, excluding missing values.

Parameters

other : Series

min_periods : int, optional

Minimum number of observations needed to have a valid result

Returns

covariance : float

Normalized by N-1 (unbiased estimator).

cummax(self, axis=None, skipna=True, *args, **kwargs)

Return cumulative maximum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative maximum.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0

The index or the name of the axis. 0 is equivalent to None or 'index'.

skipna : boolean, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

*args, **kwargs :

Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

cummax : scalar or Series

See Also

core.window.Expanding.max : Similar functionality but ignores ``NaN`` values.

Series.max : Return the maximum over Series axis.

Series.cummax : Return cumulative maximum over Series axis.

Series.cummin : Return cumulative minimum over Series axis.

Series.cumsum : Return cumulative sum over Series axis.

Series.cumprod : Return cumulative product over Series axis.

Examples

****Series****

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
```

```
>>> s
```

```
0    2.0
```

```
1    NaN
```

```
2    5.0
```

```
3   -1.0
```

```
4    0.0
```

```
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cummax()
```

```
0    2.0
```

```
1    NaN
```

```
2    5.0
```

```
3    5.0
```

```
4    5.0
```

```
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cummax(skipna=False)
```

```
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                    [3.0, np.nan],
...                    [1.0, 0.0]],
...                    columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the maximum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cummax()
   A    B
0  2.0  1.0
1  3.0  NaN
2  3.0  1.0
```

To iterate over columns and find the maximum in each row, use ``axis=1``

```
>>> df.cummax(axis=1)
   A    B
0  2.0  2.0
1  3.0  NaN
2  1.0  1.0
```

`cummin(self, axis=None, skipna=True, *args, **kwargs)`

Return cumulative minimum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative minimum.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0

The index or the name of the axis. 0 is equivalent to None or 'index'.

`skipna` : boolean, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

`*args, **kwargs` :

Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

`cummin` : scalar or Series

See Also

`core.window.Expanding.min` : Similar functionality but ignores ``NaN`` values.

`Series.min` : Return the minimum over Series axis.

`Series.cummax` : Return cumulative maximum over Series axis.

`Series.cummin` : Return cumulative minimum over Series axis.

`Series.cumsum` : Return cumulative sum over Series axis.

`Series.cumprod` : Return cumulative product over Series axis.

Examples

****Series****

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0    2.0
1    NaN
2    5.0
3   -1.0
4    0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cummin()
0    2.0
1    NaN
2    2.0
3   -1.0
4   -1.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cummin(skipna=False)
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                    [3.0, np.nan],
...                    [1.0, 0.0]],
...                    columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the minimum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cummin()
   A    B
0  2.0  1.0
1  2.0  NaN
2  1.0  0.0
```

To iterate over columns and find the minimum in each row, use ``axis=1``

```
>>> df.cummin(axis=1)
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

`cumprod(self, axis=None, skipna=True, *args, **kwargs)`

Return cumulative product over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative product.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0

The index or the name of the axis. 0 is equivalent to None or 'index'.

`skipna` : boolean, default True

Exclude NA/null values. If an entire row/column is NA, the result

```
will be NA.
*args, **kwargs :
    Additional keywords have no effect but might be accepted for
    compatibility with NumPy.
```

Returns

```
-----
cumprod : scalar or Series
```

See Also

```
-----
core.window.Expanding.prod : Similar functionality
    but ignores ``NaN`` values.
Series.prod : Return the product over
    Series axis.
Series.cummax : Return cumulative maximum over Series axis.
Series.cummin : Return cumulative minimum over Series axis.
Series.cumsum : Return cumulative sum over Series axis.
Series.cumprod : Return cumulative product over Series axis.
```

Examples

```
-----
**Series**

>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0      2.0
1      NaN
2      5.0
3     -1.0
4      0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cumprod()
0      2.0
1      NaN
2     10.0
3    -10.0
4     -0.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cumprod(skipna=False)
0      2.0
1      NaN
2      NaN
3      NaN
4      NaN
dtype: float64
```

DataFrame

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                    [3.0, np.nan],
...                    [1.0, 0.0]],
...                    columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the product in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cumprod()
   A    B
0  2.0  1.0
1  6.0  NaN
```

```
2  6.0  0.0
```

To iterate over columns and find the product in each row,
use ``axis=1``

```
>>> df.cumprod(axis=1)
      A      B
0  2.0  2.0
1  3.0  NaN
2  1.0  0.0
```

`cumsum(self, axis=None, skipna=True, *args, **kwargs)`
Return cumulative sum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative sum.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0
The index or the name of the axis. 0 is equivalent to None or 'index'.
`skipna` : boolean, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.
`*args, **kwargs` :
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

`cumsum` : scalar or Series

See Also

`core.window.Expanding.sum` : Similar functionality but ignores ``NaN`` values.
`Series.sum` : Return the sum over Series axis.
`Series.cummax` : Return cumulative maximum over Series axis.
`Series.cummin` : Return cumulative minimum over Series axis.
`Series.cumsum` : Return cumulative sum over Series axis.
`Series.cumprod` : Return cumulative product over Series axis.

Examples

Series

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0    2.0
1    NaN
2    5.0
3   -1.0
4    0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cumsum()
0    2.0
1    NaN
2    7.0
3    6.0
4    6.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cumsum(skipna=False)
0    2.0
1    NaN
2    NaN
```

```
3     NaN
4     NaN
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                     [3.0, np.nan],
...                     [1.0, 0.0]],
...                     columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the sum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cumsum()
   A    B
0  2.0  1.0
1  5.0  NaN
2  6.0  1.0
```

To iterate over columns and find the sum in each row, use ``axis=1``

```
>>> df.cumsum(axis=1)
   A    B
0  2.0  3.0
1  3.0  NaN
2  1.0  1.0
```

`diff(self, periods=1)`
First discrete difference of element.

Calculates the difference of a Series element compared with another element in the Series (default is element in previous row).

Parameters

periods : int, default 1
Periods to shift for calculating difference, accepts negative values.

Returns

diffed : Series

See Also

Series.pct_change: Percent change over given number of periods.
Series.shift: Shift index by desired number of periods with an optional time freq.
DataFrame.diff: First discrete difference of object.

Examples

Difference with previous row

```
>>> s = pd.Series([1, 1, 2, 3, 5, 8])
>>> s.diff()
0     NaN
1     0.0
2     1.0
3     1.0
4     2.0
5     3.0
dtype: float64
```

Difference with 3rd previous row

```

|     >>> s.diff(periods=3)
|     0    NaN
|     1    NaN
|     2    NaN
|     3    2.0
|     4    4.0
|     5    6.0
|     dtype: float64
|
|     Difference with following row
|
|     >>> s.diff(periods=-1)
|     0    0.0
|     1   -1.0
|     2   -1.0
|     3   -2.0
|     4   -3.0
|     5    NaN
|     dtype: float64
|
|     div = truediv(self, other, level=None, fill_value=None, axis=0)
|
|     divide = truediv(self, other, level=None, fill_value=None, axis=0)
|
|     divmod(self, other, level=None, fill_value=None, axis=0)
|     Integer division and modulo of series and other, element-wise (binary operator `divmod`).
|
|     Equivalent to ``series divmod other``, but with support to substitute a fill_value for
|     missing data in one of the inputs.
|
|     Parameters
|     -----
|     other : Series or scalar value
|     fill_value : None or float value, default None (NaN)
|         Fill existing missing (NaN) values, and any new element needed for
|         successful Series alignment, with this value before computation.
|         If data in both corresponding Series locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : Series
|
|     See Also
|     -----
|     Series.rdivmod
|
|     Examples
|     -----
|     >>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
|     >>> a
|     a    1.0
|     b    1.0
|     c    1.0
|     d    NaN
|     dtype: float64
|     >>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
|     >>> b
|     a    1.0
|     b    NaN
|     d    1.0
|     e    NaN
|     dtype: float64
|     >>> a.add(b, fill_value=0)
|     a    2.0
|     b    1.0

```



```
c    1.0
d    1.0
e    NaN
dtype: float64
```

`dot(self, other)`

Compute the dot product between the Series and the columns of other.

This method computes the dot product between the Series and another one, or the Series and each columns of a DataFrame, or the Series and each columns of an array.

It can also be called using ``self @ other`` in Python `>= 3.5`.

Parameters

`other` : Series, DataFrame or array-like

The other object to compute the dot product with its columns.

Returns

scalar, Series or `numpy.ndarray`

Return the dot product of the Series and other if other is a Series, the Series of the dot product of Series and each rows of other if other is a DataFrame or a `numpy.ndarray` between the Series and each columns of the numpy array.

See Also

`DataFrame.dot`: Compute the matrix product with the DataFrame.

`Series.mul`: Multiplication of series and other, element-wise.

Notes

The Series and other has to share the same index if other is a Series or a DataFrame.

Examples

```
>>> s = pd.Series([0, 1, 2, 3])
```

```
>>> other = pd.Series([-1, 2, -3, 4])
```

```
>>> s.dot(other)
```

```
8
```

```
>>> s @ other
```

```
8
```

```
>>> df = pd.DataFrame([[0, 1], [-2, 3], [4, -5], [6, 7]])
```

```
>>> s.dot(df)
```

```
0    24
```

```
1    14
```

```
dtype: int64
```

```
>>> arr = np.array([[0, 1], [-2, 3], [4, -5], [6, 7]])
```

```
>>> s.dot(arr)
```

```
array([24, 14])
```

`drop(self, labels=None, axis=0, index=None, columns=None, level=None, inplace=False, errors='raise')`

Return Series with specified index labels removed.

Remove elements of a Series based on specifying the index labels.

When using a multi-index, labels on different levels can be removed by specifying the level.

Parameters

`labels` : single label or list-like

Index labels to drop.

`axis` : 0, default 0

Redundant for application on Series.

`index, columns` : None

Redundant for application on Series, but index can be used instead of labels.

```
.. versionadded:: 0.21.0
level : int or level name, optional
    For MultiIndex, level for which the labels will be removed.
inplace : bool, default False
    If True, do operation inplace and return None.
errors : {'ignore', 'raise'}, default 'raise'
    If 'ignore', suppress error and only existing labels are dropped.
```

Returns

```
-----
dropped : pandas.Series
```

Raises

```
-----
KeyError
    If none of the labels are found in the index.
```

See Also

```
-----
Series.reindex : Return only specified index labels of Series.
Series.dropna : Return series without null values.
Series.drop_duplicates : Return Series with duplicate values removed.
DataFrame.drop : Drop specified labels from rows or columns.
```

Examples

```
-----
>>> s = pd.Series(data=np.arange(3), index=['A','B','C'])
>>> s
A    0
B    1
C    2
dtype: int64
```

Drop labels B en C

```
>>> s.drop(labels=['B','C'])
A    0
dtype: int64
```

Drop 2nd level label in MultiIndex Series

```
>>> midx = pd.MultiIndex(levels=[['lama', 'cow', 'falcon'],
...                               ['speed', 'weight', 'length']],
...                       codes=[[0, 0, 0, 1, 1, 1, 2, 2, 2],
...                               [0, 1, 2, 0, 1, 2, 0, 1, 2]])
>>> s = pd.Series([45, 200, 1.2, 30, 250, 1.5, 320, 1, 0.3],
...               index=midx)
>>> s
lama      speed      45.0
          weight     200.0
          length      1.2
cow       speed      30.0
          weight     250.0
          length      1.5
falcon    speed     320.0
          weight      1.0
          length      0.3
dtype: float64

>>> s.drop(labels='weight', level=1)
lama      speed      45.0
          length      1.2
cow       speed      30.0
          length      1.5
falcon    speed     320.0
          length      0.3
dtype: float64
```

```
drop_duplicates(self, keep='first', inplace=False)
    Return Series with duplicate values removed.
```

Parameters

```

-----
keep : {'first', 'last', ``False``}, default 'first'
      - 'first' : Drop duplicates except for the first occurrence.
      - 'last' : Drop duplicates except for the last occurrence.
      - ``False`` : Drop all duplicates.
inplace : boolean, default ``False``
          If ``True``, performs operation inplace and returns None.

```

Returns

```

-----
deduplicated : Series

```

See Also

```

-----
Index.drop_duplicates : Equivalent method on Index.
DataFrame.drop_duplicates : Equivalent method on DataFrame.
Series.duplicated : Related method on Series, indicating duplicate
                    Series values.

```

Examples

```

-----
Generate an Series with duplicated entries.

```

```

>>> s = pd.Series(['lama', 'cow', 'lama', 'beetle', 'lama', 'hippo'],
...                 name='animal')
>>> s
0      lama
1       cow
2      lama
3    beetle
4      lama
5     hippo
Name: animal, dtype: object

```

With the 'keep' parameter, the selection behaviour of duplicated values can be changed. The value 'first' keeps the first occurrence for each set of duplicated entries. The default value of keep is 'first'.

```

>>> s.drop_duplicates()
0      lama
1       cow
3    beetle
5     hippo
Name: animal, dtype: object

```

The value 'last' for parameter 'keep' keeps the last occurrence for each set of duplicated entries.

```

>>> s.drop_duplicates(keep='last')
1       cow
3    beetle
4      lama
5     hippo
Name: animal, dtype: object

```

The value ``False`` for parameter 'keep' discards all sets of duplicated entries. Setting the value of 'inplace' to ``True`` performs the operation inplace and returns ``None``.

```

>>> s.drop_duplicates(keep=False, inplace=True)
>>> s
1       cow
3    beetle
5     hippo
Name: animal, dtype: object

```

```

dropna(self, axis=0, inplace=False, **kwargs)
Return a new Series with missing values removed.

```

See the :ref:`User Guide <missing_data>` for more on which values are considered missing, and how to work with missing data.

Parameters

axis : {0 or 'index'}, default 0

There is only one axis to drop values from.

inplace : bool, default False

If True, do operation inplace and return None.

**kwargs

Not in use.

Returns

Series

Series with NA entries dropped from it.

See Also

Series.isna: Indicate missing values.

Series.notna : Indicate existing (non-missing) values.

Series.fillna : Replace missing values.

DataFrame.dropna : Drop rows or columns which contain NA values.

Index.dropna : Drop missing indices.

Examples

```
>>> ser = pd.Series([1., 2., np.nan])
```

```
>>> ser
```

```
0    1.0
```

```
1    2.0
```

```
2    NaN
```

```
dtype: float64
```

Drop NA values from a Series.

```
>>> ser.dropna()
```

```
0    1.0
```

```
1    2.0
```

```
dtype: float64
```

Keep the Series with valid entries in the same variable.

```
>>> ser.dropna(inplace=True)
```

```
>>> ser
```

```
0    1.0
```

```
1    2.0
```

```
dtype: float64
```

Empty strings are not considered NA values. ``None`` is considered an NA value.

```
>>> ser = pd.Series([np.NaN, 2, pd.NaT, '', None, 'I stay'])
```

```
>>> ser
```

```
0    NaN
```

```
1      2
```

```
2    NaT
```

```
3
```

```
4    None
```

```
5    I stay
```

```
dtype: object
```

```
>>> ser.dropna()
```

```
1      2
```

```
3
```

```
5    I stay
```

```
dtype: object
```

```
duplicated(self, keep='first')
```

Indicate duplicate Series values.

Duplicated values are indicated as ``True`` values in the resulting Series. Either all duplicates, all except the first or all except the last occurrence of duplicates can be indicated.

Parameters

```
-----
keep : {'first', 'last', False}, default 'first'
      - 'first' : Mark duplicates as ``True`` except for the first
        occurrence.
      - 'last' : Mark duplicates as ``True`` except for the last
        occurrence.
      - ``False`` : Mark all duplicates as ``True``.
```

Returns

```
-----
pandas.core.series.Series
```

See Also

```
-----
Index.duplicated : Equivalent method on pandas.Index.
DataFrame.duplicated : Equivalent method on pandas.DataFrame.
Series.drop_duplicates : Remove duplicate values from Series.
```

Examples

```
-----
By default, for each set of duplicated values, the first occurrence is
set on False and all others on True:
```

```
>>> animals = pd.Series(['lama', 'cow', 'lama', 'beetle', 'lama'])
>>> animals.duplicated()
0    False
1    False
2     True
3    False
4     True
dtype: bool
```

which is equivalent to

```
>>> animals.duplicated(keep='first')
0    False
1    False
2     True
3    False
4     True
dtype: bool
```

By using 'last', the last occurrence of each set of duplicated values
is set on False and all others on True:

```
>>> animals.duplicated(keep='last')
0     True
1    False
2     True
3    False
4    False
dtype: bool
```

By setting keep on ``False``, all duplicates are True:

```
>>> animals.duplicated(keep=False)
0     True
1    False
2     True
3    False
4     True
dtype: bool
```

```
eq(self, other, level=None, fill_value=None, axis=0)
Equal to of series and other, element-wise (binary operator `eq`).
```

Equivalent to ``series == other``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

```
-----
```

other : Series or scalar value
 fill_value : None or float value, default None (NaN)
 Fill existing missing (NaN) values, and any new element needed for
 successful Series alignment, with this value before computation.
 If data in both corresponding Series locations is missing
 the result will be missing
 level : int or name
 Broadcast across a level, matching Index values on the
 passed MultiIndex level

Returns

result : Series

See Also

Series.None

Examples

```

>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a    1.0
b    1.0
c    1.0
d    NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a    1.0
b    NaN
d    1.0
e    NaN
dtype: float64
>>> a.add(b, fill_value=0)
a    2.0
b    1.0
c    1.0
d    1.0
e    NaN
dtype: float64

```

ewm(self, com=None, span=None, halflife=None, alpha=None, min_periods=0, adjust=True,
 ignore_na=False, axis=0)

Provides exponential weighted functions.

.. versionadded:: 0.18.0

Parameters

com : float, optional
 Specify decay in terms of center of mass,
 :math:\alpha = 1 / (1 + com), \text{ for } com \geq 0`
 span : float, optional
 Specify decay in terms of span,
 :math:\alpha = 2 / (span + 1), \text{ for } span \geq 1`
 halflife : float, optional
 Specify decay in terms of half-life,
 :math:\alpha = 1 - \exp(\log(0.5) / halflife), \text{ for } halflife > 0`
 alpha : float, optional
 Specify smoothing factor :math:\alpha` directly,
 :math:0 < \alpha \leq 1`

.. versionadded:: 0.18.0

min_periods : int, default 0

Minimum number of observations in window required to have a value
 (otherwise result is NA).

adjust : bool, default True

Divide by decaying adjustment factor in beginning periods to account
 for imbalance in relative weightings (viewing EWMA as a moving average)

ignore_na : bool, default False

Ignore missing values when calculating weights;
specify True to reproduce pre-0.15.0 behavior

Returns

a Window sub-classed for the particular operation

See Also

rolling : Provides rolling window calculations.

expanding : Provides expanding transformations.

Notes

Exactly one of center of mass, span, half-life, and alpha must be provided. Allowed values and relationship between the parameters are specified in the parameter descriptions above; see the link at the end of this section for a detailed explanation.

When adjust is True (default), weighted averages are calculated using weights $(1-\alpha)^{(n-1)}$, $(1-\alpha)^{(n-2)}$, ..., $1-\alpha$, 1.

When adjust is False, weighted averages are calculated recursively as:

```
weighted_average[0] = arg[0];
```

```
weighted_average[i] = (1-alpha)*weighted_average[i-1] + alpha*arg[i].
```

When ignore_na is False (default), weights are based on absolute positions. For example, the weights of x and y used in calculating the final weighted average of [x, None, y] are $(1-\alpha)^2$ and 1 (if adjust is True), and $(1-\alpha)^2$ and alpha (if adjust is False).

When ignore_na is True (reproducing pre-0.15.0 behavior), weights are based on relative positions. For example, the weights of x and y used in calculating the final weighted average of [x, None, y] are $1-\alpha$ and 1 (if adjust is True), and $1-\alpha$ and alpha (if adjust is False).

More details can be found at

[http://pandas.pydata.org/pandas-docs/stable/computation.html#exponentially-weight](http://pandas.pydata.org/pandas-docs/stable/computation.html#exponentially-weighted-windows)

ed-windows

Examples

```
>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
```

```
      B
0  0.0
1  1.0
2  2.0
3  NaN
4  4.0
```

```
>>> df.ewm(com=0.5).mean()
```

```
      B
0  0.000000
1  0.750000
2  1.615385
3  1.615385
4  3.670213
```

```
expanding(self, min_periods=1, center=False, axis=0)
```

Provides expanding transformations.

.. versionadded:: 0.18.0

Parameters

min_periods : int, default 1

Minimum number of observations in window required to have a value (otherwise result is NA).

center : bool, default False

Set the labels at the center of the window.

axis: int or str, default 0

Returns

a Window sub-classed for the particular operation

See Also

rolling : Provides rolling window calculations.

ewm : Provides exponential weighted functions.

Notes

By default, the result is set to the right edge of the window. This can be changed to the center of the window by setting ``center=True``.

Examples

```
>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
```

```
      B
0  0.0
1  1.0
2  2.0
3  NaN
4  4.0
```

```
>>> df.expanding(2).sum()
```

```
      B
0  NaN
1  1.0
2  3.0
3  3.0
4  7.0
```

fillna(self, value=None, method=None, axis=None, inplace=False, limit=None, downcast=None, **kwargs)

Fill NA/NaN values using the specified method.

Parameters

value : scalar, dict, Series, or DataFrame

Value to use to fill holes (e.g. 0), alternately a dict/Series/DataFrame of values specifying which value to use for each index (for a Series) or column (for a DataFrame). (values not in the dict/Series/DataFrame will not be filled). This value cannot be a list.

method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None

Method to use for filling holes in reindexed Series

pad / ffill: propagate last valid observation forward to next valid

backfill / bfill: use NEXT valid observation to fill gap

axis : {0 or 'index'}

inplace : boolean, default False

If True, fill in place. Note: this will modify any other views on this object, (e.g. a no-copy slice for a column in a DataFrame).

limit : int, default None

If method is specified, this is the maximum number of consecutive NaN values to forward/backward fill. In other words, if there is a gap with more than this number of consecutive NaNs, it will only be partially filled. If method is not specified, this is the maximum number of entries along the entire axis where NaNs will be filled. Must be greater than 0 if not None.

downcast : dict, default is None

a dict of item->dtype of what to downcast if possible, or the string 'infer' which will try to downcast to an appropriate equal type (e.g. float64 to int64 if possible)

Returns

filled : Series

See Also


```
-----
interpolate : Fill NaN values using interpolation.
reindex, asfreq
```

Examples

```
-----
>>> df = pd.DataFrame([[np.nan, 2, np.nan, 0],
...                     [3, 4, np.nan, 1],
...                     [np.nan, np.nan, np.nan, 5],
...                     [np.nan, 3, np.nan, 4]],
...                     columns=list('ABCD'))
```

```
>>> df
   A    B    C    D
0 NaN  2.0 NaN  0
1  3.0  4.0 NaN  1
2 NaN  NaN NaN  5
3 NaN  3.0 NaN  4
```

Replace all NaN elements with 0s.

```
>>> df.fillna(0)
   A    B    C    D
0  0.0  2.0  0.0  0
1  3.0  4.0  0.0  1
2  0.0  0.0  0.0  5
3  0.0  3.0  0.0  4
```

We can also propagate non-null values forward or backward.

```
>>> df.fillna(method='ffill')
   A    B    C    D
0 NaN  2.0 NaN  0
1  3.0  4.0 NaN  1
2  3.0  4.0 NaN  5
3  3.0  3.0 NaN  4
```

Replace all NaN elements in column 'A', 'B', 'C', and 'D', with 0, 1, 2, and 3 respectively.

```
>>> values = {'A': 0, 'B': 1, 'C': 2, 'D': 3}
>>> df.fillna(value=values)
   A    B    C    D
0  0.0  2.0  2.0  0
1  3.0  4.0  2.0  1
2  0.0  1.0  2.0  5
3  0.0  3.0  2.0  4
```

Only replace the first NaN element.

```
>>> df.fillna(value=values, limit=1)
   A    B    C    D
0  0.0  2.0  2.0  0
1  3.0  4.0 NaN  1
2  NaN  1.0 NaN  5
3  NaN  3.0 NaN  4
```

`floordiv(self, other, level=None, fill_value=None, axis=0)`

Integer division of series and other, element-wise (binary operator `'floordiv'`).

Equivalent to ```series // other```, but with support to substitute a `fill_value` for missing data in one of the inputs.

Parameters

`other` : Series or scalar value

`fill_value` : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

`level` : int or name

Broadcast across a level, matching Index values on the
passed MultiIndex level

Returns

result : Series

See Also

Series.rfloordiv

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a    1.0
```

```
b    1.0
```

```
c    1.0
```

```
d    NaN
```

```
dtype: float64
```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
```

```
>>> b
```

```
a    1.0
```

```
b    NaN
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```
>>> a.add(b, fill_value=0)
```

```
a    2.0
```

```
b    1.0
```

```
c    1.0
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```
ge(self, other, level=None, fill_value=None, axis=0)
```

Greater than or equal to of series and other, element-wise (binary operator `ge`)

Equivalent to ``series >= other``, but with support to substitute a fill_value for

missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for
successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing
the result will be missing

level : int or name

Broadcast across a level, matching Index values on the
passed MultiIndex level

Returns

result : Series

See Also

Series.None

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a    1.0
```

```
b    1.0
```

```
c    1.0
```

```
d    NaN
```

```
dtype: float64
```

```

>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a    1.0
b    NaN
d    1.0
e    NaN
dtype: float64
>>> a.add(b, fill_value=0)
a    2.0
b    1.0
c    1.0
d    1.0
e    NaN
dtype: float64

```

`get_value(self, label, takeable=False)`

Quickly retrieve single value at passed index label.

.. deprecated:: 0.21.0

Please use `.at[]` or `.iat[]` accessors.

Parameters

label : object

takeable : interpret the index as indexers, default False

Returns

value : scalar value

`get_values(self)`

Same as `values` (but handles sparseness conversions); is a view.

`gt(self, other, level=None, fill_value=None, axis=0)`

Greater than of series and other, element-wise (binary operator ``gt``).

Equivalent to ```series > other```, but with support to substitute a `fill_value` for missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

result : Series

See Also

`Series.None`

Examples

```

>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a    1.0
b    1.0
c    1.0
d    NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a    1.0
b    NaN

```

```

|     d    1.0
|     e    NaN
|     dtype: float64
|     >>> a.add(b, fill_value=0)
|     a    2.0
|     b    1.0
|     c    1.0
|     d    1.0
|     e    NaN
|     dtype: float64
|
|     hist = hist_series(self, by=None, ax=None, grid=True, xlabelsize=None, xrot=None, ylabelsize=None, yrot=None, figsize=None, bins=10, **kws)
|     Draw histogram of the input series using matplotlib.
|
|     Parameters
|     -----
|     by : object, optional
|         If passed, then used to form histograms for separate groups
|     ax : matplotlib axis object
|         If not passed, uses gca()
|     grid : boolean, default True
|         Whether to show axis grid lines
|     xlabelsize : int, default None
|         If specified changes the x-axis label size
|     xrot : float, default None
|         rotation of x axis labels
|     ylabelsize : int, default None
|         If specified changes the y-axis label size
|     yrot : float, default None
|         rotation of y axis labels
|     figsize : tuple, default None
|         figure size in inches by default
|     bins : integer or sequence, default 10
|         Number of histogram bins to be used. If an integer is given, bins + 1
|         bin edges are calculated and returned. If bins is a sequence, gives
|         bin edges, including left edge of first bin and right edge of last
|         bin. In this case, bins is returned unmodified.
|     bins : integer, default 10
|         Number of histogram bins to be used
|     `**kws` : keywords
|         To be passed to the actual plotting function
|
|     See Also
|     -----
|     matplotlib.axes.Axes.hist : Plot a histogram using matplotlib.
|
|     idxmax(self, axis=0, skipna=True, *args, **kwargs)
|     Return the row label of the maximum value.
|
|     If multiple values equal the maximum, the first row label with that
|     value is returned.
|
|     Parameters
|     -----
|     skipna : boolean, default True
|         Exclude NA/null values. If the entire Series is NA, the result
|         will be NA.
|     axis : int, default 0
|         For compatibility with DataFrame.idxmax. Redundant for application
|         on Series.
|     *args, **kwargs
|         Additional keywords have no effect but might be accepted
|         for compatibility with NumPy.
|
|     Returns
|     -----
|     idxmax : Index of maximum of values.
|
|     Raises
|     -----
|     ValueError

```

If the Series is empty.

See Also

`numpy.argmax` : Return indices of the maximum values
along the given axis.

`DataFrame.idxmax` : Return index of first occurrence of maximum
over requested axis.

`Series.idxmin` : Return index *label* of the first occurrence
of minimum of values.

Notes

This method is the Series version of ```ndarray.argmax```. This method
returns the label of the maximum, while ```ndarray.argmax``` returns
the position. To get the position, use ```series.values.argmax()```.

Examples

```
>>> s = pd.Series(data=[1, None, 4, 3, 4],  
...                index=['A', 'B', 'C', 'D', 'E'])
```

```
>>> s
```

```
A    1.0
```

```
B    NaN
```

```
C    4.0
```

```
D    3.0
```

```
E    4.0
```

```
dtype: float64
```

```
>>> s.idxmax()
```

```
'C'
```

If ``skipna`` is False and there is an NA value in the data,
the function returns ``nan``.

```
>>> s.idxmax(skipna=False)
```

```
nan
```

`idxmin(self, axis=0, skipna=True, *args, **kwargs)`

Return the row label of the minimum value.

If multiple values equal the minimum, the first row label with that
value is returned.

Parameters

`skipna` : boolean, default True

Exclude NA/null values. If the entire Series is NA, the result
will be NA.

`axis` : int, default 0

For compatibility with `DataFrame.idxmin`. Redundant for application
on Series.

`*args, **kwargs`

Additional keywords have no effect but might be accepted
for compatibility with NumPy.

Returns

`idxmin` : Index of minimum of values.

Raises

`ValueError`

If the Series is empty.

See Also

`numpy.argmin` : Return indices of the minimum values
along the given axis.

`DataFrame.idxmin` : Return index of first occurrence of minimum
over requested axis.

`Series.idxmax` : Return index *label* of the first occurrence

of maximum of values.

Notes

This method is the Series version of ``ndarray.argmin``. This method returns the label of the minimum, while ``ndarray.argmin`` returns the position. To get the position, use ``series.values.argmin()``.

Examples

```
>>> s = pd.Series(data=[1, None, 4, 1],
...                 index=['A', 'B', 'C', 'D'])
```

```
>>> s
```

```
A    1.0
```

```
B    NaN
```

```
C    4.0
```

```
D    1.0
```

```
dtype: float64
```

```
>>> s.idxmin()
```

```
'A'
```

If ``skipna`` is False and there is an NA value in the data, the function returns ``nan``.

```
>>> s.idxmin(skipna=False)
```

```
nan
```

```
isin(self, values)
```

Check whether ``values`` are contained in Series.

Return a boolean Series showing whether each element in the Series matches an element in the passed sequence of ``values`` exactly.

Parameters

values : set or list-like

The sequence of values to test. Passing in a single string will raise a ``TypeError``. Instead, turn a single string into a list of one element.

.. versionadded:: 0.18.1

Support for values as a set.

Returns

isin : Series (bool dtype)

Raises

TypeError

* If ``values`` is a string

See Also

DataFrame.isin : Equivalent method on DataFrame.

Examples

```
>>> s = pd.Series(['lama', 'cow', 'lama', 'beetle', 'lama',
...                'hippo'], name='animal')
```

```
>>> s.isin(['cow', 'lama'])
```

```
0     True
```

```
1     True
```

```
2     True
```

```
3    False
```

```
4     True
```

```
5    False
```

```
Name: animal, dtype: bool
```

Passing a single string as ``s.isin('lama')`` will raise an error. Use

a list of one element instead:

```
>>> s.isin(['lama'])
0      True
1     False
2      True
3     False
4      True
5     False
Name: animal, dtype: bool
```

isna(self)

Detect missing values.

Return a boolean same-sized object indicating if the values are NA. NA values, such as None or :attr:`numpy.NaN`, gets mapped to True values.

Everything else gets mapped to False values. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``).

Returns

Series

Mask of bool values for each element in Series that indicates whether an element is not an NA value.

See Also

Series.isnull : Alias of isna.

Series.notna : Boolean inverse of isna.

Series.dropna : Omit axes labels with missing values.

isna : Top-level isna.

Examples

Show which entries in a DataFrame are NA.

```
>>> df = pd.DataFrame({'age': [5, 6, np.NaN],
...                    'born': [pd.NaT, pd.Timestamp('1939-05-27'),
...                             pd.Timestamp('1940-04-25')],
...                    'name': ['Alfred', 'Batman', ''],
...                    'toy': [None, 'Batmobile', 'Joker']})
>>> df
   age      born    name      toy
0  5.0      NaT  Alfred     None
1  6.0 1939-05-27  Batman  Batmobile
2  NaN 1940-04-25      Joker
```

```
>>> df.isna()
   age      born    name      toy
0  False    True  False    True
1  False    False  False   False
2    True    False  False   False
```

Show which entries in a Series are NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0      5.0
1      6.0
2      NaN
dtype: float64
```

```
>>> ser.isna()
0      False
1      False
2        True
dtype: bool
```

isnull(self)

Detect missing values.

Return a boolean same-sized object indicating if the values are NA. NA values, such as None or :attr:`numpy.NaN`, gets mapped to True values. Everything else gets mapped to False values. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``).

Returns

Series

Mask of bool values for each element in Series that indicates whether an element is not an NA value.

See Also

Series.isnull : Alias of isna.

Series.notna : Boolean inverse of isna.

Series.dropna : Omit axes labels with missing values.

isna : Top-level isna.

Examples

Show which entries in a DataFrame are NA.

```
>>> df = pd.DataFrame({'age': [5, 6, np.NaN],
...                    'born': [pd.NaT, pd.Timestamp('1939-05-27'),
...                             pd.Timestamp('1940-04-25')],
...                    'name': ['Alfred', 'Batman', ''],
...                    'toy': [None, 'Batmobile', 'Joker']})
```

```
>>> df
   age      born   name      toy
0  5.0      NaT  Alfred     None
1  6.0 1939-05-27  Batman  Batmobile
2  NaN 1940-04-25      Joker
```

```
>>> df.isna()
   age  born  name  toy
0  False  True  False  True
1  False  False  False  False
2   True  False  False  False
```

Show which entries in a Series are NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
```

```
>>> ser
0    5.0
1    6.0
2    NaN
dtype: float64
```

```
>>> ser.isna()
0    False
1    False
2     True
dtype: bool
```

items = iteritems(self)

iteritems(self)

Lazily iterate over (index, value) tuples.

keys(self)

Alias for index.

kurt(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)

Return unbiased kurtosis over requested axis using Fisher's definition of kurtosis (kurtosis of normal == 0.0). Normalized by N-1.

Parameters

axis: {index (0)}

Axis for the function to be applied on.

skipna : bool, default True
Exclude NA/null values when computing the result.

level : int or level name, default None
If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric_only : bool, default None
Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

****kwargs**
Additional keyword arguments to be passed to the function.

Returns

kurt : scalar or Series (if level specified)

kurtosis = kurt(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs

)
le(self, other, level=None, fill_value=None, axis=0)

Less than or equal to of series and other, element-wise (binary operator `le`).

Equivalent to ``series <= other``, but with support to substitute a fill_value for

missing data in one of the inputs.

Parameters

other : Series or scalar value
fill_value : None or float value, default None (NaN)
Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.
If data in both corresponding Series locations is missing the result will be missing
level : int or name
Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

result : Series

See Also

Series.None

Examples

>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a 1.0
b 1.0
c 1.0
d NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a 1.0
b NaN
d 1.0
e NaN
dtype: float64
>>> a.add(b, fill_value=0)
a 2.0
b 1.0
c 1.0
d 1.0
e NaN
dtype: float64

lt(self, other, level=None, fill_value=None, axis=0)

Less than of series and other, element-wise (binary operator `lt`).

Equivalent to ``series < other``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

result : Series

See Also

Series.None

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a    1.0
```

```
b    1.0
```

```
c    1.0
```

```
d    NaN
```

```
dtype: float64
```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
```

```
>>> b
```

```
a    1.0
```

```
b    NaN
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```
>>> a.add(b, fill_value=0)
```

```
a    2.0
```

```
b    1.0
```

```
c    1.0
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```
mad(self, axis=None, skipna=None, level=None)
```

Return the mean absolute deviation of the values for the requested axis.

Parameters

axis : {index (0)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

**kwargs

Additional keyword arguments to be passed to the function.

Returns

mad : scalar or Series (if level specified)

```
map(self, arg, na_action=None)
```

Map values of Series according to input correspondence.

Used for substituting each value in a Series with another value, that may be derived from a function, a ``dict`` or a :class:`Series`.

Parameters

arg : function, dict, or Series

Mapping correspondence.

na_action : {None, 'ignore'}, default None

If 'ignore', propagate NaN values, without passing them to the mapping correspondence.

Returns

Series

Same index as caller.

See Also

Series.apply : For applying more complex functions on a Series.

DataFrame.apply : Apply a function row-/column-wise.

DataFrame.applymap : Apply a function elementwise on a whole DataFrame.

Notes

When ``arg`` is a dictionary, values in Series that are not in the dictionary (as keys) are converted to ``NaN``. However, if the dictionary is a ``dict`` subclass that defines ``__missing__`` (i.e. provides a method for default values), then this default is used rather than ``NaN``.

Examples

```
>>> s = pd.Series(['cat', 'dog', np.nan, 'rabbit'])
```

```
>>> s
```

```
0      cat
```

```
1      dog
```

```
2      NaN
```

```
3  rabbit
```

```
dtype: object
```

``map`` accepts a ``dict`` or a ``Series``. Values that are not found in the ``dict`` are converted to ``NaN``, unless the dict has a default value (e.g. ``defaultdict``):

```
>>> s.map({'cat': 'kitten', 'dog': 'puppy'})
```

```
0  kitten
```

```
1  puppy
```

```
2      NaN
```

```
3      NaN
```

```
dtype: object
```

It also accepts a function:

```
>>> s.map('I am a {}'.format)
```

```
0      I am a cat
```

```
1      I am a dog
```

```
2      I am a nan
```

```
3  I am a rabbit
```

```
dtype: object
```

To avoid applying the function to missing values (and keep them as ``NaN``) ``na_action='ignore'`` can be used:

```
>>> s.map('I am a {}'.format, na_action='ignore')
```

```
0      I am a cat
```

```
1      I am a dog
```

```
2              NaN
```

```
3  I am a rabbit
```

```
dtype: object
```

```
max(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
    Return the maximum of the values for the requested axis.
```

If you want the **index** of the maximum, use ``idxmax``. This is the equivalent of the ``numpy.ndarray`` method ``argmax``.

Parameters

axis : {index (0)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

**kwargs

Additional keyword arguments to be passed to the function.

Returns

max : scalar or Series (if level specified)

See Also

Series.sum : Return the sum.

Series.min : Return the minimum.

Series.max : Return the maximum.

Series.idxmin : Return the index of the minimum.

Series.idxmax : Return the index of the maximum.

DataFrame.min : Return the sum over the requested axis.

DataFrame.min : Return the minimum over the requested axis.

DataFrame.max : Return the maximum over the requested axis.

DataFrame.idxmin : Return the index of the minimum over the requested axis.

DataFrame.idxmax : Return the index of the maximum over the requested axis.

Examples

```
>>> idx = pd.MultiIndex.from_arrays([
...     ['warm', 'warm', 'cold', 'cold'],
...     ['dog', 'falcon', 'fish', 'spider']],
...     names=['blooded', 'animal'])
>>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
>>> s
```

```
blooded  animal
warm     dog      4
         falcon   2
cold     fish     0
         spider   8
Name: legs, dtype: int64
```

```
>>> s.max()
8
```

Max using level names, as well as indices.

```
>>> s.max(level='blooded')
blooded
warm     4
cold     8
Name: legs, dtype: int64
```

```
>>> s.max(level=0)
blooded
warm     4
cold     8
Name: legs, dtype: int64
```

```
mean(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
```

Return the mean of the values for the requested axis.

Parameters

axis : {index (0)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

**kwargs

Additional keyword arguments to be passed to the function.

Returns

mean : scalar or Series (if level specified)

median(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)

Return the median of the values for the requested axis.

Parameters

axis : {index (0)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

**kwargs

Additional keyword arguments to be passed to the function.

Returns

median : scalar or Series (if level specified)

memory_usage(self, index=True, deep=False)

Return the memory usage of the Series.

The memory usage can optionally include the contribution of the index and of elements of `object` dtype.

Parameters

index : bool, default True

Specifies whether to include the memory usage of the Series index.

deep : bool, default False

If True, introspect the data deeply by interrogating `object` dtypes for system-level memory consumption, and include it in the returned value.

Returns

int

Bytes of memory consumed.

See Also

numpy.ndarray.nbytes : Total bytes consumed by the elements of the array.

DataFrame.memory_usage : Bytes consumed by a DataFrame.

Examples

>>> s = pd.Series(range(3))

```
>>> s.memory_usage()
104
```

Not including the index gives the size of the rest of the data, which is necessarily smaller:

```
>>> s.memory_usage(index=False)
24
```

The memory footprint of `object` values is ignored by default:

```
>>> s = pd.Series(["a", "b"])
>>> s.values
array(['a', 'b'], dtype=object)
>>> s.memory_usage()
96
>>> s.memory_usage(deep=True)
212
```

```
min(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
Return the minimum of the values for the requested axis.
```

If you want the *index* of the minimum, use ``idxmin``. This is the equivalent of the ``numpy.ndarray`` method ``argmin``.

Parameters

axis : {index (0)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

****kwargs**

Additional keyword arguments to be passed to the function.

Returns

min : scalar or Series (if level specified)

See Also

Series.sum : Return the sum.

Series.min : Return the minimum.

Series.max : Return the maximum.

Series.idxmin : Return the index of the minimum.

Series.idxmax : Return the index of the maximum.

DataFrame.min : Return the sum over the requested axis.

DataFrame.min : Return the minimum over the requested axis.

DataFrame.max : Return the maximum over the requested axis.

DataFrame.idxmin : Return the index of the minimum over the requested axis.

DataFrame.idxmax : Return the index of the maximum over the requested axis.

Examples

```
>>> idx = pd.MultiIndex.from_arrays([
...     ['warm', 'warm', 'cold', 'cold'],
...     ['dog', 'falcon', 'fish', 'spider']],
...     names=['blooded', 'animal'])
>>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
>>> s
blooded animal
warm      dog      4
          falcon   2
cold     fish     0
          spider   8
Name: legs, dtype: int64
```

```
>>> s.min()
0
```

Min using level names, as well as indices.

```
>>> s.min(level='blooded')
blooded
warm      2
cold      0
Name: legs, dtype: int64
```

```
>>> s.min(level=0)
blooded
warm      2
cold      0
Name: legs, dtype: int64
```

```
mod(self, other, level=None, fill_value=None, axis=0)
Modulo of series and other, element-wise (binary operator `mod`).
```

Equivalent to ``series % other``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

other : Series or scalar value
fill_value : None or float value, default None (NaN)
Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.
If data in both corresponding Series locations is missing the result will be missing
level : int or name
Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

result : Series

See Also

Series.rmod

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a      1.0
b      1.0
c      1.0
d      NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a      1.0
b      NaN
d      1.0
e      NaN
dtype: float64
>>> a.add(b, fill_value=0)
a      2.0
b      1.0
c      1.0
d      1.0
e      NaN
dtype: float64
```

```
mode(self, dropna=True)
Return the mode(s) of the dataset.
```

Always returns Series even if only one value is returned.

Parameters

dropna : boolean, default True
Don't consider counts of NaN/NaT.

.. versionadded:: 0.24.0

Returns

modes : Series (sorted)

mul(self, other, level=None, fill_value=None, axis=0)

Multiplication of series and other, element-wise (binary operator `mul`).

Equivalent to ``series * other``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

result : Series

See Also

Series.rmul

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a    1.0
```

```
b    1.0
```

```
c    1.0
```

```
d    NaN
```

```
dtype: float64
```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
```

```
>>> b
```

```
a    1.0
```

```
b    NaN
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```
>>> a.add(b, fill_value=0)
```

```
a    2.0
```

```
b    1.0
```

```
c    1.0
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

multiply = mul(self, other, level=None, fill_value=None, axis=0)

ne(self, other, level=None, fill_value=None, axis=0)

Not equal to of series and other, element-wise (binary operator `ne`).

Equivalent to ``series != other``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

result : Series

See Also

Series.None

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a    1.0
```

```
b    1.0
```

```
c    1.0
```

```
d    NaN
```

```
dtype: float64
```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
```

```
>>> b
```

```
a    1.0
```

```
b    NaN
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```
>>> a.add(b, fill_value=0)
```

```
a    2.0
```

```
b    1.0
```

```
c    1.0
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

nlargest(self, n=5, keep='first')

Return the largest `n` elements.

Parameters

n : int, default 5

Return this many descending sorted values.

keep : {'first', 'last', 'all'}, default 'first'

When there are duplicate values that cannot all fit in a Series of `n` elements:

- ``first`` : take the first occurrences based on the index order
- ``last`` : take the last occurrences based on the index order
- ``all`` : keep all occurrences. This can result in a Series of size larger than `n`.

Returns

Series

The `n` largest values in the Series, sorted in decreasing order.

See Also

Series.nsmallest: Get the `n` smallest elements.

Series.sort_values: Sort Series by values.

Series.head: Return the first `n` rows.

Notes

Faster than ```.sort_values(ascending=False).head(n)``` for small ``n`` relative to the size of the ```Series``` object.

Examples

```
>>> countries_population = {"Italy": 59000000, "France": 65000000,
...                          "Malta": 434000, "Maldives": 434000,
...                          "Brunei": 434000, "Iceland": 337000,
...                          "Nauru": 11300, "Tuvalu": 11300,
...                          "Anguilla": 11300, "Monserat": 5200}
>>> s = pd.Series(countries_population)
>>> s
Italy      59000000
France     65000000
Malta       434000
Maldives    434000
Brunei      434000
Iceland     337000
Nauru       11300
Tuvalu      11300
Anguilla    11300
Monserat    5200
dtype: int64
```

The ``n`` largest elements where ```n=5``` by default.

```
>>> s.nlargest()
France     65000000
Italy      59000000
Malta       434000
Maldives    434000
Brunei      434000
dtype: int64
```

The ``n`` largest elements where ```n=3```. Default ``keep`` value is `'first'` so Malta will be kept.

```
>>> s.nlargest(3)
France     65000000
Italy      59000000
Malta       434000
dtype: int64
```

The ``n`` largest elements where ```n=3``` and keeping the last duplicates. Brunei will be kept since it is the last with value 434000 based on the index order.

```
>>> s.nlargest(3, keep='last')
France     65000000
Italy      59000000
Brunei      434000
dtype: int64
```

The ``n`` largest elements where ```n=3``` with all duplicates kept. Note that the returned Series has five elements due to the three duplicates.

```
>>> s.nlargest(3, keep='all')
France     65000000
Italy      59000000
Malta       434000
Maldives    434000
Brunei      434000
dtype: int64
```

`nonzero(self)`

Return the `*integer*` indices of the elements that are non-zero.

.. deprecated:: 0.24.0

Please use `.to_numpy().nonzero()` as a replacement.

This method is equivalent to calling ``numpy.nonzero`` on the series data. For compatibility with NumPy, the return value is

the same (a tuple with an array of indices for each dimension), but it will always be a one-item tuple because series only have one dimension.

See Also

`numpy.nonzero`

Examples

```
>>> s = pd.Series([0, 3, 0, 4])
```

```
>>> s.nonzero()
```

```
(array([1, 3]),)
```

```
>>> s.iloc[s.nonzero()[0]]
```

```
1    3
```

```
3    4
```

```
dtype: int64
```

```
>>> s = pd.Series([0, 3, 0, 4], index=['a', 'b', 'c', 'd'])
```

```
# same return although index of s is different
```

```
>>> s.nonzero()
```

```
(array([1, 3]),)
```

```
>>> s.iloc[s.nonzero()[0]]
```

```
b    3
```

```
d    4
```

```
dtype: int64
```

`notna(self)`

Detect existing (non-missing) values.

Return a boolean same-sized object indicating if the values are not NA. Non-missing values get mapped to True. Characters such as empty strings `''` or `:attr:`numpy.inf`` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``). NA values, such as `None` or `:attr:`numpy.NaN``, get mapped to False values.

Returns

Series

Mask of bool values for each element in Series that indicates whether an element is not an NA value.

See Also

`Series.notnull` : Alias of `notna`.

`Series.isna` : Boolean inverse of `notna`.

`Series.dropna` : Omit axes labels with missing values.

`notna` : Top-level `notna`.

Examples

Show which entries in a DataFrame are not NA.

```
>>> df = pd.DataFrame({'age': [5, 6, np.NaN],
...                    'born': [pd.NaT, pd.Timestamp('1939-05-27'),
...                             pd.Timestamp('1940-04-25')],
...                    'name': ['Alfred', 'Batman', ''],
...                    'toy': [None, 'Batmobile', 'Joker']})
```

```
>>> df
```

	age	born	name	toy
0	5.0	NaT	Alfred	None
1	6.0	1939-05-27	Batman	Batmobile
2	NaN	1940-04-25		Joker

```
>>> df.notna()
```

	age	born	name	toy
0	True	False	True	False
1	True	True	True	True
2	False	True	True	True

Show which entries in a Series are not NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
1    6.0
2    NaN
dtype: float64

>>> ser.notna()
0     True
1     True
2    False
dtype: bool
```

notnull(self)

Detect existing (non-missing) values.

Return a boolean same-sized object indicating if the values are not NA. Non-missing values get mapped to True. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``). NA values, such as None or :attr:`numpy.NaN`, get mapped to False values.

Returns

Series

Mask of bool values for each element in Series that indicates whether an element is not an NA value.

See Also

Series.notnull : Alias of notna.

Series.isna : Boolean inverse of notna.

Series.dropna : Omit axes labels with missing values.

notna : Top-level notna.

Examples

Show which entries in a DataFrame are not NA.

```
>>> df = pd.DataFrame({'age': [5, 6, np.NaN],
...                    'born': [pd.NaT, pd.Timestamp('1939-05-27'),
...                             pd.Timestamp('1940-04-25')],
...                    'name': ['Alfred', 'Batman', ''],
...                    'toy': [None, 'Batmobile', 'Joker']})
>>> df
   age      born   name      toy
0  5.0      NaT  Alfred     None
1  6.0  1939-05-27  Batman  Batmobile
2  NaN  1940-04-25     Joker     Joker
```

```
>>> df.notna()
   age      born   name      toy
0  True     False  True     False
1  True     True   True     True
2  False    True   True     True
```

Show which entries in a Series are not NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
1    6.0
2    NaN
dtype: float64

>>> ser.notna()
0     True
1     True
2    False
dtype: bool
```

```
nsmallest(self, n=5, keep='first')
    Return the smallest `n` elements.
```

Parameters

`n` : int, default 5

Return this many ascending sorted values.

`keep` : {'first', 'last', 'all'}, default 'first'

When there are duplicate values that cannot all fit in a Series of `n` elements:

- ``first`` : take the first occurrences based on the index order
- ``last`` : take the last occurrences based on the index order
- ``all`` : keep all occurrences. This can result in a Series of size larger than `n`.

Returns

Series

The `n` smallest values in the Series, sorted in increasing order.

See Also

`Series.nlargest`: Get the `n` largest elements.

`Series.sort_values`: Sort Series by values.

`Series.head`: Return the first `n` rows.

Notes

Faster than ``.sort_values().head(n)`` for small `n` relative to the size of the ``Series`` object.

Examples

```
>>> countries_population = {"Italy": 59000000, "France": 65000000,
...                          "Brunei": 434000, "Malta": 434000,
...                          "Maldives": 434000, "Iceland": 337000,
...                          "Nauru": 11300, "Tuvalu": 11300,
...                          "Anguilla": 11300, "Monserat": 5200}
```

```
>>> s = pd.Series(countries_population)
```

```
>>> s
```

```
Italy      59000000
France     65000000
Brunei      434000
Malta       434000
Maldives    434000
Iceland     337000
Nauru       11300
Tuvalu      11300
Anguilla    11300
Monserat    5200
```

```
dtype: int64
```

The `n` largest elements where ``n=5`` by default.

```
>>> s.nsmallest()
```

```
Monserat    5200
Nauru       11300
Tuvalu      11300
Anguilla    11300
Iceland     337000
```

```
dtype: int64
```

The `n` smallest elements where ``n=3``. Default `keep` value is 'first' so Nauru and Tuvalu will be kept.

```
>>> s.nsmallest(3)
```

```
Monserat    5200
Nauru       11300
Tuvalu      11300
```

```
dtype: int64
```

The `n` smallest elements where ``n=3`` and keeping the last duplicates. Anguilla and Tuvalu will be kept since they are the last with value 11300 based on the index order.

```
>>> s.nsmallest(3, keep='last')
Monserat      5200
Anguilla      11300
Tuvalu        11300
dtype: int64
```

The `n` smallest elements where ``n=3`` with all duplicates kept. Note that the returned Series has four elements due to the three duplicates.

```
>>> s.nsmallest(3, keep='all')
Monserat      5200
Nauru         11300
Tuvalu        11300
Anguilla      11300
dtype: int64
```

`pow(self, other, level=None, fill_value=None, axis=0)`

Exponential power of series and other, element-wise (binary operator `pow`).

Equivalent to ``series ** other``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

result : Series

See Also

Series.rpow

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a    1.0
```

```
b    1.0
```

```
c    1.0
```

```
d    NaN
```

```
dtype: float64
```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
```

```
>>> b
```

```
a    1.0
```

```
b    NaN
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```
>>> a.add(b, fill_value=0)
```

```
a    2.0
```

```
b    1.0
```

```
c    1.0
```

```
d    1.0
```

```
e    NaN
```

```
dtype: float64
```

```

| prod(self, axis=None, skipna=None, level=None, numeric_only=None, min_count=0, **kwar
gs)
|
|     Return the product of the values for the requested axis.
|
|     Parameters
|     -----
|     axis : {index (0)}
|         Axis for the function to be applied on.
|     skipna : bool, default True
|         Exclude NA/null values when computing the result.
|     level : int or level name, default None
|         If the axis is a MultiIndex (hierarchical), count along a
|         particular level, collapsing into a scalar.
|     numeric_only : bool, default None
|         Include only float, int, boolean columns. If None, will attempt to use
|         everything, then use only numeric data. Not implemented for Series.
|     min_count : int, default 0
|         The required number of valid values to perform the operation. If fewer than
|         ``min_count`` non-NA values are present the result will be NA.
|
|     .. versionadded :: 0.22.0
|
|         Added with the default being 0. This means the sum of an all-NA
|         or empty Series is 0, and the product of an all-NA or empty
|         Series is 1.
|
|     **kwargs
|         Additional keyword arguments to be passed to the function.
|
|     Returns
|     -----
|     prod : scalar or Series (if level specified)
|
|     Examples
|     -----
|     By default, the product of an empty or all-NA Series is ``1``
|
|     >>> pd.Series([]).prod()
|     1.0
|
|     This can be controlled with the ``min_count`` parameter
|
|     >>> pd.Series([]).prod(min_count=1)
|     nan
|
|     Thanks to the ``skipna`` parameter, ``min_count`` handles all-NA and
|     empty series identically.
|
|     >>> pd.Series([np.nan]).prod()
|     1.0
|
|     >>> pd.Series([np.nan]).prod(min_count=1)
|     nan
|
|     product = prod(self, axis=None, skipna=None, level=None, numeric_only=None, min_count
=0, **kwargs)
|
|     ptp(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
|         Returns the difference between the maximum value and the
|         minimum value in the object. This is the equivalent of the
|         ``numpy.ndarray`` method ``ptp``.
|
|     .. deprecated:: 0.24.0
|         Use numpy.ptp instead
|
|     Parameters
|     -----
|     axis : {index (0)}
|         Axis for the function to be applied on.
|     skipna : bool, default True
|         Exclude NA/null values when computing the result.
|     level : int or level name, default None
|         If the axis is a MultiIndex (hierarchical), count along a

```

particular level, collapsing into a scalar.
numeric_only : bool, default None
Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

****kwargs**

Additional keyword arguments to be passed to the function.

Returns

ptp : scalar or Series (if level specified)

put(self, *args, **kwargs)

Applies the `put` method to its `values` attribute if it has one.

See Also

numpy.ndarray.put

quantile(self, q=0.5, interpolation='linear')

Return value at the given quantile.

Parameters

q : float or array-like, default 0.5 (50% quantile)

0 <= q <= 1, the quantile(s) to compute

interpolation : {'linear', 'lower', 'higher', 'midpoint', 'nearest'}

.. versionadded:: 0.18.0

This optional parameter specifies the interpolation method to use, when the desired quantile lies between two data points `i` and `j`:

- * linear: $i + (j - i) * \text{fraction}$, where `fraction` is the fractional part of the index surrounded by `i` and `j`.

- * lower: `i`.

- * higher: `j`.

- * nearest: `i` or `j` whichever is nearest.

- * midpoint: $(i + j) / 2$.

Returns

quantile : float or Series

if ``q`` is an array, a Series will be returned where the index is ``q`` and the values are the quantiles.

See Also

core.window.Rolling.quantile

numpy.percentile

Examples

```
>>> s = pd.Series([1, 2, 3, 4])
```

```
>>> s.quantile(.5)
```

```
2.5
```

```
>>> s.quantile([.25, .5, .75])
```

```
0.25    1.75
```

```
0.50    2.50
```

```
0.75    3.25
```

```
dtype: float64
```

radd(self, other, level=None, fill_value=None, axis=0)

Addition of series and other, element-wise (binary operator `radd`).

Equivalent to ``other + series``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing
the result will be missing
level : int or name
Broadcast across a level, matching Index values on the
passed MultiIndex level

Returns

result : Series

See Also

Series.add

Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a    1.0
b    1.0
c    1.0
d    NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a    1.0
b    NaN
d    1.0
e    NaN
dtype: float64
>>> a.add(b, fill_value=0)
a    2.0
b    1.0
c    1.0
d    1.0
e    NaN
dtype: float64
```

ravel(self, order='C')

Return the flattened underlying data as an ndarray.

See Also

numpy.ndarray.ravel

rdiv = rtruediv(self, other, level=None, fill_value=None, axis=0)

rdivmod(self, other, level=None, fill_value=None, axis=0)

Integer division and modulo of series and other, element-wise (binary operator `rdivmod`).

Equivalent to ``other divmod series``, but with support to substitute a fill_value for missing data in one of the inputs.

Parameters

other : Series or scalar value

fill_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing
the result will be missing

level : int or name

Broadcast across a level, matching Index values on the
passed MultiIndex level

Returns

result : Series

See Also

```
-----  
Series.divmod
```

Examples

```
-----  
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])  
>>> a  
a    1.0  
b    1.0  
c    1.0  
d    NaN  
dtype: float64  
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])  
>>> b  
a    1.0  
b    NaN  
d    1.0  
e    NaN  
dtype: float64  
>>> a.add(b, fill_value=0)  
a    2.0  
b    1.0  
c    1.0  
d    1.0  
e    NaN  
dtype: float64
```

```
reindex(self, index=None, **kwargs)
```

Conform Series to new index with optional filling logic, placing NA/NaN in locations having no value in the previous index. A new object is produced unless the new index is equivalent to the current one and ``copy=False``.

Parameters

```
-----
```

index : array-like, optional

New labels / index to conform to, should be specified using keywords. Preferably an Index object to avoid duplicating data

method : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}

Method to use for filling holes in reindexed DataFrame.

Please note: this is only applicable to DataFrames/Series with a monotonically increasing/decreasing index.

- * None (default): don't fill gaps

- * pad / ffill: propagate last valid observation forward to next valid

- * backfill / bfill: use next valid observation to fill gap

- * nearest: use nearest valid observations to fill gap

copy : bool, default True

Return a new object, even if the passed indexes are the same.

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level.

fill_value : scalar, default np.NaN

Value to use for missing values. Defaults to NaN, but can be any "compatible" value.

limit : int, default None

Maximum number of consecutive elements to forward or backward fill.

tolerance : optional

Maximum distance between original and new labels for inexact matches. The values of the index at the matching locations must satisfy the equation ``abs(index[indexer] - target) <= tolerance``.

Tolerance may be a scalar value, which applies the same tolerance to all values, or list-like, which applies variable tolerance per element. List-like includes list, tuple, array, Series, and must be the same size as the index and its dtype must exactly match the index's type.

```
.. versionadded:: 0.21.0 (list-like tolerance)
```

Returns

Series with changed index.

See Also

DataFrame.set_index : Set row labels.

DataFrame.reset_index : Remove row labels or move them to new columns.

DataFrame.reindex_like : Change to same indices as other DataFrame.

Examples

```DataFrame.reindex``` supports two calling conventions

\* ```(index=index_labels, columns=column_labels, ...)```

\* ```(labels, axis={'index', 'columns'}, ...)```

We *highly* recommend using keyword arguments to clarify your intent.

Create a dataframe with some fictional data.

```
>>> index = ['Firefox', 'Chrome', 'Safari', 'IE10', 'Konqueror']
>>> df = pd.DataFrame({
... 'http_status': [200, 200, 404, 404, 301],
... 'response_time': [0.04, 0.02, 0.07, 0.08, 1.0]},
... index=index)
>>> df
```

	http_status	response_time
Firefox	200	0.04
Chrome	200	0.02
Safari	404	0.07
IE10	404	0.08
Konqueror	301	1.00

Create a new index and reindex the dataframe. By default values in the new index that do not have corresponding records in the dataframe are assigned ```NaN```.

```
>>> new_index= ['Safari', 'Iceweasel', 'Comodo Dragon', 'IE10',
... 'Chrome']
>>> df.reindex(new_index)
```

	http_status	response_time
Safari	404.0	0.07
Iceweasel	NaN	NaN
Comodo Dragon	NaN	NaN
IE10	404.0	0.08
Chrome	200.0	0.02

We can fill in the missing values by passing a value to the keyword ```fill_value```. Because the index is not monotonically increasing or decreasing, we cannot use arguments to the keyword ```method``` to fill the ```NaN``` values.

```
>>> df.reindex(new_index, fill_value=0)
```

	http_status	response_time
Safari	404	0.07
Iceweasel	0	0.00
Comodo Dragon	0	0.00
IE10	404	0.08
Chrome	200	0.02

```
>>> df.reindex(new_index, fill_value='missing')
```

	http_status	response_time
Safari	404	0.07
Iceweasel	missing	missing
Comodo Dragon	missing	missing
IE10	404	0.08
Chrome	200	0.02

We can also reindex the columns.

```
>>> df.reindex(columns=['http_status', 'user_agent'])
 http_status user_agent
Firefox 200 NaN
Chrome 200 NaN
Safari 404 NaN
IE10 404 NaN
Konqueror 301 NaN
```

Or we can use "axis-style" keyword arguments

```
>>> df.reindex(['http_status', 'user_agent'], axis="columns")
 http_status user_agent
Firefox 200 NaN
Chrome 200 NaN
Safari 404 NaN
IE10 404 NaN
Konqueror 301 NaN
```

To further illustrate the filling functionality in ``reindex``, we will create a dataframe with a monotonically increasing index (for example, a sequence of dates).

```
>>> date_index = pd.date_range('1/1/2010', periods=6, freq='D')
>>> df2 = pd.DataFrame({"prices": [100, 101, np.nan, 100, 89, 88]},
... index=date_index)
>>> df2
 prices
2010-01-01 100.0
2010-01-02 101.0
2010-01-03 NaN
2010-01-04 100.0
2010-01-05 89.0
2010-01-06 88.0
```

Suppose we decide to expand the dataframe to cover a wider date range.

```
>>> date_index2 = pd.date_range('12/29/2009', periods=10, freq='D')
>>> df2.reindex(date_index2)
 prices
2009-12-29 NaN
2009-12-30 NaN
2009-12-31 NaN
2010-01-01 100.0
2010-01-02 101.0
2010-01-03 NaN
2010-01-04 100.0
2010-01-05 89.0
2010-01-06 88.0
2010-01-07 NaN
```

The index entries that did not have a value in the original data frame (for example, '2009-12-29') are by default filled with ``NaN``. If desired, we can fill in the missing values using one of several options.

For example, to back-propagate the last valid value to fill the ``NaN`` values, pass ``bfill`` as an argument to the ``method`` keyword.

```
>>> df2.reindex(date_index2, method='bfill')
 prices
2009-12-29 100.0
2009-12-30 100.0
2009-12-31 100.0
2010-01-01 100.0
2010-01-02 101.0
2010-01-03 NaN
2010-01-04 100.0
```

```

2010-01-05 89.0
2010-01-06 88.0
2010-01-07 NaN

```

Please note that the ``NaN`` value present in the original dataframe (at index value 2010-01-03) will not be filled by any of the value propagation schemes. This is because filling while reindexing does not look at dataframe values, but only compares the original and desired indexes. If you do want to fill in the ``NaN`` values present in the original dataframe, use the ``fillna()`` method.

See the :ref:`user guide <basics.reindexing>` for more.

```

reindex_axis(self, labels, axis=0, **kwargs)
 Conform Series to new index with optional filling logic.

```

```

.. deprecated:: 0.21.0
 Use ``Series.reindex`` instead.

```

```

rename(self, index=None, **kwargs)
 Alter Series index labels or name.

```

Function / dict values must be unique (1-to-1). Labels not contained in a dict / Series will be left as-is. Extra labels listed don't throw an error.

Alternatively, change ``Series.name`` with a scalar value.

See the :ref:`user guide <basics.rename>` for more.

#### Parameters

```

index : scalar, hashable sequence, dict-like or function, optional
 dict-like or functions are transformations to apply to
 the index.
 Scalar or hashable sequence-like will alter the ``Series.name``
 attribute.
copy : bool, default True
 Also copy underlying data
inplace : bool, default False
 Whether to return a new Series. If True then value of copy is
 ignored.
level : int or level name, default None
 In case of a MultiIndex, only rename labels in the specified
 level.

```

#### Returns

```

renamed : Series (new object)

```

#### See Also

```

Series.rename_axis

```

#### Examples

```

>>> s = pd.Series([1, 2, 3])
>>> s
0 1
1 2
2 3
dtype: int64
>>> s.rename("my_name") # scalar, changes Series.name
0 1
1 2
2 3
Name: my_name, dtype: int64
>>> s.rename(lambda x: x ** 2) # function, changes labels
0 1
1 2
4 3
dtype: int64

```

```
>>> s.rename({1: 3, 2: 5}) # mapping, changes labels
```

```
0 1
```

```
3 2
```

```
5 3
```

```
dtype: int64
```

```
reorder_levels(self, order)
```

Rearrange index levels using input order.

May not drop or duplicate levels.

Parameters

-----

order : list of int representing new level order  
(reference level by number or key)

Returns

-----

type of caller (new object)

```
repeat(self, repeats, axis=None)
```

Repeat elements of a Series.

Returns a new Series where each element of the current Series  
is repeated consecutively a given number of times.

Parameters

-----

repeats : int or array of ints

The number of repetitions for each element. This should be a  
non-negative integer. Repeating 0 times will return an empty  
Series.

axis : None

Must be ``None``. Has no effect but is accepted for compatibility  
with numpy.

Returns

-----

repeated\_series : Series

Newly created Series with repeated elements.

See Also

-----

Index.repeat : Equivalent function for Index.

numpy.repeat : Similar method for :class:`numpy.ndarray`.

Examples

-----

```
>>> s = pd.Series(['a', 'b', 'c'])
```

```
>>> s
```

```
0 a
```

```
1 b
```

```
2 c
```

```
dtype: object
```

```
>>> s.repeat(2)
```

```
0 a
```

```
0 a
```

```
1 b
```

```
1 b
```

```
2 c
```

```
2 c
```

```
dtype: object
```

```
>>> s.repeat([1, 2, 3])
```

```
0 a
```

```
1 b
```

```
1 b
```

```
2 c
```

```
2 c
```

```
2 c
```

```
dtype: object
```

```
replace(self, to_replace=None, value=None, inplace=False, limit=None, regex=False, me
```

```
thod='pad')
```

```
 Replace values given in `to_replace` with `value`.
```

```
 Values of the Series are replaced with other values dynamically.
 This differs from updating with ``.loc`` or ``.iloc``, which require
 you to specify a location to update with some value.
```

```
Parameters
```

```

```

```
to_replace : str, regex, list, dict, Series, int, float, or None
```

```
 How to find the values that will be replaced.
```

```
 * numeric, str or regex:
```

- numeric: numeric values equal to `to\_replace` will be replaced with `value`
- str: string exactly matching `to\_replace` will be replaced with `value`
- regex: regexs matching `to\_replace` will be replaced with `value`

```
 * list of str, regex, or numeric:
```

- First, if `to\_replace` and `value` are both lists, they **must** be the same length.
- Second, if ``regex=True`` then all of the strings in **both** lists will be interpreted as regexs otherwise they will match directly. This doesn't matter much for `value` since there are only a few possible substitution regexes you can use.
- str, regex and numeric rules apply as above.

```
 * dict:
```

- Dicts can be used to specify different replacement values for different existing values. For example, ``{'a': 'b', 'y': 'z'}`` replaces the value 'a' with 'b' and 'y' with 'z'. To use a dict in this way the `value` parameter should be `None`.
- For a DataFrame a dict can specify that different values should be replaced in different columns. For example, ``{'a': 1, 'b': 'z'}`` looks for the value 1 in column 'a' and the value 'z' in column 'b' and replaces these values with whatever is specified in `value`. The `value` parameter should not be ``None`` in this case. You can treat this as a special case of passing two lists except that you are specifying the column to search in.
- For a DataFrame nested dictionaries, e.g., ``{'a': {'b': np.nan}}``, are read as follows: look in column 'a' for the value 'b' and replace it with NaN. The `value` parameter should be ``None`` to use a nested dict in this way. You can nest regular expressions as well. Note that column names (the top-level dictionary keys in a nested dictionary) **cannot** be regular expressions.

```
 * None:
```

- This means that the `regex` argument must be a string, compiled regular expression, or list, dict, ndarray or Series of such elements. If `value` is also ``None`` then this **must** be a nested dictionary or Series.

```
 See the examples section for examples of each of these.
```

```
value : scalar, dict, list, str, regex, default None
```

```
 Value to replace any values matching `to_replace` with.
```

```
 For a DataFrame a dict of values can be used to specify which
 value to use for each column (columns not in the dict will not be
 filled). Regular expressions, strings and lists or dicts of such
 objects are also allowed.
```

```
inplace : bool, default False
```

```
 If True, in place. Note: this will modify any
```

```
 other views on this object (e.g. a column from a DataFrame).
```

```
 Returns the caller if this is True.
```

```

limit : int, default None
 Maximum size gap to forward or backward fill.
regex : bool or same types as `to_replace`, default False
 Whether to interpret `to_replace` and/or `value` as regular
 expressions. If this is ``True`` then `to_replace` *must* be a
 string. Alternatively, this could be a regular expression or a
 list, dict, or array of regular expressions in which case
 `to_replace` must be ``None``.
method : {'pad', 'ffill', 'bfill', 'None'}
 The method to use when for replacement, when `to_replace` is a
 scalar, list or tuple and `value` is ``None``.

 .. versionchanged:: 0.23.0
 Added to DataFrame.

Returns

Series
 Object after replacement.

Raises

AssertionError
 * If `regex` is not a ``bool`` and `to_replace` is not
 ``None``.
TypeError
 * If `to_replace` is a ``dict`` and `value` is not a ``list``,
 ``dict``, ``ndarray``, or ``Series``
 * If `to_replace` is ``None`` and `regex` is not compilable
 into a regular expression or is a list, dict, ndarray, or
 Series.
 * When replacing multiple ``bool`` or ``datetime64`` objects and
 the arguments to `to_replace` does not match the type of the
 value being replaced
ValueError
 * If a ``list`` or an ``ndarray`` is passed to `to_replace` and
 `value` but they are not the same length.

See Also

Series.fillna : Fill NA values.
Series.where : Replace values based on boolean condition.
Series.str.replace : Simple string replacement.

Notes

* Regex substitution is performed under the hood with ``re.sub``. The
 rules for substitution for ``re.sub`` are the same.
* Regular expressions will only substitute on strings, meaning you
 cannot provide, for example, a regular expression matching floating
 point numbers and expect the columns in your frame that have a
 numeric dtype to be matched. However, if those floating point
 numbers *are* strings, then you can do this.
* This method has *a lot* of options. You are encouraged to experiment
 and play with this method to gain intuition about how it works.
* When dict is used as the `to_replace` value, it is like
 key(s) in the dict are the to_replace part and
 value(s) in the dict are the value parameter.

Examples

Scalar `to_replace` and `value`

>>> s = pd.Series([0, 1, 2, 3, 4])
>>> s.replace(0, 5)
0 5
1 1
2 2
3 3
4 4
dtype: int64

```



```

>>> df = pd.DataFrame({'A': [0, 1, 2, 3, 4],
... 'B': [5, 6, 7, 8, 9],
... 'C': ['a', 'b', 'c', 'd', 'e']})
>>> df.replace(0, 5)
 A B C
0 5 5 a
1 1 6 b
2 2 7 c
3 3 8 d
4 4 9 e

List-like `to_replace`

>>> df.replace([0, 1, 2, 3], 4)
 A B C
0 4 5 a
1 4 6 b
2 4 7 c
3 4 8 d
4 4 9 e

>>> df.replace([0, 1, 2, 3], [4, 3, 2, 1])
 A B C
0 4 5 a
1 3 6 b
2 2 7 c
3 1 8 d
4 4 9 e

>>> s.replace([1, 2], method='bfill')
0 0
1 3
2 3
3 3
4 4
dtype: int64

dict-like `to_replace`

>>> df.replace({0: 10, 1: 100})
 A B C
0 10 5 a
1 100 6 b
2 2 7 c
3 3 8 d
4 4 9 e

>>> df.replace({'A': 0, 'B': 5}, 100)
 A B C
0 100 100 a
1 1 6 b
2 2 7 c
3 3 8 d
4 4 9 e

>>> df.replace({'A': {0: 100, 4: 400}})
 A B C
0 100 5 a
1 1 6 b
2 2 7 c
3 3 8 d
4 400 9 e

Regular expression `to_replace`

>>> df = pd.DataFrame({'A': ['bat', 'foo', 'bait'],
... 'B': ['abc', 'bar', 'xyz']})
>>> df.replace(to_replace=r'^ba.$', value='new', regex=True)
 A B
0 new abc
1 foo new

```

```

2 bait xyz

>>> df.replace({'A': r'^ba.$'}, {'A': 'new'}, regex=True)
 A B
0 new abc
1 foo bar
2 bait xyz

>>> df.replace(regex=r'^ba.$', value='new')
 A B
0 new abc
1 foo new
2 bait xyz

```

```

>>> df.replace(regex={r'^ba.$': 'new', 'foo': 'xyz'})
 A B
0 new abc
1 xyz new
2 bait xyz

>>> df.replace(regex=[r'^ba.$', 'foo'], value='new')
 A B
0 new abc
1 new new
2 bait xyz

```

Note that when replacing multiple ``bool`` or ``datetime64`` objects, the data types in the `to\_replace` parameter must match the data type of the value being replaced:

```

>>> df = pd.DataFrame({'A': [True, False, True],
... 'B': [False, True, False]})
>>> df.replace({'a string': 'new value', True: False}) # raises
Traceback (most recent call last):

```

```

...
TypeError: Cannot compare types 'ndarray(dtype=bool)' and 'str'

```

This raises a ``TypeError`` because one of the ``dict`` keys is not of the correct type for replacement.

Compare the behavior of ``s.replace({'a': None})`` and ``s.replace('a', None)`` to understand the peculiarities of the `to\_replace` parameter:

```

>>> s = pd.Series([10, 'a', 'a', 'b', 'a'])

```

When one uses a dict as the `to\_replace` value, it is like the value(s) in the dict are equal to the `value` parameter.

```

`s.replace({'a': None})` is equivalent to
`s.replace(to_replace={'a': None}, value=None, method=None)`:

```

```

>>> s.replace({'a': None})
0 10
1 None
2 None
3 b
4 None
dtype: object

```

When ``value=None`` and `to\_replace` is a scalar, list or tuple, `replace` uses the method parameter (default 'pad') to do the replacement. So this is why the 'a' values are being replaced by 10 in rows 1 and 2 and 'b' in row 4 in this case.

The command ``s.replace('a', None)`` is actually equivalent to ``s.replace(to\_replace='a', value=None, method='pad')``:

```

>>> s.replace('a', None)
0 10
1 10
2 10
3 b
4 b

```

dtype: object

```
reset_index(self, level=None, drop=False, name=None, inplace=False)
 Generate a new DataFrame or Series with the index reset.
```

This is useful when the index needs to be treated as a column, or when the index is meaningless and needs to be reset to the default before another operation.

#### Parameters

-----

level : int, str, tuple, or list, default optional  
For a Series with a MultiIndex, only remove the specified levels from the index. Removes all levels by default.

drop : bool, default False  
Just reset the index, without inserting it as a column in the new DataFrame.

name : object, optional  
The name to use for the column containing the original Series values. Uses ``self.name`` by default. This argument is ignored when `drop` is True.

inplace : bool, default False  
Modify the Series in place (do not create a new object).

#### Returns

-----

##### Series or DataFrame

When `drop` is False (the default), a DataFrame is returned. The newly created columns will come first in the DataFrame, followed by the original Series values.

When `drop` is True, a `Series` is returned.

In either case, if ``inplace=True``, no value is returned.

#### See Also

-----

DataFrame.reset\_index: Analogous function for DataFrame.

#### Examples

-----

```
>>> s = pd.Series([1, 2, 3, 4], name='foo',
... index=pd.Index(['a', 'b', 'c', 'd'], name='idx'))
```

Generate a DataFrame with default index.

```
>>> s.reset_index()
 idx foo
0 a 1
1 b 2
2 c 3
3 d 4
```

To specify the name of the new column use `name`.

```
>>> s.reset_index(name='values')
 idx values
0 a 1
1 b 2
2 c 3
3 d 4
```

To generate a new Series with the default set `drop` to True.

```
>>> s.reset_index(drop=True)
0 1
1 2
2 3
3 4
Name: foo, dtype: int64
```

To update the Series in place, without generating a new one set `inplace` to True. Note that it also requires ``drop=True``.

```
>>> s.reset_index(inplace=True, drop=True)
>>> s
0 1
1 2
2 3
3 4
Name: foo, dtype: int64
```

The ``level`` parameter is interesting for Series with a multi-level index.

```
>>> arrays = [np.array(['bar', 'bar', 'baz', 'baz']),
... np.array(['one', 'two', 'one', 'two'])]
>>> s2 = pd.Series(
... range(4), name='foo',
... index=pd.MultiIndex.from_arrays(arrays,
... names=['a', 'b']))
```

To remove a specific level from the Index, use ``level``.

```
>>> s2.reset_index(level='a')
 a foo
b
one bar 0
two bar 1
one baz 2
two baz 3
```

If ``level`` is not set, all levels are removed from the Index.

```
>>> s2.reset_index()
 a b foo
0 bar one 0
1 bar two 1
2 baz one 2
3 baz two 3
```

`rfloordiv(self, other, level=None, fill_value=None, axis=0)`

Integer division of series and other, element-wise (binary operator ``rfloordiv``).

Equivalent to ``other // series``, but with support to substitute a `fill_value` for missing data in one of the inputs.

Parameters

-----

`other` : Series or scalar value

`fill_value` : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

`level` : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

-----

result : Series

See Also

-----

`Series.floordiv`

Examples

-----

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a 1.0
```

```
b 1.0
```

```
c 1.0
```

```
d NaN
```

```

dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a 1.0
b NaN
d 1.0
e NaN
dtype: float64
>>> a.add(b, fill_value=0)
a 2.0
b 1.0
c 1.0
d 1.0
e NaN
dtype: float64

```

```

rmod(self, other, level=None, fill_value=None, axis=0)

```

Modulo of series and other, element-wise (binary operator ``rmod``).

Equivalent to ``other % series``, but with support to substitute a `fill_value` for missing data in one of the inputs.

Parameters

-----

`other` : Series or scalar value

`fill_value` : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

`level` : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

-----

result : Series

See Also

-----

`Series.mod`

Examples

-----

```

>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a 1.0
b 1.0
c 1.0
d NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a 1.0
b NaN
d 1.0
e NaN
dtype: float64
>>> a.add(b, fill_value=0)
a 2.0
b 1.0
c 1.0
d 1.0
e NaN
dtype: float64

```

```

rmul(self, other, level=None, fill_value=None, axis=0)

```

Multiplication of series and other, element-wise (binary operator ``rmul``).

Equivalent to ``other * series``, but with support to substitute a `fill_value` for missing data in one of the inputs.

## Parameters

other : Series or scalar value

fill\_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

## Returns

result : Series

## See Also

Series.mul

## Examples

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a 1.0
```

```
b 1.0
```

```
c 1.0
```

```
d NaN
```

```
dtype: float64
```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
```

```
>>> b
```

```
a 1.0
```

```
b NaN
```

```
d 1.0
```

```
e NaN
```

```
dtype: float64
```

```
>>> a.add(b, fill_value=0)
```

```
a 2.0
```

```
b 1.0
```

```
c 1.0
```

```
d 1.0
```

```
e NaN
```

```
dtype: float64
```

```
rolling(self, window, min_periods=None, center=False, win_type=None, on=None, axis=0, closed=None)
```

Provides rolling window calculations.

.. versionadded:: 0.18.0

## Parameters

window : int, or offset

Size of the moving window. This is the number of observations used for calculating the statistic. Each window will be a fixed size.

If its an offset then this will be the time period of each window. Each window will be a variable sized based on the observations included in the time-period. This is only valid for datetimelike indexes. This is new in 0.19.0

min\_periods : int, default None

Minimum number of observations in window required to have a value (otherwise result is NA). For a window that is specified by an offset, 'min\_periods' will default to 1. Otherwise, 'min\_periods' will default to the size of the window.

center : bool, default False

Set the labels at the center of the window.

win\_type : str, default None

Provide a window type. If ``None``, all points are evenly weighted.

See the notes below for further information.

on : str, optional

For a DataFrame, column on which to calculate

the rolling window, rather than the index  
axis : int or str, default 0  
closed : str, default None  
    Make the interval closed on the 'right', 'left', 'both' or  
    'neither' endpoints.  
    For offset-based windows, it defaults to 'right'.  
    For fixed windows, defaults to 'both'. Remaining cases not implemented  
    for fixed windows.  
  
.. versionadded:: 0.20.0

#### Returns

-----  
a Window or Rolling sub-classed for the particular operation

#### See Also

-----  
expanding : Provides expanding transformations.  
ewm : Provides exponential weighted functions.

#### Notes

-----  
By default, the result is set to the right edge of the window. This can be  
changed to the center of the window by setting ``center=True``.

To learn more about the offsets & frequency strings, please see `this link  
<<http://pandas.pydata.org/pandas-docs/stable/timeseries.html#offset-aliases>>`\_\_.

The recognized win\_types are:

- \* ``boxcar``
- \* ``triang``
- \* ``blackman``
- \* ``hamming``
- \* ``bartlett``
- \* ``parzen``
- \* ``bohman``
- \* ``blackmanharris``
- \* ``nuttall``
- \* ``barthann``
- \* ``kaiser`` (needs beta)
- \* ``gaussian`` (needs std)
- \* ``general\_gaussian`` (needs power, width)
- \* ``slepian`` (needs width).

If ``win\_type=None`` all points are evenly weighted. To learn more about  
different window types see `scipy.signal window functions  
<<https://docs.scipy.org/doc/scipy/reference/signal.html#window-functions>>`\_\_.

#### Examples

-----  
  
>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})  
>>> df  
      B  
0 0.0  
1 1.0  
2 2.0  
3 NaN  
4 4.0

Rolling sum with a window length of 2, using the 'triang'  
window type.

>>> df.rolling(2, win\_type='triang').sum()  
      B  
0 NaN  
1 1.0  
2 2.5  
3 NaN  
4 NaN

Rolling sum with a window length of 2, min\_periods defaults to the window length.

```
>>> df.rolling(2).sum()
 B
0 NaN
1 1.0
2 3.0
3 NaN
4 NaN
```

Same as above, but explicitly set the min\_periods

```
>>> df.rolling(2, min_periods=1).sum()
 B
0 0.0
1 1.0
2 3.0
3 2.0
4 4.0
```

A ragged (meaning not-a-regular frequency), time-indexed DataFrame

```
>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]},
... index = [pd.Timestamp('20130101 09:00:00'),
... pd.Timestamp('20130101 09:00:02'),
... pd.Timestamp('20130101 09:00:03'),
... pd.Timestamp('20130101 09:00:05'),
... pd.Timestamp('20130101 09:00:06')])
```

```
>>> df
 B
2013-01-01 09:00:00 0.0
2013-01-01 09:00:02 1.0
2013-01-01 09:00:03 2.0
2013-01-01 09:00:05 NaN
2013-01-01 09:00:06 4.0
```

Contrasting to an integer rolling window, this will roll a variable length window corresponding to the time period. The default for min\_periods is 1.

```
>>> df.rolling('2s').sum()
 B
2013-01-01 09:00:00 0.0
2013-01-01 09:00:02 1.0
2013-01-01 09:00:03 3.0
2013-01-01 09:00:05 NaN
2013-01-01 09:00:06 4.0
```

```
round(self, decimals=0, *args, **kwargs)
Round each value in a Series to the given number of decimals.
```

Parameters

-----

decimals : int

Number of decimal places to round to (default: 0).

If decimals is negative, it specifies the number of positions to the left of the decimal point.

Returns

-----

Series object

See Also

-----

numpy.around

DataFrame.round

```
rpow(self, other, level=None, fill_value=None, axis=0)
Exponential power of series and other, element-wise (binary operator `rpow`).
```



Equivalent to ``other \*\* series``, but with support to substitute a fill\_value for missing data in one of the inputs.

#### Parameters

-----

other : Series or scalar value

fill\_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

#### Returns

-----

result : Series

#### See Also

-----

Series.pow

#### Examples

-----

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
```

```
>>> a
```

```
a 1.0
```

```
b 1.0
```

```
c 1.0
```

```
d NaN
```

```
dtype: float64
```

```
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
```

```
>>> b
```

```
a 1.0
```

```
b NaN
```

```
d 1.0
```

```
e NaN
```

```
dtype: float64
```

```
>>> a.add(b, fill_value=0)
```

```
a 2.0
```

```
b 1.0
```

```
c 1.0
```

```
d 1.0
```

```
e NaN
```

```
dtype: float64
```

```
rsub(self, other, level=None, fill_value=None, axis=0)
```

Subtraction of series and other, element-wise (binary operator `rsub`).

Equivalent to ``other - series``, but with support to substitute a fill\_value for missing data in one of the inputs.

#### Parameters

-----

other : Series or scalar value

fill\_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

#### Returns

-----

result : Series

#### See Also

-----

Series.sub

#### Examples

-----

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a 1.0
b 1.0
c 1.0
d NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a 1.0
b NaN
d 1.0
e NaN
dtype: float64
>>> a.add(b, fill_value=0)
a 2.0
b 1.0
c 1.0
d 1.0
e NaN
dtype: float64
```

rtruediv(self, other, level=None, fill\_value=None, axis=0)

Floating division of series and other, element-wise (binary operator `rtruediv`).

Equivalent to ``other / series``, but with support to substitute a fill\_value for missing data in one of the inputs.

#### Parameters

-----

other : Series or scalar value

fill\_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

#### Returns

-----

result : Series

#### See Also

-----

Series.truediv

#### Examples

-----

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a 1.0
b 1.0
c 1.0
d NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a 1.0
b NaN
d 1.0
e NaN
dtype: float64
>>> a.add(b, fill_value=0)
a 2.0
b 1.0
c 1.0
```

```
d 1.0
e NaN
dtype: float64
```

```
searchsorted(self, value, side='left', sorter=None)
 Find indices where elements should be inserted to maintain order.
```

Find the indices into a sorted Series `self` such that, if the corresponding elements in `value` were inserted before the indices, the order of `self` would be preserved.

#### Parameters

-----

value : array\_like

Values to insert into `self`.

side : {'left', 'right'}, optional

If 'left', the index of the first suitable location found is given.

If 'right', return the last such index. If there is no suitable index, return either 0 or N (where N is the length of `self`).

sorter : 1-D array\_like, optional

Optional array of integer indices that sort `self` into ascending order. They are typically the result of ``np.argsort``.

#### Returns

-----

int or array of int

A scalar or array of insertion points with the same shape as `value`.

.. versionchanged :: 0.24.0

If `value` is a scalar, an int is now always returned.

Previously, scalar inputs returned an 1-item array for :class:`Series` and :class:`Categorical`.

#### See Also

-----

numpy.searchsorted

#### Notes

-----

Binary search is used to find the required insertion points.

#### Examples

-----

```
>>> x = pd.Series([1, 2, 3])
```

```
>>> x
```

```
0 1
```

```
1 2
```

```
2 3
```

```
dtype: int64
```

```
>>> x.searchsorted(4)
```

```
3
```

```
>>> x.searchsorted([0, 4])
```

```
array([0, 3])
```

```
>>> x.searchsorted([1, 3], side='left')
```

```
array([0, 2])
```

```
>>> x.searchsorted([1, 3], side='right')
```

```
array([1, 3])
```

```
>>> x = pd.Categorical(['apple', 'bread', 'bread',
 'cheese', 'milk'], ordered=True)
```

```
[apple, bread, bread, cheese, milk]
```

```
Categories (4, object): [apple < bread < cheese < milk]
```

```
>>> x.searchsorted('bread')
```

```
1
```

```
>>> x.searchsorted(['bread'], side='right')
array([3])
```

```
sem(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None, **kwargs)
 Return unbiased standard error of the mean over requested axis.
```

Normalized by N-1 by default. This can be changed using the ddof argument

#### Parameters

-----

axis : {index (0)}

skipna : boolean, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar

ddof : int, default 1

Delta Degrees of Freedom. The divisor used in calculations is N - ddof, where N represents the number of elements.

numeric\_only : boolean, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

#### Returns

-----

sem : scalar or Series (if level specified)

```
set_value(self, label, value, takeable=False)
```

Quickly set single value at passed label.

.. deprecated:: 0.21.0

Please use .at[] or .iat[] accessors.

If label is not contained, a new object is created with the label placed at the end of the result index.

#### Parameters

-----

label : object

Partial indexing with MultiIndex not allowed

value : object

Scalar value

takeable : interpret the index as indexers, default False

#### Returns

-----

series : Series

If label is contained, will be reference to calling Series, otherwise a new object

```
shift(self, periods=1, freq=None, axis=0, fill_value=None)
```

Shift index by desired number of periods with an optional time `freq`.

When `freq` is not passed, shift the index without realigning the data.

If `freq` is passed (in this case, the index must be date or datetime, or it will raise a `NotImplementedError`), the index will be

increased using the periods and the `freq`.

#### Parameters

-----

periods : int

Number of periods to shift. Can be positive or negative.

freq : DateOffset, tseries.offsets, timedelta, or str, optional

Offset to use from the tseries module or time rule (e.g. 'EOM').

If `freq` is specified then the index values are shifted but the data is not realigned. That is, use `freq` if you would like to extend the index when shifting and preserve the original data.

axis : {0 or 'index', 1 or 'columns', None}, default None

Shift direction.

fillvalue : object, optional

The scalar value to use for newly introduced missing values.

the default depends on the dtype of `self`.  
For numeric data, ``np.nan`` is used.  
For datetime, timedelta, or period data, etc. :attr:`NaT` is used.  
For extension dtypes, ``self.dtype.na\_value`` is used.

.. versionchanged:: 0.24.0

#### Returns

-----

##### Series

Copy of input object, shifted.

#### See Also

-----

Index.shift : Shift values of Index.

DatetimeIndex.shift : Shift values of DatetimeIndex.

PeriodIndex.shift : Shift values of PeriodIndex.

tshift : Shift the time index, using the index's frequency if available.

#### Examples

-----

```
>>> df = pd.DataFrame({'Col1': [10, 20, 15, 30, 45],
... 'Col2': [13, 23, 18, 33, 48],
... 'Col3': [17, 27, 22, 37, 52]})
```

```
>>> df.shift(periods=3)
```

	Col1	Col2	Col3
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	10.0	13.0	17.0
4	20.0	23.0	27.0

```
>>> df.shift(periods=1, axis='columns')
```

	Col1	Col2	Col3
0	NaN	10.0	13.0
1	NaN	20.0	23.0
2	NaN	15.0	18.0
3	NaN	30.0	33.0
4	NaN	45.0	48.0

```
>>> df.shift(periods=3, fill_value=0)
```

	Col1	Col2	Col3
0	0	0	0
1	0	0	0
2	0	0	0
3	10	13	17
4	20	23	27

skew(self, axis=None, skipna=None, level=None, numeric\_only=None, \*\*kwargs)  
Return unbiased skew over requested axis  
Normalized by N-1.

#### Parameters

-----

axis : {index (0)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar.

numeric\_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

\*\*kwargs

Additional keyword arguments to be passed to the function.

#### Returns

-----

skew : scalar or Series (if level specified)

```
| sort_index(self, axis=0, level=None, ascending=True, inplace=False, kind='quicksort',
| na_position='last', sort_remaining=True)
```

```
| Sort Series by index labels.
```

```
|
| Returns a new Series sorted by label if `inplace` argument is
| ``False``, otherwise updates the original series and returns None.
```

```
| Parameters
```

```
| -----
```

```
| axis : int, default 0
```

```
| Axis to direct sorting. This can only be 0 for Series.
```

```
| level : int, optional
```

```
| If not None, sort on values in specified index level(s).
```

```
| ascending : bool, default true
```

```
| Sort ascending vs. descending.
```

```
| inplace : bool, default False
```

```
| If True, perform operation in-place.
```

```
| kind : {'quicksort', 'mergesort', 'heapsort'}, default 'quicksort'
```

```
| Choice of sorting algorithm. See also :func:`numpy.sort` for more
| information. 'mergesort' is the only stable algorithm. For
| DataFrames, this option is only applied when sorting on a single
| column or label.
```

```
| na_position : {'first', 'last'}, default 'last'
```

```
| If 'first' puts NaNs at the beginning, 'last' puts NaNs at the end.
```

```
| Not implemented for MultiIndex.
```

```
| sort_remaining : bool, default True
```

```
| If true and sorting by level and index is multilevel, sort by other
| levels too (in order) after sorting by specified level.
```

```
| Returns
```

```
| -----
```

```
| pandas.Series
```

```
| The original Series sorted by the labels
```

```
| See Also
```

```
| -----
```

```
| DataFrame.sort_index: Sort DataFrame by the index.
```

```
| DataFrame.sort_values: Sort DataFrame by the value.
```

```
| Series.sort_values : Sort Series by the value.
```

```
| Examples
```

```
| -----
```

```
| >>> s = pd.Series(['a', 'b', 'c', 'd'], index=[3, 2, 1, 4])
```

```
| >>> s.sort_index()
```

```
| 1 c
```

```
| 2 b
```

```
| 3 a
```

```
| 4 d
```

```
| dtype: object
```

```
| Sort Descending
```

```
| >>> s.sort_index(ascending=False)
```

```
| 4 d
```

```
| 3 a
```

```
| 2 b
```

```
| 1 c
```

```
| dtype: object
```

```
| Sort Inplace
```

```
| >>> s.sort_index(inplace=True)
```

```
| >>> s
```

```
| 1 c
```

```
| 2 b
```

```
| 3 a
```

```
| 4 d
```

```
| dtype: object
```

```
| By default NaNs are put at the end, but use `na_position` to place
| them at the beginning
```

```
>>> s = pd.Series(['a', 'b', 'c', 'd'], index=[3, 2, 1, np.nan])
>>> s.sort_index(na_position='first')
NaN d
1.0 c
2.0 b
3.0 a
dtype: object
```

Specify index level to sort

```
>>> arrays = [np.array(['qux', 'qux', 'foo', 'foo',
... 'baz', 'baz', 'bar', 'bar']),
... np.array(['two', 'one', 'two', 'one',
... 'two', 'one', 'two', 'one'])]
>>> s = pd.Series([1, 2, 3, 4, 5, 6, 7, 8], index=arrays)
>>> s.sort_index(level=1)
bar one 8
baz one 6
foo one 4
qux one 2
bar two 7
baz two 5
foo two 3
qux two 1
dtype: int64
```

Does not sort by remaining levels when sorting by levels

```
>>> s.sort_index(level=1, sort_remaining=False)
qux one 2
foo one 4
baz one 6
bar one 8
qux two 1
foo two 3
baz two 5
bar two 7
dtype: int64
```

```
sort_values(self, axis=0, ascending=True, inplace=False, kind='quicksort', na_positio
n='last')
```

Sort by the values.

Sort a Series in ascending or descending order by some criterion.

Parameters

-----

axis : {0 or 'index'}, default 0

Axis to direct sorting. The value 'index' is accepted for compatibility with DataFrame.sort\_values.

ascending : bool, default True

If True, sort values in ascending order, otherwise descending.

inplace : bool, default False

If True, perform operation in-place.

kind : {'quicksort', 'mergesort' or 'heapsort'}, default 'quicksort'

Choice of sorting algorithm. See also :func:`numpy.sort` for more information. 'mergesort' is the only stable algorithm.

na\_position : {'first' or 'last'}, default 'last'

Argument 'first' puts NaNs at the beginning, 'last' puts NaNs at the end.

Returns

-----

Series

Series ordered by values.

See Also

-----

Series.sort\_index : Sort by the Series indices.

DataFrame.sort\_values : Sort DataFrame by the values along either axis.

```
DataFrame.sort_index : Sort DataFrame by indices.
```

## Examples

-----

```
>>> s = pd.Series([np.nan, 1, 3, 10, 5])
```

```
>>> s
```

```
0 NaN
```

```
1 1.0
```

```
2 3.0
```

```
3 10.0
```

```
4 5.0
```

```
dtype: float64
```

Sort values ascending order (default behaviour)

```
>>> s.sort_values(ascending=True)
```

```
1 1.0
```

```
2 3.0
```

```
4 5.0
```

```
3 10.0
```

```
0 NaN
```

```
dtype: float64
```

Sort values descending order

```
>>> s.sort_values(ascending=False)
```

```
3 10.0
```

```
4 5.0
```

```
2 3.0
```

```
1 1.0
```

```
0 NaN
```

```
dtype: float64
```

Sort values inplace

```
>>> s.sort_values(ascending=False, inplace=True)
```

```
>>> s
```

```
3 10.0
```

```
4 5.0
```

```
2 3.0
```

```
1 1.0
```

```
0 NaN
```

```
dtype: float64
```

Sort values putting NAs first

```
>>> s.sort_values(na_position='first')
```

```
0 NaN
```

```
1 1.0
```

```
2 3.0
```

```
4 5.0
```

```
3 10.0
```

```
dtype: float64
```

Sort a series of strings

```
>>> s = pd.Series(['z', 'b', 'd', 'a', 'c'])
```

```
>>> s
```

```
0 z
```

```
1 b
```

```
2 d
```

```
3 a
```

```
4 c
```

```
dtype: object
```

```
>>> s.sort_values()
```

```
3 a
```

```
1 b
```

```
4 c
```

```
2 d
```

```
0 z
```

```
dtype: object
```



```
std(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None, **kwargs)
 Return sample standard deviation over requested axis.
```

Normalized by N-1 by default. This can be changed using the ddof argument

#### Parameters

-----

```
axis : {index (0)}
skipna : boolean, default True
 Exclude NA/null values. If an entire row/column is NA, the result
 will be NA
level : int or level name, default None
 If the axis is a MultiIndex (hierarchical), count along a
 particular level, collapsing into a scalar
ddof : int, default 1
 Delta Degrees of Freedom. The divisor used in calculations is N - ddof,
 where N represents the number of elements.
numeric_only : boolean, default None
 Include only float, int, boolean columns. If None, will attempt to use
 everything, then use only numeric data. Not implemented for Series.
```

#### Returns

-----

```
std : scalar or Series (if level specified)
```

```
sub(self, other, level=None, fill_value=None, axis=0)
 Subtraction of series and other, element-wise (binary operator `sub`).
```

Equivalent to ``series - other``, but with support to substitute a fill\_value for missing data in one of the inputs.

#### Parameters

-----

```
other : Series or scalar value
fill_value : None or float value, default None (NaN)
 Fill existing missing (NaN) values, and any new element needed for
 successful Series alignment, with this value before computation.
 If data in both corresponding Series locations is missing
 the result will be missing
level : int or name
 Broadcast across a level, matching Index values on the
 passed MultiIndex level
```

#### Returns

-----

```
result : Series
```

#### See Also

-----

```
Series.rsub
```

#### Examples

-----

```
>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a 1.0
b 1.0
c 1.0
d NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a 1.0
b NaN
d 1.0
e NaN
dtype: float64
>>> a.add(b, fill_value=0)
a 2.0
b 1.0
c 1.0
```

```

| d 1.0
| e NaN
| dtype: float64
|
| subtract = sub(self, other, level=None, fill_value=None, axis=0)
|
| sum(self, axis=None, skipna=None, level=None, numeric_only=None, min_count=0, **kwarg
s)
| Return the sum of the values for the requested axis.
|
| This is equivalent to the method ``numpy.sum``.
|
| Parameters
| -----
| axis : {index (0)}
| Axis for the function to be applied on.
| skipna : bool, default True
| Exclude NA/null values when computing the result.
| level : int or level name, default None
| If the axis is a MultiIndex (hierarchical), count along a
| particular level, collapsing into a scalar.
| numeric_only : bool, default None
| Include only float, int, boolean columns. If None, will attempt to use
| everything, then use only numeric data. Not implemented for Series.
| min_count : int, default 0
| The required number of valid values to perform the operation. If fewer than
| ``min_count`` non-NA values are present the result will be NA.
|
| .. versionadded :: 0.22.0
|
| Added with the default being 0. This means the sum of an all-NA
| or empty Series is 0, and the product of an all-NA or empty
| Series is 1.
| **kwargs
| Additional keyword arguments to be passed to the function.
|
| Returns
| -----
| sum : scalar or Series (if level specified)
|
| See Also
| -----
| Series.sum : Return the sum.
| Series.min : Return the minimum.
| Series.max : Return the maximum.
| Series.idxmin : Return the index of the minimum.
| Series.idxmax : Return the index of the maximum.
| DataFrame.min : Return the sum over the requested axis.
| DataFrame.min : Return the minimum over the requested axis.
| DataFrame.max : Return the maximum over the requested axis.
| DataFrame.idxmin : Return the index of the minimum over the requested axis.
| DataFrame.idxmax : Return the index of the maximum over the requested axis.
|
| Examples
| -----
|
| >>> idx = pd.MultiIndex.from_arrays([
| ... ['warm', 'warm', 'cold', 'cold'],
| ... ['dog', 'falcon', 'fish', 'spider']],
| ... names=['blooded', 'animal'])
| >>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
| >>> s
| blooded animal
| warm dog 4
| falcon 2
| cold fish 0
| spider 8
| Name: legs, dtype: int64
|
| >>> s.sum()
| 14

```

Sum using level names, as well as indices.

```
>>> s.sum(level='blooded')
blooded
warm 6
cold 8
Name: legs, dtype: int64
```

```
>>> s.sum(level=0)
blooded
warm 6
cold 8
Name: legs, dtype: int64
```

By default, the sum of an empty or all-NA Series is ``0``.

```
>>> pd.Series([]).sum() # min_count=0 is the default
0.0
```

This can be controlled with the ``min\_count`` parameter. For example, if you'd like the sum of an empty series to be NaN, pass ``min\_count=1``.

```
>>> pd.Series([]).sum(min_count=1)
nan
```

Thanks to the ``skipna`` parameter, ``min\_count`` handles all-NA and empty series identically.

```
>>> pd.Series([np.nan]).sum()
0.0
```

```
>>> pd.Series([np.nan]).sum(min_count=1)
nan
```

```
swaplevel(self, i=-2, j=-1, copy=True)
Swap levels i and j in a MultiIndex.
```

Parameters  
-----

i, j : int, string (can be mixed)  
Level of index to be swapped. Can pass level name as string.

Returns  
-----

swapped : Series

.. versionchanged:: 0.18.1

The indexes ``i`` and ``j`` are now optional, and default to the two innermost levels of the index.

```
to_csv(self, *args, **kwargs)
Write object to a comma-separated values (csv) file.
```

.. versionchanged:: 0.24.0  
The order of arguments for Series was changed.

Parameters  
-----

path\_or\_buf : str or file handle, default None  
File path or object, if None is provided the result is returned as a string. If a file object is passed it should be opened with ``newline=''``, disabling universal newlines.

.. versionchanged:: 0.24.0

Was previously named "path" for Series.

sep : str, default ','  
String of length 1. Field delimiter for the output file.  
na\_rep : str, default ''  
Missing data representation.

```

float_format : str, default None
 Format string for floating point numbers.
columns : sequence, optional
 Columns to write.
header : bool or list of str, default True
 Write out the column names. If a list of strings is given it is
 assumed to be aliases for the column names.

 .. versionchanged:: 0.24.0

 Previously defaulted to False for Series.

index : bool, default True
 Write row names (index).
index_label : str or sequence, or False, default None
 Column label for index column(s) if desired. If None is given, and
 `header` and `index` are True, then the index names are used. A
 sequence should be given if the object uses MultiIndex. If
 False do not print fields for index names. Use index_label=False
 for easier importing in R.
mode : str
 Python write mode, default 'w'.
encoding : str, optional
 A string representing the encoding to use in the output file,
 defaults to 'ascii' on Python 2 and 'utf-8' on Python 3.
compression : str, default 'infer'
 Compression mode among the following possible values: {'infer',
 'gzip', 'bz2', 'zip', 'xz', None}. If 'infer' and `path_or_buf`
 is path-like, then detect compression from the following
 extensions: '.gz', '.bz2', '.zip' or '.xz'. (otherwise no
 compression).

 .. versionchanged:: 0.24.0

 'infer' option added and set to default.

quoting : optional constant from csv module
 Defaults to csv.QUOTE_MINIMAL. If you have set a `float_format`
 then floats are converted to strings and thus csv.QUOTE_NONNUMERIC
 will treat them as non-numeric.
quotechar : str, default '"'
 String of length 1. Character used to quote fields.
line_terminator : string, optional
 The newline character or character sequence to use in the output
 file. Defaults to `os.linesep`, which depends on the OS in which
 this method is called ('\n' for linux, '\r\n' for Windows, i.e.).

 .. versionchanged:: 0.24.0

chunksize : int or None
 Rows to write at a time.
tupleize_cols : bool, default False
 Write MultiIndex columns as a list of tuples (if True) or in
 the new, expanded format, where each MultiIndex column is a row
 in the CSV (if False).

 .. deprecated:: 0.21.0
 This argument will be removed and will always write each row
 of the multi-index as a separate row in the CSV file.
date_format : str, default None
 Format string for datetime objects.
doublequote : bool, default True
 Control quoting of `quotechar` inside a field.
escapechar : str, default None
 String of length 1. Character used to escape `sep` and `quotechar`
 when appropriate.
decimal : str, default '.'
 Character recognized as decimal separator. E.g. use ',' for
 European data.

```

Returns

-----

None or str

If `path_or_buf` is `None`, returns the resulting csv format as a string. Otherwise returns `None`.

See Also

-----

`read_csv` : Load a CSV file into a `DataFrame`.

`to_excel` : Load an Excel file into a `DataFrame`.

Examples

-----

```
>>> df = pd.DataFrame({'name': ['Raphael', 'Donatello'],
... 'mask': ['red', 'purple'],
... 'weapon': ['sai', 'bo staff']})
>>> df.to_csv(index=False)
'name,mask,weapon\nRaphael,red,sai\nDonatello,purple,bo staff\n'
```

`to_dict(self, into=<class 'dict'>)`

Convert Series to {label -> value} dict or dict-like object.

Parameters

-----

`into` : class, default dict

The `collections.Mapping` subclass to use as the return object. Can be the actual class or an empty instance of the mapping type you want. If you want a `collections.defaultdict`, you must pass it initialized.

.. versionadded:: 0.21.0

Returns

-----

`value_dict` : `collections.Mapping`

Examples

-----

```
>>> s = pd.Series([1, 2, 3, 4])
>>> s.to_dict()
{0: 1, 1: 2, 2: 3, 3: 4}
>>> from collections import OrderedDict, defaultdict
>>> s.to_dict(OrderedDict)
OrderedDict([(0, 1), (1, 2), (2, 3), (3, 4)])
>>> dd = defaultdict(list)
>>> s.to_dict(dd)
defaultdict(<type 'list'>, {0: 1, 1: 2, 2: 3, 3: 4})
```

`to_frame(self, name=None)`

Convert Series to `DataFrame`.

Parameters

-----

`name` : object, default None

The passed name should substitute for the series name (if it has one).

Returns

-----

`data_frame` : `DataFrame`

`to_period(self, freq=None, copy=True)`

Convert Series from `DatetimeIndex` to `PeriodIndex` with desired frequency (inferred from index if not passed).

Parameters

-----

`freq` : string, default

Returns

-----

`ts` : Series with `PeriodIndex`

`to_sparse(self, kind='block', fill_value=None)`

Convert Series to `SparseSeries`.

Parameters

-----

kind : {'block', 'integer'}  
fill\_value : float, defaults to NaN (missing)

Returns

-----

sp : SparseSeries

to\_string(self, buf=None, na\_rep='NaN', float\_format=None, header=True, index=True, length=False, dtype=False, name=False, max\_rows=None)

Render a string representation of the Series.

Parameters

-----

buf : StringIO-like, optional  
buffer to write to  
na\_rep : string, optional  
string representation of NAN to use, default 'NaN'  
float\_format : one-parameter function, optional  
formatter function to apply to columns' elements if they are floats  
default None  
header : boolean, default True  
Add the Series header (index name)  
index : bool, optional  
Add index (row) labels, default True  
length : boolean, default False  
Add the Series length  
dtype : boolean, default False  
Add the Series dtype  
name : boolean, default False  
Add the Series name if not None  
max\_rows : int, optional  
Maximum number of rows to show before truncating. If None, show all.

Returns

-----

formatted : string (if not buffer passed)

to\_timestamp(self, freq=None, how='start', copy=True)

Cast to DatetimeIndex of timestamps, at \*beginning\* of period.

Parameters

-----

freq : string, default frequency of PeriodIndex  
Desired frequency  
how : {'s', 'e', 'start', 'end'}  
Convention for converting period to timestamp; start of period  
vs. end

Returns

-----

ts : Series with DatetimeIndex

transform(self, func, axis=0, \*args, \*\*kwargs)

Call ``func`` on self producing a Series with transformed values  
and that has the same axis length as self.

.. versionadded:: 0.20.0

Parameters

-----

func : function, str, list or dict  
Function to use for transforming the data. If a function, must either  
work when passed a Series or when passed to Series.apply.

Accepted combinations are:

- function
- string function name

- list of functions and/or function names, e.g. ``[np.exp, 'sqrt']``
- dict of axis labels -> functions, function names or list of such.

axis : {0 or 'index'}

Parameter needed for compatibility with DataFrame.

\*args

Positional arguments to pass to `func`.

\*\*kwargs

Keyword arguments to pass to `func`.

Returns

-----

Series

A Series that must have the same length as self.

Raises

-----

ValueError : If the returned Series has a different length than self.

See Also

-----

Series.agg : Only perform aggregating type operations.

Series.apply : Invoke function on a Series.

Examples

-----

```
>>> df = pd.DataFrame({'A': range(3), 'B': range(1, 4)})
```

```
>>> df
```

```
 A B
0 0 1
1 1 2
2 2 3
```

```
>>> df.transform(lambda x: x + 1)
```

```
 A B
0 1 2
1 2 3
2 3 4
```

Even though the resulting Series must have the same length as the input Series, it is possible to provide several input functions:

```
>>> s = pd.Series(range(3))
```

```
>>> s
```

```
0 0
1 1
2 2
```

```
dtype: int64
```

```
>>> s.transform([np.sqrt, np.exp])
```

```
 sqrt exp
0 0.000000 1.000000
1 1.000000 2.718282
2 1.414214 7.389056
```

```
truediv(self, other, level=None, fill_value=None, axis=0)
```

Floating division of series and other, element-wise (binary operator `truediv`).

Equivalent to ``series / other``, but with support to substitute a fill\_value for missing data in one of the inputs.

Parameters

-----

other : Series or scalar value

fill\_value : None or float value, default None (NaN)

Fill existing missing (NaN) values, and any new element needed for successful Series alignment, with this value before computation.

If data in both corresponding Series locations is missing the result will be missing

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level

Returns

-----

```
result : Series
```

See Also

Series.rtruediv

Examples

```

>>> a = pd.Series([1, 1, 1, np.nan], index=['a', 'b', 'c', 'd'])
>>> a
a 1.0
b 1.0
c 1.0
d NaN
dtype: float64
>>> b = pd.Series([1, np.nan, 1, np.nan], index=['a', 'b', 'd', 'e'])
>>> b
a 1.0
b NaN
d 1.0
e NaN
dtype: float64
>>> a.add(b, fill_value=0)
a 2.0
b 1.0
c 1.0
d 1.0
e NaN
dtype: float64
```

unique(self)

Return unique values of Series object.

Uniques are returned in order of appearance. Hash table-based unique, therefore does NOT sort.

Returns

-----  
ndarray or ExtensionArray

The unique values returned as a NumPy array. In case of an extension-array backed Series, a new :class:`~api.extensions.ExtensionArray` of that type with just the unique values is returned. This includes

- \* Categorical
- \* Period
- \* Datetime with Timezone
- \* Interval
- \* Sparse
- \* IntegerNA

See Also

-----  
unique : Top-level unique method for any 1-d array-like object.  
Index.unique : Return Index with unique values from an Index object.

Examples

```

>>> pd.Series([2, 1, 3, 3], name='A').unique()
array([2, 1, 3])

>>> pd.Series([pd.Timestamp('2016-01-01') for _ in range(3)]).unique()
array(['2016-01-01T00:00:00.000000000', dtype='datetime64[ns]'])

>>> pd.Series([pd.Timestamp('2016-01-01', tz='US/Eastern')
... for _ in range(3)]).unique()
<DatetimeArray>
['2016-01-01 00:00:00-05:00']
Length: 1, dtype: datetime64[ns, US/Eastern]
```

An unordered Categorical will return categories in the order of appearance.



```
>>> pd.Series(pd.Categorical(list('baabc'))).unique()
[b, a, c]
Categories (3, object): [b, a, c]
```

An ordered Categorical preserves the category ordering.

```
>>> pd.Series(pd.Categorical(list('baabc'), categories=list('abc'),
... ordered=True)).unique()
[b, a, c]
Categories (3, object): [a < b < c]
```

```
unstack(self, level=-1, fill_value=None)
Unstack, a.k.a. pivot, Series with MultiIndex to produce DataFrame.
The level involved will automatically get sorted.
```

Parameters

-----  
level : int, string, or list of these, default last level  
Level(s) to unstack, can pass level name  
fill\_value : replace NaN with this value if the unstack produces  
missing values

.. versionadded:: 0.18.0

Returns

-----  
unstacked : DataFrame

Examples

```
>>> s = pd.Series([1, 2, 3, 4],
... index=pd.MultiIndex.from_product(['one', 'two'], ['a', 'b']))
>>> s
one a 1
 b 2
two a 3
 b 4
dtype: int64
```

```
>>> s.unstack(level=-1)
 a b
one 1 2
two 3 4
```

```
>>> s.unstack(level=0)
 one two
a 1 3
b 2 4
```

```
update(self, other)
Modify Series in place using non-NA values from passed
Series. Aligns on index.
```

Parameters

-----  
other : Series

Examples

```
>>> s = pd.Series([1, 2, 3])
>>> s.update(pd.Series([4, 5, 6]))
>>> s
0 4
1 5
2 6
dtype: int64
```

```
>>> s = pd.Series(['a', 'b', 'c'])
>>> s.update(pd.Series(['d', 'e'], index=[0, 2]))
>>> s
0 d
```

```
1 b
2 e
dtype: object
```

```
>>> s = pd.Series([1, 2, 3])
>>> s.update(pd.Series([4, 5, 6, 7, 8]))
>>> s
0 4
1 5
2 6
dtype: int64
```

If ``other`` contains NaNs the corresponding values are not updated in the original Series.

```
>>> s = pd.Series([1, 2, 3])
>>> s.update(pd.Series([4, np.nan, 6]))
>>> s
0 4
1 2
2 6
dtype: int64
```

```
valid(self, inplace=False, **kwargs)
 Return Series without null values.
```

```
.. deprecated:: 0.23.0
 Use :meth:`Series.dropna` instead.
```

```
var(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None, **kwargs)
 Return unbiased variance over requested axis.
```

Normalized by N-1 by default. This can be changed using the ddof argument

Parameters

-----

axis : {index (0)}

skipna : boolean, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a scalar

ddof : int, default 1

Delta Degrees of Freedom. The divisor used in calculations is N - ddof, where N represents the number of elements.

numeric\_only : boolean, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

Returns

-----

var : scalar or Series (if level specified)

```
view(self, dtype=None)
 Create a new view of the Series.
```

This function will return a new Series with a view of the same underlying values in memory, optionally reinterpreted with a new data type. The new data type must preserve the same size in bytes as to not cause index misalignment.

Parameters

-----

dtype : data type

Data type object or one of their string representations.

Returns

-----

Series

A new Series object as a view of the same data in memory.

See Also

`numpy.ndarray.view` : Equivalent numpy function to create a new view of the same data in memory.

Notes

Series are instantiated with `dtype=float64` by default. While `numpy.ndarray.view()` will return a view with the same data type as the original array, `Series.view()` (without specified dtype) will try using `float64` and may fail if the original data type size in bytes is not the same.

Examples

```
>>> s = pd.Series([-2, -1, 0, 1, 2], dtype='int8')
>>> s
0 -2
1 -1
2 0
3 1
4 2
dtype: int8
```

The 8 bit signed integer representation of `-1` is `0b11111111`, but the same bytes represent 255 if read as an 8 bit unsigned integer:

```
>>> us = s.view('uint8')
>>> us
0 254
1 255
2 0
3 1
4 2
dtype: uint8
```

The views share the same underlying values:

```
>>> us[0] = 128
>>> s
0 -128
1 -1
2 0
3 1
4 2
dtype: int8
```

---

Class methods defined here:

`from_array(arr, index=None, name=None, dtype=None, copy=False, fastpath=False)` from `builtins.type`

Construct Series from array.

.. deprecated :: 0.23.0

Use `pd.Series(...)` constructor instead.

`from_csv(path, sep=',', parse_dates=True, header=None, index_col=0, encoding=None, infer_datetime_format=False)` from `builtins.type`

Read CSV file.

.. deprecated:: 0.21.0

Use `:func:`pandas.read_csv`` instead.

It is preferable to use the more powerful `:func:`pandas.read_csv`` for most general purposes, but `from_csv` makes for an easy roundtrip to and from a file (the exact counterpart of `to_csv`), especially with a time Series.

This method only differs from `:func:`pandas.read_csv`` in some defaults:

- `index_col` is `0` instead of `None` (take first column as index

```
by default)
- `header` is ``None`` instead of ``0`` (the first row is not used as
 the column names)
- `parse_dates` is ``True`` instead of ``False`` (try parsing the index
 as datetime by default)
```

With `:func:`pandas.read_csv``, the option ```squeeze=True``` can be used to return a Series like ```from_csv```.

#### Parameters

-----

```
path : string file path or file handle / StringIO
sep : string, default ',',
 Field delimiter
parse_dates : boolean, default True
 Parse dates. Different default from read_table
header : int, default None
 Row to use as header (skip prior rows)
index_col : int or sequence, default 0
 Column to use for index. If a sequence is given, a MultiIndex
 is used. Different default from read_table
encoding : string, optional
 a string representing the encoding to use if the contents are
 non-ascii, for python versions prior to 3
infer_datetime_format : boolean, default False
 If True and `parse_dates` is True for a column, try to infer the
 datetime format based on the first datetime string. If the format
 can be inferred, there often will be a large parsing speed-up.
```

#### Returns

-----

```
y : Series
```

#### See Also

-----

```
read_csv
```

-----  
Data descriptors defined here:

#### asobject

Return object Series which contains boxed values.

.. deprecated :: 0.23.0

Use ```astype(object)``` instead.

*\*this is an internal non-public method\**

#### axes

Return a list of the row axis labels.

#### dtype

Return the dtype object of the underlying data.

#### dtypes

Return the dtype object of the underlying data.

#### ftype

Return if the data is sparse|dense.

#### ftypes

Return if the data is sparse|dense.

#### hasnans

Return if I have any nans; enables various perf speedups.

#### imag

Return imag value of vector.

#### index

The index (axis labels) of the Series.

name  
Return name of the Series.

real  
Return the real value of vector.

values  
Return Series as ndarray or ndarray-like depending on the dtype.

.. warning::

We recommend using :attr:`Series.array` or  
:meth:`Series.to\_numpy`, depending on whether you need  
a reference to the underlying data or a NumPy array.

#### Returns

-----

arr : numpy.ndarray or ndarray-like

#### See Also

-----

Series.array : Reference to the underlying data.

Series.to\_numpy : A NumPy array representing the underlying data.

#### Examples

-----

```
>>> pd.Series([1, 2, 3]).values
array([1, 2, 3])
```

```
>>> pd.Series(list('aabc')).values
array(['a', 'a', 'b', 'c'], dtype=object)
```

```
>>> pd.Series(list('aabc')).astype('category').values
[a, a, b, c]
Categories (3, object): [a, b, c]
```

Timezone aware datetime data is converted to UTC:

```
>>> pd.Series(pd.date_range('20130101', periods=3,
... tz='US/Eastern')).values
array(['2013-01-01T05:00:00.000000000',
 '2013-01-02T05:00:00.000000000',
 '2013-01-03T05:00:00.000000000'], dtype='datetime64[ns]')
```

-----  
Data and other attributes defined here:

cat = <class 'pandas.core.arrays.categorical.CategoricalAccessor'>  
Accessor object for categorical properties of the Series values.

Be aware that assigning to `categories` is a inplace operation, while all  
methods return new categorical data per default (but can be called with  
`inplace=True`).

#### Parameters

-----

data : Series or CategoricalIndex

#### Examples

-----

```
>>> s.cat.categories
>>> s.cat.categories = list('abc')
>>> s.cat.rename_categories(list('cab'))
>>> s.cat.reorder_categories(list('cab'))
>>> s.cat.add_categories(['d', 'e'])
>>> s.cat.remove_categories(['d'])
>>> s.cat.remove_unused_categories()
>>> s.cat.set_categories(list('abcde'))
>>> s.cat.as_ordered()
>>> s.cat.as_unordered()
```

```
dt = <class 'pandas.core.indexes.accessors.CombinedDatetimelikePropert...
 Accessor object for datetimelike properties of the Series values.
```

#### Examples

-----

```
>>> s.dt.hour
>>> s.dt.second
>>> s.dt.quarter
```

Returns a Series indexed like the original Series.

Raises TypeError if the Series does not contain datetimelike values.

```
plot = <class 'pandas.plotting._core.SeriesPlotMethods'>
```

Series plotting accessor and method.

#### Examples

-----

```
>>> s.plot.line()
>>> s.plot.bar()
>>> s.plot.hist()
```

Plotting methods can also be accessed by calling the accessor as a method with the ``kind`` argument:

``s.plot(kind='line')`` is equivalent to ``s.plot.line()``

```
sparse = <class 'pandas.core.arrays.sparse.SparseAccessor'>
```

Accessor for SparseSparse from other sparse matrix data types.

```
str = <class 'pandas.core.strings.StringMethods'>
```

Vectorized string functions for Series and Index. NAs stay NA unless handled otherwise by a particular method. Patterned after Python's string methods, with some inspiration from R's stringr package.

#### Examples

-----

```
>>> s.str.split('_')
>>> s.str.replace('_', '')
```

-----  
Methods inherited from pandas.core.base.IndexOpsMixin:

```
__iter__(self)
```

Return an iterator of the values.

These are each a scalar type, which is a Python scalar  
(for str, int, float) or a pandas scalar  
(for Timestamp/Timedelta/Interval/Period)

```
factorize(self, sort=False, na_sentinel=-1)
```

Encode the object as an enumerated type or categorical variable.

This method is useful for obtaining a numeric representation of an array when all that matters is identifying distinct values. ``factorize`` is available as both a top-level function :func:`pandas.factorize`, and as a method :meth:`Series.factorize` and :meth:`Index.factorize`.

#### Parameters

-----

sort : boolean, default False

Sort ``uniques`` and shuffle ``labels`` to maintain the relationship.

na\_sentinel : int, default -1

Value to mark "not found".

#### Returns

-----

labels : ndarray

An integer ndarray that's an indexer into ``uniques``.

``uniques.take(labels)`` will have the same values as ``values``.

uniques : ndarray, Index, or Categorical

The unique valid values. When ``values`` is Categorical, ``uniques``

is a Categorical. When `values` is some other pandas object, an `Index` is returned. Otherwise, a 1-D ndarray is returned.

.. note ::

Even if there's a missing value in `values`, `uniques` will *not* contain an entry for it.

See Also

-----

cut : Discretize continuous-valued array.  
unique : Find the unique value in an array.

Examples

-----

These examples all show factorize as a top-level method like ``pd.factorize(values)``. The results are identical for methods like :meth:`Series.factorize`.

```
>>> labels, uniques = pd.factorize(['b', 'b', 'a', 'c', 'b'])
>>> labels
array([0, 0, 1, 2, 0])
>>> uniques
array(['b', 'a', 'c'], dtype=object)
```

With ``sort=True``, the `uniques` will be sorted, and `labels` will be shuffled so that the relationship is the maintained.

```
>>> labels, uniques = pd.factorize(['b', 'b', 'a', 'c', 'b'], sort=True)
>>> labels
array([1, 1, 0, 2, 1])
>>> uniques
array(['a', 'b', 'c'], dtype=object)
```

Missing values are indicated in `labels` with `na\_sentinel` (`-1` by default). Note that missing values are never included in `uniques`.

```
>>> labels, uniques = pd.factorize(['b', None, 'a', 'c', 'b'])
>>> labels
array([0, -1, 1, 2, 0])
>>> uniques
array(['b', 'a', 'c'], dtype=object)
```

Thus far, we've only factorized lists (which are internally coerced to NumPy arrays). When factorizing pandas objects, the type of `uniques` will differ. For Categoricals, a `Categorical` is returned.

```
>>> cat = pd.Categorical(['a', 'a', 'c'], categories=['a', 'b', 'c'])
>>> labels, uniques = pd.factorize(cat)
>>> labels
array([0, 0, 1])
>>> uniques
[a, c]
Categories (3, object): [a, b, c]
```

Notice that ``'b'`` is in ``uniques.categories``, despite not being present in ``cat.values``.

For all other pandas objects, an Index of the appropriate type is returned.

```
>>> cat = pd.Series(['a', 'a', 'c'])
>>> labels, uniques = pd.factorize(cat)
>>> labels
array([0, 0, 1])
>>> uniques
Index(['a', 'c'], dtype='object')
```

item(self)

Return the first element of the underlying data as a python scalar.

```
nunique(self, dropna=True)
 Return number of unique elements in the object.
```

Excludes NA values by default.

Parameters

-----

dropna : boolean, default True  
 Don't include NaN in the count.

Returns

-----

nunique : int

```
to_list = tolist(self)
 Return a list of the values.
```

These are each a scalar type, which is a Python scalar  
(for str, int, float) or a pandas scalar  
(for Timestamp/Timedelta/Interval/Period)

See Also

-----

numpy.ndarray.tolist

```
to_numpy(self, dtype=None, copy=False)
 A NumPy ndarray representing the values in this Series or Index.
```

.. versionadded:: 0.24.0

Parameters

-----

dtype : str or numpy.dtype, optional  
 The dtype to pass to :meth:`numpy.asarray`  
copy : bool, default False  
 Whether to ensure that the returned value is a not a view on  
 another array. Note that ``copy=False`` does not *ensure* that  
 ``to\_numpy()`` is no-copy. Rather, ``copy=True`` ensure that  
 a copy is made, even if not strictly necessary.

Returns

-----

numpy.ndarray

See Also

-----

Series.array : Get the actual data stored within.  
Index.array : Get the actual data stored within.  
DataFrame.to\_numpy : Similar method for DataFrame.

Notes

-----

The returned array will be the same up to equality (values equal  
in `self` will be equal in the returned array; likewise for values  
that are not equal). When `self` contains an ExtensionArray, the  
dtype may be different. For example, for a category-dtype Series,  
``to\_numpy()`` will return a NumPy array and the categorical dtype  
will be lost.

For NumPy dtypes, this will be a reference to the actual data stored  
in this Series or Index (assuming ``copy=False``). Modifying the result  
in place will modify the data stored in the Series or Index (not that  
we recommend doing that).

For extension types, ``to\_numpy()`` *may* require copying data and  
coercing the result to a NumPy type (possibly object), which may be  
expensive. When you need a no-copy reference to the underlying data,  
:attr:`Series.array` should be used instead.

This table lays out the different dtypes and default return types of  
``to\_numpy()`` for various dtypes within pandas.



```

=====
dtype array type
=====
category[T] ndarray[T] (same dtype as input)
period ndarray[object] (Periods)
interval ndarray[object] (Intervals)
IntegerNA ndarray[object]
datetime64[ns] datetime64[ns]
datetime64[ns, tz] ndarray[object] (Timestamps)
=====

```

#### Examples

```

>>> ser = pd.Series(pd.Categorical(['a', 'b', 'a']))
>>> ser.to_numpy()
array(['a', 'b', 'a'], dtype=object)

```

Specify the `dtype` to control how datetime-aware data is represented. Use ``dtype=object`` to return an ndarray of pandas :class:`Timestamp` objects, each with the correct ``tz``.

```

>>> ser = pd.Series(pd.date_range('2000', periods=2, tz="CET"))
>>> ser.to_numpy(dtype=object)
array([Timestamp('2000-01-01 00:00:00+0100', tz='CET', freq='D'),
 Timestamp('2000-01-02 00:00:00+0100', tz='CET', freq='D')],
 dtype=object)

```

Or ``dtype='datetime64[ns]'`` to return an ndarray of native datetime64 values. The values are converted to UTC and the timezone info is dropped.

```

>>> ser.to_numpy(dtype="datetime64[ns]")
... # doctest: +ELLIPSIS
array(['1999-12-31T23:00:00.000000000', '2000-01-01T23:00:00...'],
 dtype='datetime64[ns]')

```

`tolist(self)`

Return a list of the values.

These are each a scalar type, which is a Python scalar (for str, int, float) or a pandas scalar (for Timestamp/Timedelta/Interval/Period)

See Also

-----

`numpy.ndarray.tolist`

`transpose(self, *args, **kwargs)`

Return the transpose, which is by definition self.

`value_counts(self, normalize=False, sort=True, ascending=False, bins=None, dropna=True)`

Return a Series containing counts of unique values.

The resulting object will be in descending order so that the first element is the most frequently-occurring element. Excludes NA values by default.

Parameters

-----

`normalize` : boolean, default False

If True then the object returned will contain the relative frequencies of the unique values.

`sort` : boolean, default True

Sort by values.

`ascending` : boolean, default False

Sort in ascending order.

`bins` : integer, optional

Rather than count values, group them into half-open bins, a convenience for ``pd.cut``, only works with numeric data.

`dropna` : boolean, default True

Don't include counts of NaN.

Returns

-----

counts : Series

See Also

-----

Series.count: Number of non-NA elements in a Series.

DataFrame.count: Number of non-NA elements in a DataFrame.

Examples

-----

```
>>> index = pd.Index([3, 1, 2, 3, 4, np.nan])
```

```
>>> index.value_counts()
```

```
3.0 2
```

```
4.0 1
```

```
2.0 1
```

```
1.0 1
```

```
dtype: int64
```

With `normalize` set to `True`, returns the relative frequency by dividing all values by the sum of values.

```
>>> s = pd.Series([3, 1, 2, 3, 4, np.nan])
```

```
>>> s.value_counts(normalize=True)
```

```
3.0 0.4
```

```
4.0 0.2
```

```
2.0 0.2
```

```
1.0 0.2
```

```
dtype: float64
```

**\*\*bins\*\***

Bins can be useful for going from a continuous variable to a categorical variable; instead of counting unique apparitions of values, divide the index in the specified number of half-open bins.

```
>>> s.value_counts(bins=3)
```

```
(2.0, 3.0] 2
```

```
(0.996, 2.0] 2
```

```
(3.0, 4.0] 1
```

```
dtype: int64
```

**\*\*dropna\*\***

With `dropna` set to `False` we can also see NaN index values.

```
>>> s.value_counts(dropna=False)
```

```
3.0 2
```

```
NaN 1
```

```
4.0 1
```

```
2.0 1
```

```
1.0 1
```

```
dtype: int64
```

---

Data descriptors inherited from pandas.core.base.IndexOpsMixin:

T

Return the transpose, which is by definition self.

\_\_dict\_\_

dictionary for instance variables (if defined)

\_\_weakref\_\_

list of weak references to the object (if defined)

array

TheExtensionArray of the data backing this Series or Index.

```
.. versionadded:: 0.24.0
```

## Returns

-----

array : ExtensionArray

An ExtensionArray of the values stored within. For extension types, this is the actual array. For NumPy native types, this is a thin (no copy) wrapper around :class:`numpy.ndarray`.

```.array``` differs ```.values``` which may require converting the data to a different form.

See Also

`Index.to_numpy` : Similar method that always returns a NumPy array.

`Series.to_numpy` : Similar method that always returns a NumPy array.

Notes

This table lays out the different array types for each extension dtype within pandas.

dtype	array type
category	Categorical
period	PeriodArray
interval	IntervalArray
IntegerNA	IntegerArray
datetime64[ns, tz]	DatetimeArray

For any 3rd-party extension types, the array type will be an ExtensionArray.

For all remaining dtypes ```.array``` will be a :class:`arrays.NumpyExtensionArray` wrapping the actual ndarray stored within. If you absolutely need a NumPy array (possibly with copying / coercing data), then use :meth:`Series.to_numpy` instead.

Examples

For regular NumPy types like int, and float, a PandasArray is returned.

```
>>> pd.Series([1, 2, 3]).array
<PandasArray>
[1, 2, 3]
Length: 3, dtype: int64
```

For extension types, like Categorical, the actual ExtensionArray is returned

```
>>> ser = pd.Series(pd.Categorical(['a', 'b', 'a']))
>>> ser.array
[a, b, a]
Categories (2, object): [a, b]
```

base

Return the base object if the memory of the underlying data is shared.

data

Return the data pointer of the underlying data.

empty

flags

Return the ndarray.flags for the underlying data.

is_monotonic

Return boolean if values in the object are

```

monotonic_increasing.

.. versionadded:: 0.19.0

Returns
-----
is_monotonic : boolean

is_monotonic_decreasing
    Return boolean if values in the object are
    monotonic_decreasing.

.. versionadded:: 0.19.0

Returns
-----
is_monotonic_decreasing : boolean

is_monotonic_increasing
    Return boolean if values in the object are
    monotonic_increasing.

.. versionadded:: 0.19.0

Returns
-----
is_monotonic : boolean

is_unique
    Return boolean if values in the object are unique.

Returns
-----
is_unique : boolean

itemsizesize
    Return the size of the dtype of the item of the underlying data.

nbytes
    Return the number of bytes in the underlying data.

ndim
    Number of dimensions of the underlying data, by definition 1.

shape
    Return a tuple of the shape of the underlying data.

size
    Return the number of elements in the underlying data.

strides
    Return the strides of the underlying data.

-----
Data and other attributes inherited from pandas.core.base.IndexOpsMixin:
__array_priority__ = 1000
-----

Methods inherited from pandas.core.generic.NDFrame:

__abs__(self)

__bool__ = __nonzero__(self)

__contains__(self, key)
    True if the key is in the info axis

__copy__(self, deep=True)

__deepcopy__(self, memo=None)
    Parameters

```

```
-----
memo, default None
Standard signature. Unused
```

```
__delitem__(self, key)
Delete item
```

```
__finalize__(self, other, method=None, **kwargs)
Propagate metadata from other to self.
```

```
Parameters
-----
```

```
other : the object from which to get the attributes that we are going
to propagate
method : optional, a passed method name ; possibly to take different
types of propagation actions based on this
```

```
__getattr__(self, name)
After regular attribute access, try looking up the name
This allows simpler access to columns for interactive use.
```

```
__getstate__(self)
```

```
__hash__(self)
Return hash(self).
```

```
__invert__(self)
```

```
__neg__(self)
```

```
__nonzero__(self)
```

```
__pos__(self)
```

```
__round__(self, decimals=0)
```

```
__setattr__(self, name, value)
After regular attribute access, try setting the name
This allows simpler access to columns for interactive use.
```

```
__setstate__(self, state)
```

```
abs(self)
Return a Series/DataFrame with absolute numeric value of each element.
```

This function only applies to elements that are all numeric.

```
Returns
-----
```

```
abs
Series/DataFrame containing the absolute value of each element.
```

```
See Also
-----
```

```
numpy.absolute : Calculate the absolute value element-wise.
```

```
Notes
-----
```

```
For ``complex`` inputs, ``1.2 + 1j``, the absolute value is
:math:\sqrt{a^2 + b^2}.
```

```
Examples
-----
```

```
Absolute numeric values in a Series.
```

```
>>> s = pd.Series([-1.10, 2, -3.33, 4])
>>> s.abs()
0    1.10
1    2.00
2    3.33
3    4.00
dtype: float64
```

Absolute numeric values in a Series with complex numbers.

```
>>> s = pd.Series([1.2 + 1j])
>>> s.abs()
0    1.56205
dtype: float64
```

Absolute numeric values in a Series with a Timedelta element.

```
>>> s = pd.Series([pd.Timedelta('1 days')])
>>> s.abs()
0    1 days
dtype: timedelta64[ns]
```

Select rows with data closest to certain value using argsort (from [StackOverflow <https://stackoverflow.com/a/17758115>](https://stackoverflow.com/a/17758115) `__`).

```
>>> df = pd.DataFrame({
...     'a': [4, 5, 6, 7],
...     'b': [10, 20, 30, 40],
...     'c': [100, 50, -30, -50]
... })
>>> df
   a  b  c
0  4 10 100
1  5 20  50
2  6 30 -30
3  7 40 -50
>>> df.loc[(df.c - 43).abs().argsort()]
   a  b  c
1  5 20  50
0  4 10 100
2  6 30 -30
3  7 40 -50
```

```
add_prefix(self, prefix)
    Prefix labels with string `prefix`.
```

For Series, the row labels are prefixed.
For DataFrame, the column labels are prefixed.

Parameters

prefix : str
 The string to add before each label.

Returns

Series or DataFrame

New Series or DataFrame with updated labels.

See Also

Series.add_suffix: Suffix row labels with string `suffix`.
DataFrame.add_suffix: Suffix column labels with string `suffix`.

Examples

```
>>> s = pd.Series([1, 2, 3, 4])
>>> s
0    1
1    2
2    3
3    4
dtype: int64
```

```
>>> s.add_prefix('item_')
item_0    1
item_1    2
item_2    3
item_3    4
```

```
dtype: int64
```

```
>>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
```

```
>>> df
```

```
   A  B
0  1  3
1  2  4
2  3  5
3  4  6
```

```
>>> df.add_prefix('col_')
```

```
   col_A  col_B
0      1      3
1      2      4
2      3      5
3      4      6
```

```
add_suffix(self, suffix)
```

```
    Suffix labels with string `suffix`.
```

```
For Series, the row labels are suffixed.
```

```
For DataFrame, the column labels are suffixed.
```

```
Parameters
```

```
-----
```

```
suffix : str
```

```
    The string to add after each label.
```

```
Returns
```

```
-----
```

```
Series or DataFrame
```

```
    New Series or DataFrame with updated labels.
```

```
See Also
```

```
-----
```

```
Series.add_prefix: Prefix row labels with string `prefix`.
```

```
DataFrame.add_prefix: Prefix column labels with string `prefix`.
```

```
Examples
```

```
-----
```

```
>>> s = pd.Series([1, 2, 3, 4])
```

```
>>> s
```

```
0    1
1    2
2    3
3    4
```

```
dtype: int64
```

```
>>> s.add_suffix('_item')
```

```
0_item    1
1_item    2
2_item    3
3_item    4
dtype: int64
```

```
>>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
```

```
>>> df
```

```
   A  B
0  1  3
1  2  4
2  3  5
3  4  6
```

```
>>> df.add_suffix('_col')
```

```
   A_col  B_col
0      1      3
1      2      4
2      3      5
3      4      6
```

```
as_blocks(self, copy=True)
```

```
    Convert the frame to a dict of dtype -> Constructor Types that each has
```

a homogeneous dtype.

.. deprecated:: 0.21.0

NOTE: the dtypes of the blocks WILL BE PRESERVED HERE (unlike in `as_matrix`)

Parameters

`copy` : boolean, default True

Returns

values : a dict of dtype -> Constructor Types

`as_matrix(self, columns=None)`

Convert the frame to its Numpy-array representation.

.. deprecated:: 0.23.0

Use `:meth:`DataFrame.values`` instead.

Parameters

`columns` : list, optional, default:None

If None, return all columns, otherwise, returns specified columns.

Returns

values : ndarray

If the caller is heterogeneous and contains booleans or objects, the result will be of dtype=object. See Notes.

See Also

`DataFrame.values`

Notes

Return is NOT a Numpy-matrix, rather, a Numpy-array.

The dtype will be a lower-common-denominator dtype (implicit upcasting); that is to say if the dtypes (even of numeric types) are mixed, the one that accommodates all will be chosen. Use this with care if you are not dealing with the blocks.

e.g. If the dtypes are float16 and float32, dtype will be upcast to float32. If dtypes are int32 and uint8, dtype will be upcase to int32. By `numpy.find_common_type` convention, mixing int64 and uint64 will result in a float64 dtype.

This method is provided for backwards compatibility. Generally, it is recommended to use `'.values'`.

`asfreq(self, freq, method=None, how=None, normalize=False, fill_value=None)`

Convert TimeSeries to specified frequency.

Optionally provide filling method to pad/backfill missing values.

Returns the original data conformed to a new index with the specified frequency. ```resample``` is more appropriate if an operation, such as summarization, is necessary to represent the data at the new frequency.

Parameters

`freq` : DateOffset object, or string

`method` : {'backfill'/'bfill', 'pad'/'ffill'}, default None

Method to use for filling holes in reindexed Series (note this does not fill NaNs that already were present):

- * 'pad' / 'ffill': propagate last valid observation forward to next valid

- * 'backfill' / 'bfill': use NEXT valid observation to fill


```
how : {'start', 'end'}, default end
      For PeriodIndex only, see PeriodIndex.asfreq
normalize : bool, default False
          Whether to reset output index to midnight
fill_value : scalar, optional
            Value to use for missing values, applied during upsampling (note
            this does not fill NaNs that already were present).
```

```
.. versionadded:: 0.20.0
```

Returns

```
-----
converted : same type as caller
```

See Also

```
-----
reindex
```

Notes

```
-----
To learn more about the frequency strings, please see `this link
<http://pandas.pydata.org/pandas-docs/stable/timeseries.html#offset-aliases>`__.
```

Examples

```
-----
```

Start by creating a series with 4 one minute timestamps.

```
>>> index = pd.date_range('1/1/2000', periods=4, freq='T')
>>> series = pd.Series([0.0, None, 2.0, 3.0], index=index)
>>> df = pd.DataFrame({'s':series})
>>> df
```

	s
2000-01-01 00:00:00	0.0
2000-01-01 00:01:00	NaN
2000-01-01 00:02:00	2.0
2000-01-01 00:03:00	3.0

Upsample the series into 30 second bins.

```
>>> df.asfreq(freq='30S')
s
2000-01-01 00:00:00    0.0
2000-01-01 00:00:30    NaN
2000-01-01 00:01:00    NaN
2000-01-01 00:01:30    NaN
2000-01-01 00:02:00    2.0
2000-01-01 00:02:30    NaN
2000-01-01 00:03:00    3.0
```

Upsample again, providing a ``fill value``.

```
>>> df.asfreq(freq='30S', fill_value=9.0)
s
2000-01-01 00:00:00    0.0
2000-01-01 00:00:30    9.0
2000-01-01 00:01:00    NaN
2000-01-01 00:01:30    9.0
2000-01-01 00:02:00    2.0
2000-01-01 00:02:30    9.0
2000-01-01 00:03:00    3.0
```

Upsample again, providing a ``method``.

```
>>> df.asfreq(freq='30S', method='bfill')
s
2000-01-01 00:00:00    0.0
2000-01-01 00:00:30    NaN
2000-01-01 00:01:00    NaN
2000-01-01 00:01:30    2.0
2000-01-01 00:02:00    2.0
2000-01-01 00:02:30    3.0
```

```
asof(self, where, subset=None)
```

Return the last row(s) without any NaNs before `where`.

The last row (for each element in `where`, if list) without any NaN is taken.

In case of a :class:`~pandas.DataFrame`, the last row without NaN considering only the subset of columns (if not `None`)

.. versionadded:: 0.19.0 For DataFrame

If there is no good value, NaN is returned for a Series or a Series of NaN values for a DataFrame

Parameters

where : date or array-like of dates

Date(s) before which the last row(s) are returned.

subset : str or array-like of str, default `None`

For DataFrame, if not `None`, only use these columns to check for NaNs.

Returns

scalar, Series, or DataFrame

- * scalar : when `self` is a Series and `where` is a scalar
- * Series: when `self` is a Series and `where` is an array-like, or when `self` is a DataFrame and `where` is a scalar
- * DataFrame : when `self` is a DataFrame and `where` is an array-like

See Also

merge_asof : Perform an asof merge. Similar to left join.

Notes

Dates are assumed to be sorted. Raises if this is not the case.

Examples

A Series and a scalar `where`.

```
>>> s = pd.Series([1, 2, np.nan, 4], index=[10, 20, 30, 40])
```

```
>>> s
```

```
10    1.0
```

```
20    2.0
```

```
30    NaN
```

```
40    4.0
```

```
dtype: float64
```

```
>>> s.asof(20)
```

```
2.0
```

For a sequence `where`, a Series is returned. The first value is NaN, because the first element of `where` is before the first index value.

```
>>> s.asof([5, 20])
```

```
5      NaN
```

```
20    2.0
```

```
dtype: float64
```

Missing values are not considered. The following is ``2.0``, not NaN, even though NaN is at the index location for ``30``.

```
>>> s.asof(30)
```

```
2.0
```

Take all columns into consideration

```
>>> df = pd.DataFrame({'a': [10, 20, 30, 40, 50],
...                     'b': [None, None, None, None, 500]},
...                     index=pd.DatetimeIndex(['2018-02-27 09:01:00',
...                                             '2018-02-27 09:02:00',
...                                             '2018-02-27 09:03:00',
...                                             '2018-02-27 09:04:00',
...                                             '2018-02-27 09:05:00']))
>>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
...                             '2018-02-27 09:04:30']))
               a    b
2018-02-27 09:03:30 NaN NaN
2018-02-27 09:04:30 NaN NaN
```

Take a single column into consideration

```
>>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
...                             '2018-02-27 09:04:30']),
...          subset=['a'])
               a    b
2018-02-27 09:03:30  30.0 NaN
2018-02-27 09:04:30  40.0 NaN
```

`astype(self, dtype, copy=True, errors='raise', **kwargs)`
 Cast a pandas object to a specified dtype ``dtype``.

Parameters

`dtype` : data type, or dict of column name -> data type

Use a `numpy.dtype` or Python type to cast entire pandas object to the same type. Alternatively, use `{col: dtype, ...}`, where `col` is a column label and `dtype` is a `numpy.dtype` or Python type to cast one or more of the DataFrame's columns to column-specific types.

`copy` : bool, default True

Return a copy when ``copy=True`` (be very careful setting ``copy=False`` as changes to values then may propagate to other pandas objects).

`errors` : {'raise', 'ignore'}, default 'raise'

Control raising of exceptions on invalid data for provided dtype.

- ``raise`` : allow exceptions to be raised

- ``ignore`` : suppress exceptions. On error return original object

.. versionadded:: 0.20.0

`kwargs` : keyword arguments to pass on to the constructor

Returns

casted : same type as caller

See Also

`to_datetime` : Convert argument to datetime.

`to_timedelta` : Convert argument to timedelta.

`to_numeric` : Convert argument to a numeric type.

`numpy.ndarray.astype` : Cast a numpy array to a specified type.

Examples

```
>>> ser = pd.Series([1, 2], dtype='int32')
```

```
>>> ser
```

```
0    1
```

```
1    2
```

```
dtype: int32
```

```
>>> ser.astype('int64')
```

```
0    1
```

```
1    2
```

```
dtype: int64
```

Convert to categorical type:

```
>>> ser.astype('category')
0    1
1    2
dtype: category
Categories (2, int64): [1, 2]
```

Convert to ordered categorical type with custom ordering:

```
>>> cat_dtype = pd.api.types.CategoricalDtype(
...             categories=[2, 1], ordered=True)
>>> ser.astype(cat_dtype)
0    1
1    2
dtype: category
Categories (2, int64): [2 < 1]
```

Note that using ``copy=False`` and changing data on a new pandas object may propagate changes:

```
>>> s1 = pd.Series([1,2])
>>> s2 = s1.astype('int64', copy=False)
>>> s2[0] = 10
>>> s1 # note that s1[0] has changed too
0     10
1      2
dtype: int64
```

```
at_time(self, time, asof=False, axis=None)
    Select values at particular time of day (e.g. 9:30AM).
```

Parameters

```
-----
time : datetime.time or string
axis : {0 or 'index', 1 or 'columns'}, default 0
```

.. versionadded:: 0.24.0

Returns

```
-----
values_at_time : same type as caller
```

Raises

```
-----
TypeError
    If the index is not a :class:`DatetimeIndex`
```

See Also

```
-----
between_time : Select values between particular times of the day.
first : Select initial periods of time series based on a date offset.
last : Select final periods of time series based on a date offset.
DatetimeIndex.indexer_at_time : Get just the index locations for
    values at particular time of the day.
```

Examples

```
-----
>>> i = pd.date_range('2018-04-09', periods=4, freq='12H')
>>> ts = pd.DataFrame({'A': [1,2,3,4]}, index=i)
>>> ts
              A
2018-04-09 00:00:00  1
2018-04-09 12:00:00  2
2018-04-10 00:00:00  3
2018-04-10 12:00:00  4

>>> ts.at_time('12:00')
              A
2018-04-09 12:00:00  2
2018-04-10 12:00:00  4
```

```
between_time(self, start_time, end_time, include_start=True, include_end=True, axis=None)
    one)
```

Select values between particular times of the day (e.g., 9:00-9:30 AM).

By setting ``start_time`` to be later than ``end_time``, you can get the times that are **not** between the two times.

Parameters

start_time : datetime.time or string
end_time : datetime.time or string
include_start : boolean, default True
include_end : boolean, default True
axis : {0 or 'index', 1 or 'columns'}, default 0

.. versionadded:: 0.24.0

Returns

values_between_time : same type as caller

Raises

TypeError

If the index is not a :class:`DatetimeIndex`

See Also

at_time : Select values at a particular time of the day.
first : Select initial periods of time series based on a date offset.
last : Select final periods of time series based on a date offset.
DatetimeIndex.indexer_between_time : Get just the index locations for values between particular times of the day.

Examples

```
>>> i = pd.date_range('2018-04-09', periods=4, freq='1D20min')
>>> ts = pd.DataFrame({'A': [1,2,3,4]}, index=i)
>>> ts
```

```

           A
2018-04-09 00:00:00  1
2018-04-10 00:20:00  2
2018-04-11 00:40:00  3
2018-04-12 01:00:00  4
```

```
>>> ts.between_time('0:15', '0:45')
```

```

           A
2018-04-10 00:20:00  2
2018-04-11 00:40:00  3
```

You get the times that are **not** between two times by setting ``start_time`` later than ``end_time``:

```
>>> ts.between_time('0:45', '0:15')
```

```

           A
2018-04-09 00:00:00  1
2018-04-12 01:00:00  4
```

bfill(self, axis=None, inplace=False, limit=None, downcast=None)
Synonym for :meth:`DataFrame.fillna` with ``method='bfill'``.

bool(self)

Return the bool of a single element PandasObject.

This must be a boolean scalar value, either True or False. Raise a ValueError if the PandasObject does not have exactly 1 element, or that element is not boolean

clip(self, lower=None, upper=None, axis=None, inplace=False, *args, **kwargs)
Trim values at input threshold(s).

Assigns values outside boundary to boundary values. Thresholds can be singular values or array like, and in the latter case theclipping is performed element-wise in the specified axis.

Parameters

`lower` : float or array_like, default None
Minimum threshold value. All values below this threshold will be set to it.

`upper` : float or array_like, default None
Maximum threshold value. All values above this threshold will be set to it.

`axis` : int or string axis name, optional
Align object with lower and upper along the given axis.

`inplace` : boolean, default False
Whether to perform the operation in place on the data.

.. versionadded:: 0.21.0

`*args, **kwargs`
Additional keywords have no effect but might be accepted for compatibility with numpy.

Returns

Series or DataFrame
Same type as calling object with the values outside the clip boundaries replaced

Examples

```
>>> data = {'col_0': [9, -3, 0, -1, 5], 'col_1': [-2, -7, 6, 8, -5]}
>>> df = pd.DataFrame(data)
>>> df
```

	col_0	col_1
0	9	-2
1	-3	-7
2	0	6
3	-1	8
4	5	-5

Clips per column using lower and upper thresholds:

```
>>> df.clip(-4, 6)
```

	col_0	col_1
0	6	-2
1	-3	-4
2	0	6
3	-1	6
4	5	-4

Clips using specific lower and upper thresholds per column element:

```
>>> t = pd.Series([2, -4, -1, 6, 3])
>>> t
```

0	2
1	-4
2	-1
3	6
4	3

dtype: int64

```
>>> df.clip(t, t + 4, axis=0)
```

	col_0	col_1
0	6	2
1	-3	-4
2	0	3
3	6	8
4	5	3

`clip_lower(self, threshold, axis=None, inplace=False)`
Trim values below a given threshold.

.. deprecated:: 0.24.0
Use `clip(lower=threshold)` instead.

Elements below the `threshold` will be changed to match the `threshold` value(s). Threshold can be a single value or an array, in the latter case it performs the truncation element-wise.

Parameters

`threshold` : numeric or array-like

Minimum value allowed. All values below threshold will be set to this value.

- * `float` : every value is compared to `threshold`.

- * `array-like` : The shape of `threshold` should match the object it's compared to. When `self` is a Series, `threshold` should be the length. When `self` is a DataFrame, `threshold` should be 2-D and the same shape as `self` for ``axis=None``, or 1-D and the same length as the axis being compared.

`axis` : {0 or 'index', 1 or 'columns'}, default 0

Align `self` with `threshold` along the given axis.

`inplace` : boolean, default False

Whether to perform the operation in place on the data.

.. versionadded:: 0.21.0

Returns

Series or DataFrame

Original data with values trimmed.

See Also

`Series.clip` : General purpose method to trim Series values to given threshold(s).

`DataFrame.clip` : General purpose method to trim DataFrame values to given threshold(s).

Examples

Series single threshold clipping:

```
>>> s = pd.Series([5, 6, 7, 8, 9])
>>> s.clip(lower=8)
0    8
1    8
2    8
3    8
4    9
dtype: int64
```

Series clipping element-wise using an array of thresholds. `threshold` should be the same length as the Series.

```
>>> elemwise_thresholds = [4, 8, 7, 2, 5]
>>> s.clip(lower=elemwise_thresholds)
0    5
1    8
2    7
3    8
4    9
dtype: int64
```

DataFrames can be compared to a scalar.

```
>>> df = pd.DataFrame({"A": [1, 3, 5], "B": [2, 4, 6]})
>>> df
   A  B
0  1  2
1  3  4
2  5  6
```

```
>>> df.clip(lower=3)
```

```
   A  B
0  3  3
1  3  4
2  5  6
```

Or to an array of values. By default, ``threshold`` should be the same shape as the DataFrame.

```
>>> df.clip(lower=np.array([[3, 4], [2, 2], [6, 2]]))
```

```
   A  B
0  3  4
1  3  4
2  6  6
```

Control how ``threshold`` is broadcast with ``axis``. In this case ``threshold`` should be the same length as the axis specified by ``axis``.

```
>>> df.clip(lower=[3, 3, 5], axis='index')
```

```
   A  B
0  3  3
1  3  4
2  5  6
```

```
>>> df.clip(lower=[4, 5], axis='columns')
```

```
   A  B
0  4  5
1  4  5
2  5  6
```

```
clip_upper(self, threshold, axis=None, inplace=False)
```

Trim values above a given threshold.

.. deprecated:: 0.24.0

Use `clip(upper=threshold)` instead.

Elements above the ``threshold`` will be changed to match the ``threshold`` value(s). Threshold can be a single value or an array, in the latter case it performs the truncation element-wise.

Parameters

`threshold` : numeric or array-like

Maximum value allowed. All values above threshold will be set to this value.

* float : every value is compared to ``threshold``.

* array-like : The shape of ``threshold`` should match the object it's compared to. When ``self`` is a Series, ``threshold`` should be the length. When ``self`` is a DataFrame, ``threshold`` should be 2-D and the same shape as ``self`` for ``axis=None``, or 1-D and the same length as the axis being compared.

`axis` : {0 or 'index', 1 or 'columns'}, default 0

Align object with ``threshold`` along the given axis.

`inplace` : boolean, default False

Whether to perform the operation in place on the data.

.. versionadded:: 0.21.0

Returns

Series or DataFrame

Original data with values trimmed.

See Also

`Series.clip` : General purpose method to trim Series values to given threshold(s).

`DataFrame.clip` : General purpose method to trim DataFrame values to given threshold(s).

Examples

```
>>> s = pd.Series([1, 2, 3, 4, 5])
```

```
>>> s
```

```
0    1
```

```
1    2
```

```
2    3
```

```
3    4
```

```
4    5
```

```
dtype: int64
```

```
>>> s.clip(upper=3)
```

```
0    1
```

```
1    2
```

```
2    3
```

```
3    3
```

```
4    3
```

```
dtype: int64
```

```
>>> elemwise_thresholds = [5, 4, 3, 2, 1]
```

```
>>> elemwise_thresholds
```

```
[5, 4, 3, 2, 1]
```

```
>>> s.clip(upper=elemwise_thresholds)
```

```
0    1
```

```
1    2
```

```
2    3
```

```
3    2
```

```
4    1
```

```
dtype: int64
```

```
convert_objects(self, convert_dates=True, convert_numeric=False, convert_timedeltas=True, copy=True)
```

Attempt to infer better dtype for object columns.

.. deprecated:: 0.21.0

Parameters

`convert_dates` : boolean, default True

If True, convert to date where possible. If 'coerce', force conversion, with unconvertible values becoming NaT.

`convert_numeric` : boolean, default False

If True, attempt to coerce to numbers (including strings), with unconvertible values becoming NaN.

`convert_timedeltas` : boolean, default True

If True, convert to timedelta where possible. If 'coerce', force conversion, with unconvertible values becoming NaT.

`copy` : boolean, default True

If True, return a copy even if no copy is necessary (e.g. no conversion was done). Note: This is meant for internal use, and should not be confused with inplace.

Returns

`converted` : same as input object

See Also

`to_datetime` : Convert argument to datetime.

`to_timedelta` : Convert argument to timedelta.

`to_numeric` : Convert argument to numeric type.

```
copy(self, deep=True)
```

Make a copy of this object's indices and data.

When ``deep=True`` (default), a new object will be created with a copy of the calling object's data and indices. Modifications to the data or indices of the copy will not be reflected in the original object (see notes below).

When ``deep=False``, a new object will be created without copying the calling object's data or index (only references to the data and index are copied). Any changes to the data of the original will be reflected in the shallow copy (and vice versa).

Parameters

deep : bool, default True

Make a deep copy, including a copy of the data and the indices.

With ``deep=False`` neither the indices nor the data are copied.

Returns

copy : Series, DataFrame or Panel

Object type matches caller.

Notes

When ``deep=True``, data is copied but actual Python objects will not be copied recursively, only the reference to the object. This is in contrast to `copy.deepcopy` in the Standard Library, which recursively copies object data (see examples below).

While ``Index`` objects are copied when ``deep=True``, the underlying numpy array is not copied for performance reasons. Since ``Index`` is immutable, the underlying data can be safely shared and a copy is not needed.

Examples

```
>>> s = pd.Series([1, 2], index=["a", "b"])
```

```
>>> s
```

```
a    1
```

```
b    2
```

```
dtype: int64
```

```
>>> s_copy = s.copy()
```

```
>>> s_copy
```

```
a    1
```

```
b    2
```

```
dtype: int64
```

****Shallow copy versus default (deep) copy:****

```
>>> s = pd.Series([1, 2], index=["a", "b"])
```

```
>>> deep = s.copy()
```

```
>>> shallow = s.copy(deep=False)
```

Shallow copy shares data and index with original.

```
>>> s is shallow
```

```
False
```

```
>>> s.values is shallow.values and s.index is shallow.index
```

```
True
```

Deep copy has own copy of data and index.

```
>>> s is deep
```

```
False
```

```
>>> s.values is deep.values or s.index is deep.index
```

```
False
```

Updates to the data shared by shallow copy and original is reflected in both; deep copy remains unchanged.

```
>>> s[0] = 3
```

```
>>> shallow[1] = 4
```

```
>>> s
```

```
a    3
```

```
b    4
```

```
dtype: int64
```

```
>>> shallow
```

```

a    3
b    4
dtype: int64
>>> deep
a    1
b    2
dtype: int64

```

Note that when copying an object containing Python objects, a deep copy will copy the data, but will not do so recursively. Updating a nested data object will be reflected in the deep copy.

```

>>> s = pd.Series([[1, 2], [3, 4]])
>>> deep = s.copy()
>>> s[0][0] = 10
>>> s
0    [10, 2]
1     [3, 4]
dtype: object
>>> deep
0    [10, 2]
1     [3, 4]
dtype: object

```

```

describe(self, percentiles=None, include=None, exclude=None)
Generate descriptive statistics that summarize the central tendency,
dispersion and shape of a dataset's distribution, excluding
`NaN` values.

```

Analyzes both numeric and object series, as well as ``DataFrame`` column sets of mixed data types. The output will vary depending on what is provided. Refer to the notes below for more detail.

Parameters

percentiles : list-like of numbers, optional

The percentiles to include in the output. All should fall between 0 and 1. The default is ``[.25, .5, .75]``, which returns the 25th, 50th, and 75th percentiles.

include : 'all', list-like of dtypes or None (default), optional

A white list of data types to include in the result. Ignored for ``Series``. Here are the options:

- 'all' : All columns of the input will be included in the output.
- A list-like of dtypes : Limits the results to the provided data types.
To limit the result to numeric types submit ``numpy.number``. To limit it instead to object columns submit the ``numpy.object`` data type. Strings can also be used in the style of ``select_dtypes`` (e.g. ``df.describe(include=['O'])``). To select pandas categorical columns, use ``'category'``
- None (default) : The result will include all numeric columns.

exclude : list-like of dtypes or None (default), optional,

A black list of data types to omit from the result. Ignored for ``Series``. Here are the options:

- A list-like of dtypes : Excludes the provided data types from the result. To exclude numeric types submit ``numpy.number``. To exclude object columns submit the data type ``numpy.object``. Strings can also be used in the style of ``select_dtypes`` (e.g. ``df.describe(include=['O'])``). To exclude pandas categorical columns, use ``'category'``
- None (default) : The result will exclude nothing.

Returns

Series or DataFrame

Summary statistics of the Series or Dataframe provided.

See Also

DataFrame.count: Count number of non-NA/null observations.
DataFrame.max: Maximum of the values in the object.
DataFrame.min: Minimum of the values in the object.
DataFrame.mean: Mean of the values.
DataFrame.std: Standard deviation of the observations.
DataFrame.select_dtypes: Subset of a DataFrame including/excluding columns based on their dtype.

Notes

For numeric data, the result's index will include ``count``, ``mean``, ``std``, ``min``, ``max`` as well as lower, ``50`` and upper percentiles. By default the lower percentile is ``25`` and the upper percentile is ``75``. The ``50`` percentile is the same as the median.

For object data (e.g. strings or timestamps), the result's index will include ``count``, ``unique``, ``top``, and ``freq``. The ``top`` is the most common value. The ``freq`` is the most common value's frequency. Timestamps also include the ``first`` and ``last`` items.

If multiple object values have the highest count, then the ``count`` and ``top`` results will be arbitrarily chosen from among those with the highest count.

For mixed data types provided via a ``DataFrame``, the default is to return only an analysis of numeric columns. If the dataframe consists only of object and categorical data without any numeric columns, the default is to return an analysis of both the object and categorical columns. If ``include='all'`` is provided as an option, the result will include a union of attributes of each type.

The ``include`` and ``exclude`` parameters can be used to limit which columns in a ``DataFrame`` are analyzed for the output. The parameters are ignored when analyzing a ``Series``.

Examples

Describing a numeric ``Series``.

```
>>> s = pd.Series([1, 2, 3])
>>> s.describe()
count      3.0
mean       2.0
std        1.0
min        1.0
25%        1.5
50%        2.0
75%        2.5
max        3.0
dtype: float64
```

Describing a categorical ``Series``.

```
>>> s = pd.Series(['a', 'a', 'b', 'c'])
>>> s.describe()
count      4
unique     3
top        a
freq       2
dtype: object
```

Describing a timestamp ``Series``.

```
>>> s = pd.Series([
...     np.datetime64("2000-01-01"),
...     np.datetime64("2010-01-01"),
...     np.datetime64("2010-01-01")
... ])
>>> s.describe()
```

```
count          3
unique         2
top      2010-01-01 00:00:00
freq          2
first      2000-01-01 00:00:00
last      2010-01-01 00:00:00
dtype: object
```

Describing a ``DataFrame``. By default only numeric fields are returned.

```
>>> df = pd.DataFrame({'categorical': pd.Categorical(['d','e','f']),
...                    'numeric': [1, 2, 3],
...                    'object': ['a', 'b', 'c']}
...                    })
>>> df.describe()
```

```
      numeric
count      3.0
mean       2.0
std        1.0
min        1.0
25%        1.5
50%        2.0
75%        2.5
max        3.0
```

Describing all columns of a ``DataFrame`` regardless of data type.

```
>>> df.describe(include='all')
      categorical  numeric object
count           3        3.0     3
unique          3        NaN     3
top             f        NaN     c
freq            1        NaN     1
mean           NaN        2.0    NaN
std            NaN        1.0    NaN
min            NaN        1.0    NaN
25%            NaN        1.5    NaN
50%            NaN        2.0    NaN
75%            NaN        2.5    NaN
max            NaN        3.0    NaN
```

Describing a column from a ``DataFrame`` by accessing it as an attribute.

```
>>> df.numeric.describe()
count      3.0
mean       2.0
std        1.0
min        1.0
25%        1.5
50%        2.0
75%        2.5
max        3.0
Name: numeric, dtype: float64
```

Including only numeric columns in a ``DataFrame`` description.

```
>>> df.describe(include=[np.number])
      numeric
count      3.0
mean       2.0
std        1.0
min        1.0
25%        1.5
50%        2.0
75%        2.5
max        3.0
```

Including only string columns in a ``DataFrame`` description.

```
>>> df.describe(include=[np.object])
```

```

|         object
| count      3
| unique     3
| top        c
| freq       1

```

Including only categorical columns from a ``DataFrame`` description.

```

| >>> df.describe(include=['category'])
|         categorical
| count      3
| unique     3
| top        f
| freq       1

```

Excluding numeric columns from a ``DataFrame`` description.

```

| >>> df.describe(exclude=[np.number])
|         categorical object
| count      3      3
| unique     3      3
| top        f      c
| freq       1      1

```

Excluding object columns from a ``DataFrame`` description.

```

| >>> df.describe(exclude=[np.object])
|         categorical  numeric
| count      3      3.0
| unique     3      NaN
| top        f      NaN
| freq       1      NaN
| mean      NaN      2.0
| std       NaN      1.0
| min       NaN      1.0
| 25%       NaN      1.5
| 50%       NaN      2.0
| 75%       NaN      2.5
| max       NaN      3.0

```

`droplevel(self, level, axis=0)`

Return DataFrame with requested index / column level(s) removed.

.. versionadded:: 0.24.0

Parameters

level : int, str, or list-like

If a string is given, must be the name of a level

If list-like, elements must be names or positional indexes of levels.

axis : {0 or 'index', 1 or 'columns'}, default 0

Returns

DataFrame.droplevel()

Examples

```

| >>> df = pd.DataFrame([
| ...     [1, 2, 3, 4],
| ...     [5, 6, 7, 8],
| ...     [9, 10, 11, 12]
| ... ]).set_index([0, 1]).rename_axis(['a', 'b'])

```

```

| >>> df.columns = pd.MultiIndex.from_tuples([
| ...     ('c', 'e'), ('d', 'f')
| ... ], names=['level_1', 'level_2'])

```

```

| >>> df
| level_1  c  d

```

```

level_2    e    f
a b
1 2        3    4
5 6        7    8
9 10       11   12

```

```
>>> df.droplevel('a')
```

```

level_1    c    d
level_2    e    f
b
2        3    4
6        7    8
10       11   12

```

```
>>> df.droplevel('level2', axis=1)
```

```

level_1    c    d
a b
1 2        3    4
5 6        7    8
9 10       11   12

```

```
equals(self, other)
```

Test whether two objects contain the same elements.

This function allows two Series or DataFrames to be compared against each other to see if they have the same shape and elements. NaNs in the same location are considered equal. The column headers do not need to have the same type, but the elements within the columns must be the same dtype.

Parameters

other : Series or DataFrame

The other Series or DataFrame to be compared with the first.

Returns

bool

True if all elements are the same in both objects, False otherwise.

See Also

Series.eq : Compare two Series objects of the same length

and return a Series where each element is True if the element in each Series is equal, False otherwise.

DataFrame.eq : Compare two DataFrame objects of the same shape and return a DataFrame where each element is True if the respective element in each DataFrame is equal, False otherwise.

assert_series_equal : Return True if left and right Series are equal, False otherwise.

assert_frame_equal : Return True if left and right DataFrames are equal, False otherwise.

numpy.array_equal : Return True if two arrays have the same shape and elements, False otherwise.

Notes

This function requires that the elements have the same dtype as their respective elements in the other Series or DataFrame. However, the column labels do not need to have the same type, as long as they are still considered equal.

Examples

```
>>> df = pd.DataFrame({1: [10], 2: [20]})
```

```
>>> df
```

```

   1  2
0 10 20

```

DataFrames `df` and `exactly_equal` have the same types and values for their elements and column labels, which will return True.

```
>>> exactly_equal = pd.DataFrame({1: [10], 2: [20]})
>>> exactly_equal
   1  2
0 10 20
>>> df.equals(exactly_equal)
True
```

DataFrames `df` and `different_column_type` have the same element types and values, but have different types for the column labels, which will still return `True`.

```
>>> different_column_type = pd.DataFrame({1.0: [10], 2.0: [20]})
>>> different_column_type
   1.0  2.0
0   10   20
>>> df.equals(different_column_type)
True
```

DataFrames `df` and `different_data_type` have different types for the same values for their elements, and will return `False` even though their column labels are the same values and types.

```
>>> different_data_type = pd.DataFrame({1: [10.0], 2: [20.0]})
>>> different_data_type
   1  2
0 10.0 20.0
>>> df.equals(different_data_type)
False
```

`ffill(self, axis=None, inplace=False, limit=None, downcast=None)`
 Synonym for `:meth:`DataFrame.fillna`` with ```method='ffill'```.

`filter(self, items=None, like=None, regex=None, axis=None)`
 Subset rows or columns of dataframe according to labels in the specified index.

Note that this routine does not filter a dataframe on its contents. The filter is applied to the labels of the index.

Parameters

`items` : list-like

List of axis to restrict to (must not all be present).

`like` : string

Keep axis where `"arg in col == True"`.

`regex` : string (regular expression)

Keep axis with `re.search(regex, col) == True`.

`axis` : int or string axis name

The axis to filter on. By default this is the info axis, `'index'` for Series, `'columns'` for DataFrame.

Returns

same type as input object

See Also

`DataFrame.loc`

Notes

The ```items```, ```like```, and ```regex``` parameters are enforced to be mutually exclusive.

```axis``` defaults to the info axis that is used when indexing with ```[]```.

#### Examples

-----

```
>>> df = pd.DataFrame(np.array([[1,2,3], [4,5,6]]),
... index=['mouse', 'rabbit'],
```



```
... columns=['one', 'two', 'three'])
```

```
>>> # select columns by name
```

```
>>> df.filter(items=['one', 'three'])
```

```
 one three
mouse 1 3
rabbit 4 6
```

```
>>> # select columns by regular expression
```

```
>>> df.filter(regex='e$', axis=1)
```

```
 one three
mouse 1 3
rabbit 4 6
```

```
>>> # select rows containing 'bbi'
```

```
>>> df.filter(like='bbi', axis=0)
```

```
 one two three
rabbit 4 5 6
```

```
first(self, offset)
```

Convenience method for subsetting initial periods of time series data based on a date offset.

Parameters

-----

offset : string, DateOffset, dateutil.relativedelta

Returns

-----

subset : same type as caller

Raises

-----

TypeError

If the index is not a :class:`DatetimeIndex`

See Also

-----

last : Select final periods of time series based on a date offset.

at\_time : Select values at a particular time of the day.

between\_time : Select values between particular times of the day.

Examples

-----

```
>>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
```

```
>>> ts = pd.DataFrame({'A': [1,2,3,4]}, index=i)
```

```
>>> ts
```

```
 A
2018-04-09 1
2018-04-11 2
2018-04-13 3
2018-04-15 4
```

Get the rows for the first 3 days:

```
>>> ts.first('3D')
```

```
 A
2018-04-09 1
2018-04-11 2
```

Notice the data for 3 first calendar days were returned, not the first 3 days observed in the dataset, and therefore data for 2018-04-13 was not returned.

```
first_valid_index(self)
```

Return index for first non-NA/null value.

Returns

-----

scalar : type of index

Notes

-----  
If all elements are non-NA/null, returns None.  
Also returns None for empty NDFrame.

get(self, key, default=None)  
Get item from object for given key (DataFrame column, Panel slice, etc.). Returns default value if not found.

Parameters

-----

key : object

Returns

-----

value : same type as items contained in object

get\_dtype\_counts(self)  
Return counts of unique dtypes in this object.

Returns

-----

dtype : Series

Series with the count of columns with each dtype.

See Also

-----

dtypes : Return the dtypes in this object.

Examples

-----

```
>>> a = [['a', 1, 1.0], ['b', 2, 2.0], ['c', 3, 3.0]]
>>> df = pd.DataFrame(a, columns=['str', 'int', 'float'])
>>> df
 str int float
0 a 1 1.0
1 b 2 2.0
2 c 3 3.0

>>> df.get_dtype_counts()
float64 1
int64 1
object 1
dtype: int64
```

get\_ftype\_counts(self)  
Return counts of unique ftypes in this object.

.. deprecated:: 0.23.0

This is useful for SparseDataFrame or for DataFrames containing sparse arrays.

Returns

-----

dtype : Series

Series with the count of columns with each type and sparsity (dense/sparse)

See Also

-----

ftypes : Return ftypes (indication of sparse/dense and dtype) in this object.

Examples

-----

```
>>> a = [['a', 1, 1.0], ['b', 2, 2.0], ['c', 3, 3.0]]
>>> df = pd.DataFrame(a, columns=['str', 'int', 'float'])
>>> df
 str int float
0 a 1 1.0
1 b 2 2.0
2 c 3 3.0
```

```
>>> df.get_ftype_counts() # doctest: +SKIP
float64:dense 1
int64:dense 1
object:dense 1
dtype: int64
```

groupby(self, by=None, axis=0, level=None, as\_index=True, sort=True, group\_keys=True, squeeze=False, observed=False, \*\*kwargs)

Group DataFrame or Series using a mapper or by a Series of columns.

A groupby operation involves some combination of splitting the object, applying a function, and combining the results. This can be used to group large amounts of data and compute operations on these groups.

#### Parameters

**by** : mapping, function, label, or list of labels

Used to determine the groups for the groupby.

If ``by`` is a function, it's called on each value of the object's index. If a dict or Series is passed, the Series or dict VALUES will be used to determine the groups (the Series' values are first aligned; see ``.align()`` method). If an ndarray is passed, the values are used as-is to determine the groups. A label or list of labels may be passed to group by the columns in ``self``. Notice that a tuple is interpreted as a (single) key.

**axis** : {0 or 'index', 1 or 'columns'}, default 0

Split along rows (0) or columns (1).

**level** : int, level name, or sequence of such, default None

If the axis is a MultiIndex (hierarchical), group by a particular level or levels.

**as\_index** : bool, default True

For aggregated output, return object with group labels as the index. Only relevant for DataFrame input. as\_index=False is effectively "SQL-style" grouped output.

**sort** : bool, default True

Sort group keys. Get better performance by turning this off.

Note this does not influence the order of observations within each group. Groupby preserves the order of rows within each group.

**group\_keys** : bool, default True

When calling apply, add group keys to index to identify pieces.

**squeeze** : bool, default False

Reduce the dimensionality of the return type if possible, otherwise return a consistent type.

**observed** : bool, default False

This only applies if any of the groupers are Categoricals.

If True: only show observed values for categorical groupers.

If False: show all values for categorical groupers.

.. versionadded:: 0.23.0

#### \*\*kwargs

Optional, only accepts keyword argument 'mutated' and is passed to groupby.

#### Returns

DataFrameGroupBy or SeriesGroupBy

Depends on the calling object and returns groupby object that contains information about the groups.

#### See Also

**resample** : Convenience method for frequency conversion and resampling of time series.

#### Notes

See the `user guide

<<http://pandas.pydata.org/pandas-docs/stable/groupby.html>>`\_ for more.

## Examples

-----

```
>>> df = pd.DataFrame({'Animal' : ['Falcon', 'Falcon',
... 'Parrot', 'Parrot'],
... 'Max Speed' : [380., 370., 24., 26.]})
```

```
>>> df
 Animal Max Speed
0 Falcon 380.0
1 Falcon 370.0
2 Parrot 24.0
3 Parrot 26.0
>>> df.groupby(['Animal']).mean()
 Max Speed
Animal
Falcon 375.0
Parrot 25.0
```

## \*\*Hierarchical Indexes\*\*

We can groupby different levels of a hierarchical index using the `level` parameter:

```
>>> arrays = [['Falcon', 'Falcon', 'Parrot', 'Parrot'],
... ['Capitve', 'Wild', 'Capitve', 'Wild']]
>>> index = pd.MultiIndex.from_arrays(arrays, names=('Animal', 'Type'))
>>> df = pd.DataFrame({'Max Speed' : [390., 350., 30., 20.]},
... index=index)
```

```
>>> df
 Max Speed
Animal Type
Falcon Capitve 390.0
 Wild 350.0
Parrot Capitve 30.0
 Wild 20.0
>>> df.groupby(level=0).mean()
 Max Speed
Animal
Falcon 370.0
Parrot 25.0
>>> df.groupby(level=1).mean()
 Max Speed
Type
Capitve 210.0
Wild 185.0
```

head(self, n=5)

Return the first `n` rows.

This function returns the first `n` rows for the object based on position. It is useful for quickly testing if your object has the right type of data in it.

## Parameters

-----

n : int, default 5  
Number of rows to select.

## Returns

-----

obj\_head : same type as caller  
The first `n` rows of the caller object.

## See Also

-----

DataFrame.tail: Returns the last `n` rows.

## Examples

-----

```
>>> df = pd.DataFrame({'animal':['alligator', 'bee', 'falcon', 'lion',
... 'monkey', 'parrot', 'shark', 'whale', 'zebra']})
```

```
>>> df
 animal
```

```
0 alligator
1 bee
2 falcon
3 lion
4 monkey
5 parrot
6 shark
7 whale
8 zebra
```

Viewing the first 5 lines

```
>>> df.head()
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
```

Viewing the first `n` lines (three in this case)

```
>>> df.head(3)
 animal
0 alligator
1 bee
2 falcon
```

`infer_objects(self)`

Attempt to infer better dtypes for object columns.

Attempts soft conversion of object-dtyped columns, leaving non-object and unconvertible columns unchanged. The inference rules are the same as during normal Series/DataFrame construction.

.. versionadded:: 0.21.0

Returns

-----

converted : same type as input object

See Also

-----

`to_datetime` : Convert argument to datetime.

`to_timedelta` : Convert argument to timedelta.

`to_numeric` : Convert argument to numeric type.

Examples

-----

```
>>> df = pd.DataFrame({"A": ["a", 1, 2, 3]})
```

```
>>> df = df.iloc[1:]
```

```
>>> df
```

```
 A
1 1
2 2
3 3
```

```
>>> df.dtypes
```

```
A object
```

```
dtype: object
```

```
>>> df.infer_objects().dtypes
```

```
A int64
```

```
dtype: object
```

`interpolate(self, method='linear', axis=0, limit=None, inplace=False, limit_direction='forward', limit_area=None, downcast=None, **kwargs)`

Interpolate values according to different methods.

Please note that only ``method='linear'`` is supported for DataFrame/Series with a MultiIndex.

## Parameters

method : str, default 'linear'

Interpolation technique to use. One of:

- \* 'linear': Ignore the index and treat the values as equally spaced. This is the only method supported on MultiIndexes.
- \* 'time': Works on daily and higher resolution data to interpolate given length of interval.
- \* 'index', 'values': use the actual numerical values of the index.
- \* 'pad': Fill in NaNs using existing values.
- \* 'nearest', 'zero', 'slinear', 'quadratic', 'cubic', 'spline', 'barycentric', 'polynomial': Passed to ``scipy.interpolate.interpld``. Both 'polynomial' and 'spline' require that you also specify an ``order`` (int), e.g. ``df.interpolate(method='polynomial', order=4)``. These use the numerical values of the index.
- \* 'krogh', 'piecewise\_polynomial', 'spline', 'pchip', 'akima': Wrappers around the SciPy interpolation methods of similar names. See ``Notes``.
- \* 'from\_derivatives': Refers to ``scipy.interpolate.BPoly.from_derivatives`` which replaces 'piecewise\_polynomial' interpolation method in scipy 0.18.

.. versionadded:: 0.18.1

Added support for the 'akima' method.

Added interpolate method 'from\_derivatives' which replaces 'piecewise\_polynomial' in SciPy 0.18; backwards-compatible with SciPy < 0.18

axis : {0 or 'index', 1 or 'columns', None}, default None

Axis to interpolate along.

limit : int, optional

Maximum number of consecutive NaNs to fill. Must be greater than 0.

inplace : bool, default False

Update the data in place if possible.

limit\_direction : {'forward', 'backward', 'both'}, default 'forward'

If limit is specified, consecutive NaNs will be filled in this direction.

limit\_area : {'None', 'inside', 'outside'}, default None

If limit is specified, consecutive NaNs will be filled with this restriction.

\* ``None``: No fill restriction.

\* 'inside': Only fill NaNs surrounded by valid values (interpolate).

\* 'outside': Only fill NaNs outside valid values (extrapolate).

.. versionadded:: 0.21.0

downcast : optional, 'infer' or None, defaults to None

Downcast dtypes if possible.

\*\*kwargs

Keyword arguments to pass on to the interpolating function.

## Returns

Series or DataFrame

Returns the same object type as the caller, interpolated at some or all ``NaN`` values

## See Also

fillna : Fill missing values using different methods.

scipy.interpolate.Akima1DInterpolator : Piecewise cubic polynomials (Akima interpolator).

scipy.interpolate.BPoly.from\_derivatives : Piecewise polynomial in the Bernstein basis.

```
scipy.interpolate.interpld : Interpolate a 1-D function.
scipy.interpolate.KroghInterpolator : Interpolate polynomial (Krogh
 interpolator).
scipy.interpolate.PchipInterpolator : PCHIP 1-d monotonic cubic
 interpolation.
scipy.interpolate.CubicSpline : Cubic spline data interpolator.
```

#### Notes

-----

The 'krogh', 'piecewise\_polynomial', 'spline', 'pchip' and 'akima' methods are wrappers around the respective SciPy implementations of similar names. These use the actual numerical values of the index.

For more information on their behavior, see the

`SciPy documentation

<<http://docs.scipy.org/doc/scipy/reference/interpolate.html#univariate-interpolat>

ion>`\_\_

and `SciPy tutorial

<<http://docs.scipy.org/doc/scipy/reference/tutorial/interpolate.html>>`\_\_.

#### Examples

-----

Filling in ``NaN`` in a :class:`~pandas.Series` via linear interpolation.

```
>>> s = pd.Series([0, 1, np.nan, 3])
```

```
>>> s
```

```
0 0.0
```

```
1 1.0
```

```
2 NaN
```

```
3 3.0
```

```
dtype: float64
```

```
>>> s.interpolate()
```

```
0 0.0
```

```
1 1.0
```

```
2 2.0
```

```
3 3.0
```

```
dtype: float64
```

Filling in ``NaN`` in a Series by padding, but filling at most two consecutive ``NaN`` at a time.

```
>>> s = pd.Series([np.nan, "single_one", np.nan,
... "fill_two_more", np.nan, np.nan, np.nan,
... 4.71, np.nan])
```

```
>>> s
```

```
0 NaN
```

```
1 single_one
```

```
2 NaN
```

```
3 fill_two_more
```

```
4 NaN
```

```
5 NaN
```

```
6 NaN
```

```
7 4.71
```

```
8 NaN
```

```
dtype: object
```

```
>>> s.interpolate(method='pad', limit=2)
```

```
0 NaN
```

```
1 single_one
```

```
2 single_one
```

```
3 fill_two_more
```

```
4 fill_two_more
```

```
5 fill_two_more
```

```
6 NaN
```

```
7 4.71
```

```
8 4.71
```

```
dtype: object
```

Filling in ``NaN`` in a Series via polynomial interpolation or splines:

Both 'polynomial' and 'spline' methods require that you also specify an ``order`` (int).

```
>>> s = pd.Series([0, 2, np.nan, 8])
```

```
>>> s.interpolate(method='polynomial', order=2)
0 0.000000
1 2.000000
2 4.666667
3 8.000000
dtype: float64
```

Fill the DataFrame forward (that is, going down) along each column using linear interpolation.

Note how the last entry in column 'a' is interpolated differently, because there is no entry after it to use for interpolation. Note how the first entry in column 'b' remains ``NaN``, because there is no entry before it to use for interpolation.

```
>>> df = pd.DataFrame([(0.0, np.nan, -1.0, 1.0),
... (np.nan, 2.0, np.nan, np.nan),
... (2.0, 3.0, np.nan, 9.0),
... (np.nan, 4.0, -4.0, 16.0)],
... columns=list('abcd'))
>>> df
 a b c d
0 0.0 NaN -1.0 1.0
1 NaN 2.0 NaN NaN
2 2.0 3.0 NaN 9.0
3 NaN 4.0 -4.0 16.0
>>> df.interpolate(method='linear', limit_direction='forward', axis=0)
 a b c d
0 0.0 NaN -1.0 1.0
1 1.0 2.0 -2.0 5.0
2 2.0 3.0 -3.0 9.0
3 2.0 4.0 -4.0 16.0
```

Using polynomial interpolation.

```
>>> df['d'].interpolate(method='polynomial', order=2)
0 1.0
1 4.0
2 9.0
3 16.0
Name: d, dtype: float64
```

`last(self, offset)`

Convenience method for subsetting final periods of time series data based on a date offset.

Parameters

-----

`offset` : string, DateOffset, dateutil.relativedelta

Returns

-----

subset : same type as caller

Raises

-----

`TypeError`

If the index is not a :class:`DatetimeIndex`

See Also

-----

`first` : Select initial periods of time series based on a date offset.

`at_time` : Select values at a particular time of the day.

`between_time` : Select values between particular times of the day.

Examples

-----

```
>>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
>>> ts = pd.DataFrame({'A': [1,2,3,4]}, index=i)
>>> ts
```

```

 A
2018-04-09 1
```



```
2018-04-11 2
2018-04-13 3
2018-04-15 4
```

Get the rows for the last 3 days:

```
>>> ts.last('3D')
 A
2018-04-13 3
2018-04-15 4
```

Notice the data for 3 last calendar days were returned, not the last 3 observed days in the dataset, and therefore data for 2018-04-11 was not returned.

`last_valid_index(self)`

Return index for last non-NA/null value.

Returns

-----

scalar : type of index

Notes

-----

If all elements are non-NA/null, returns None.

Also returns None for empty NDFrame.

`mask(self, cond, other=nan, inplace=False, axis=None, level=None, errors='raise', try_cast=False, raise_on_error=None)`

Replace values where the condition is True.

Parameters

-----

`cond` : boolean NDFrame, array-like, or callable

Where ``cond`` is False, keep the original value. Where

True, replace with corresponding value from ``other``.

If ``cond`` is callable, it is computed on the NDFrame and

should return boolean NDFrame or array. The callable must

not change input NDFrame (though pandas doesn't check it).

.. versionadded:: 0.18.1

A callable can be used as `cond`.

`other` : scalar, NDFrame, or callable

Entries where ``cond`` is True are replaced with

corresponding value from ``other``.

If `other` is callable, it is computed on the NDFrame and

should return scalar or NDFrame. The callable must not

change input NDFrame (though pandas doesn't check it).

.. versionadded:: 0.18.1

A callable can be used as `other`.

`inplace` : boolean, default False

Whether to perform the operation in place on the data.

`axis` : int, default None

Alignment axis if needed.

`level` : int, default None

Alignment level if needed.

`errors` : str, {'raise', 'ignore'}, default 'raise'

Note that currently this parameter won't affect

the results and will always coerce to a suitable dtype.

- 'raise' : allow exceptions to be raised.

- 'ignore' : suppress exceptions. On error return original object.

`try_cast` : boolean, default False

Try to cast the result back to the input type (if possible).

`raise_on_error` : boolean, default True

Whether to raise on invalid data types (e.g. trying to where on strings).

.. deprecated:: 0.21.0

Use ``errors``.

## Returns

-----

wh : same type as caller

## See Also

-----

:func:`DataFrame.where` : Return an object of same shape as  
self.

## Notes

-----

The mask method is an application of the if-then idiom. For each element in the calling DataFrame, if ``cond`` is ``False`` the element is used; otherwise the corresponding element from the DataFrame ``other`` is used.

The signature for :func:`DataFrame.where` differs from :func:`numpy.where`. Roughly ``df1.where(m, df2)`` is equivalent to ``np.where(m, df1, df2)``.

For further details and examples see the ``mask`` documentation in :ref:`indexing <indexing.where\_mask>`.

## Examples

-----

```
>>> s = pd.Series(range(5))
```

```
>>> s.where(s > 0)
```

```
0 NaN
```

```
1 1.0
```

```
2 2.0
```

```
3 3.0
```

```
4 4.0
```

```
dtype: float64
```

```
>>> s.mask(s > 0)
```

```
0 0.0
```

```
1 NaN
```

```
2 NaN
```

```
3 NaN
```

```
4 NaN
```

```
dtype: float64
```

```
>>> s.where(s > 1, 10)
```

```
0 10
```

```
1 10
```

```
2 2
```

```
3 3
```

```
4 4
```

```
dtype: int64
```

```
>>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A', 'B'])
```

```
>>> m = df % 3 == 0
```

```
>>> df.where(m, -df)
```

```
 A B
```

```
0 0 -1
```

```
1 -2 3
```

```
2 -4 -5
```

```
3 6 -7
```

```
4 -8 9
```

```
>>> df.where(m, -df) == np.where(m, df, -df)
```

```
 A B
```

```
0 True True
```

```
1 True True
```

```
2 True True
```

```
3 True True
```

```
4 True True
```

```
>>> df.where(m, -df) == df.mask(~m, -df)
```

```
 A B
```

```

0 True True
1 True True
2 True True
3 True True
4 True True

```

```

pct_change(self, periods=1, fill_method='pad', limit=None, freq=None, **kwargs)
 Percentage change between the current and a prior element.

```

Computes the percentage change from the immediately previous row by default. This is useful in comparing the percentage of change in a time series of elements.

#### Parameters

-----

```

periods : int, default 1
 Periods to shift for forming percent change.
fill_method : str, default 'pad'
 How to handle NAs before computing percent changes.
limit : int, default None
 The number of consecutive NAs to fill before stopping.
freq : DateOffset, timedelta, or offset alias string, optional
 Increment to use from time series API (e.g. 'M' or BDay()).
**kwargs
 Additional keyword arguments are passed into
 `DataFrame.shift` or `Series.shift`.

```

#### Returns

-----

```

chg : Series or DataFrame
 The same type as the calling object.

```

#### See Also

-----

```

Series.diff : Compute the difference of two elements in a Series.
DataFrame.diff : Compute the difference of two elements in a DataFrame.
Series.shift : Shift the index by some number of periods.
DataFrame.shift : Shift the index by some number of periods.

```

#### Examples

-----

```

Series

```

```

>>> s = pd.Series([90, 91, 85])
>>> s
0 90
1 91
2 85
dtype: int64

```

```

>>> s.pct_change()
0 NaN
1 0.011111
2 -0.065934
dtype: float64

```

```

>>> s.pct_change(periods=2)
0 NaN
1 NaN
2 -0.055556
dtype: float64

```

See the percentage change in a Series where filling NAs with last valid observation forward to next valid.

```

>>> s = pd.Series([90, 91, None, 85])
>>> s
0 90.0
1 91.0
2 NaN
3 85.0
dtype: float64

```

```
>>> s.pct_change(fill_method='ffill')
0 NaN
1 0.011111
2 0.000000
3 -0.065934
dtype: float64
```

**\*\*DataFrame\*\***

Percentage change in French franc, Deutsche Mark, and Italian lira from 1980-01-01 to 1980-03-01.

```
>>> df = pd.DataFrame({
... 'FR': [4.0405, 4.0963, 4.3149],
... 'GR': [1.7246, 1.7482, 1.8519],
... 'IT': [804.74, 810.01, 860.13]},
... index=['1980-01-01', '1980-02-01', '1980-03-01'])
>>> df
```

	FR	GR	IT
1980-01-01	4.0405	1.7246	804.74
1980-02-01	4.0963	1.7482	810.01
1980-03-01	4.3149	1.8519	860.13

```
>>> df.pct_change()
 FR GR IT
1980-01-01 NaN NaN NaN
1980-02-01 0.013810 0.013684 0.006549
1980-03-01 0.053365 0.059318 0.061876
```

Percentage of change in GOOG and APPL stock volume. Shows computing the percentage change between columns.

```
>>> df = pd.DataFrame({
... '2016': [1769950, 30586265],
... '2015': [1500923, 40912316],
... '2014': [1371819, 41403351]},
... index=['GOOG', 'APPL'])
>>> df
```

	2016	2015	2014
GOOG	1769950	1500923	1371819
APPL	30586265	40912316	41403351

```
>>> df.pct_change(axis='columns')
 2016 2015 2014
GOOG NaN -0.151997 -0.086016
APPL NaN 0.337604 0.012002
```

```
pipe(self, func, *args, **kwargs)
Apply func(self, *args, **kwargs).
```

Parameters

-----

```
func : function
 function to apply to the NDFrame.
 ``args``, and ``kwargs`` are passed into ``func``.
 Alternatively a ``(callable, data_keyword)`` tuple where
 ``data_keyword`` is a string indicating the keyword of
 ``callable`` that expects the NDFrame.
args : iterable, optional
 positional arguments passed into ``func``.
kwargs : mapping, optional
 a dictionary of keyword arguments passed into ``func``.
```

Returns

-----

```
object : the return type of ``func``.
```

See Also

-----

```
DataFrame.apply
DataFrame.applymap
```

```
Series.map
```

```
Notes
```

```

```

Use ``.pipe`` when chaining together functions that expect Series, DataFrames or GroupBy objects. Instead of writing

```
>>> f(g(h(df), arg1=a), arg2=b, arg3=c)
```

You can write

```
>>> (df.pipe(h)
... .pipe(g, arg1=a)
... .pipe(f, arg2=b, arg3=c)
...)
```

If you have a function that takes the data as (say) the second argument, pass a tuple indicating which keyword expects the data. For example, suppose ``f`` takes its data as ``arg2``:

```
>>> (df.pipe(h)
... .pipe(g, arg1=a)
... .pipe((f, 'arg2'), arg1=a, arg3=c)
...)
```

```
pop(self, item)
```

Return item and drop from frame. Raise KeyError if not found.

```
Parameters
```

```

```

```
item : str
```

Column label to be popped

```
Returns
```

```

```

```
popped : Series
```

```
Examples
```

```

```

```
>>> df = pd.DataFrame([('falcon', 'bird', 389.0),
... ('parrot', 'bird', 24.0),
... ('lion', 'mammal', 80.5),
... ('monkey', 'mammal', np.nan)],
... columns=('name', 'class', 'max_speed'))
```

```
>>> df
 name class max_speed
0 falcon bird 389.0
1 parrot bird 24.0
2 lion mammal 80.5
3 monkey mammal NaN
```

```
>>> df.pop('class')
0 bird
1 bird
2 mammal
3 mammal
Name: class, dtype: object
```

```
>>> df
 name max_speed
0 falcon 389.0
1 parrot 24.0
2 lion 80.5
3 monkey NaN
```

```
rank(self, axis=0, method='average', numeric_only=None, na_option='keep', ascending=True, pct=False)
```

Compute numerical data ranks (1 through n) along axis. Equal values are assigned a rank that is the average of the ranks of those values.

```
Parameters
```

```

axis : {0 or 'index', 1 or 'columns'}, default 0
 index to direct ranking
method : {'average', 'min', 'max', 'first', 'dense'}
 * average: average rank of group
 * min: lowest rank in group
 * max: highest rank in group
 * first: ranks assigned in order they appear in the array
 * dense: like 'min', but rank always increases by 1 between groups
numeric_only : boolean, default None
 Include only float, int, boolean data. Valid only for DataFrame or
 Panel objects
na_option : {'keep', 'top', 'bottom'}
 * keep: leave NA values where they are
 * top: smallest rank if ascending
 * bottom: smallest rank if descending
ascending : boolean, default True
 False for ranks by high (1) to low (N)
pct : boolean, default False
 Computes percentage rank of data

```

Returns

```

ranks : same type as caller

```

```

reindex_like(self, other, method=None, copy=True, limit=None, tolerance=None)
Return an object with matching indices as other object.

```

Conform the object to the same index on all axes. Optional filling logic, placing NaN in locations having no value in the previous index. A new object is produced unless the new index is equivalent to the current one and copy=False.

Parameters

```

other : Object of the same data type
 Its row and column indices are used to define the new indices
 of this object.
method : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}
 Method to use for filling holes in reindexed DataFrame.
 Please note: this is only applicable to DataFrames/Series with a
 monotonically increasing/decreasing index.

 * None (default): don't fill gaps
 * pad / ffill: propagate last valid observation forward to next
 valid
 * backfill / bfill: use next valid observation to fill gap
 * nearest: use nearest valid observations to fill gap

copy : bool, default True
 Return a new object, even if the passed indexes are the same.
limit : int, default None
 Maximum number of consecutive labels to fill for inexact matches.
tolerance : optional
 Maximum distance between original and new labels for inexact
 matches. The values of the index at the matching locations must
 satisfy the equation ``abs(index[indexer] - target) <= tolerance``.

 Tolerance may be a scalar value, which applies the same tolerance
 to all values, or list-like, which applies variable tolerance per
 element. List-like includes list, tuple, array, Series, and must be
 the same size as the index and its dtype must exactly match the
 index's type.

 .. versionadded:: 0.21.0 (list-like tolerance)

```

Returns

```

Series or DataFrame
 Same type as caller, but with changed indices on each axis.

```

See Also

```

DataFrame.set_index : Set row labels.
DataFrame.reset_index : Remove row labels or move them to new columns.
DataFrame.reindex : Change to new indices or expand indices.
```

#### Notes

Same as calling

```
``.reindex(index=other.index, columns=other.columns,...)``.
```

#### Examples

```

>>> df1 = pd.DataFrame([[24.3, 75.7, 'high'],
... [31, 87.8, 'high'],
... [22, 71.6, 'medium'],
... [35, 95, 'medium']],
... columns=['temp_celsius', 'temp_fahrenheit', 'windspeed'],
... index=pd.date_range(start='2014-02-12',
... end='2014-02-15', freq='D'))
```

```
>>> df1
 temp_celsius temp_fahrenheit windspeed
2014-02-12 24.3 75.7 high
2014-02-13 31.0 87.8 high
2014-02-14 22.0 71.6 medium
2014-02-15 35.0 95.0 medium
```

```
>>> df2 = pd.DataFrame([[28, 'low'],
... [30, 'low'],
... [35.1, 'medium']],
... columns=['temp_celsius', 'windspeed'],
... index=pd.DatetimeIndex(['2014-02-12', '2014-02-13',
... '2014-02-15']))
```

```
>>> df2
 temp_celsius windspeed
2014-02-12 28.0 low
2014-02-13 30.0 low
2014-02-15 35.1 medium
```

```
>>> df2.reindex_like(df1)
 temp_celsius temp_fahrenheit windspeed
2014-02-12 28.0 NaN low
2014-02-13 30.0 NaN low
2014-02-14 NaN NaN NaN
2014-02-15 35.1 NaN medium
```

```
rename_axis(self, mapper=None, index=None, columns=None, axis=None, copy=True, inplace=False)
```

Set the name of the axis for the index or columns.

#### Parameters

mapper : scalar, list-like, optional

Value to set the axis name attribute.

index, columns : scalar, list-like, dict-like or function, optional

A scalar, list-like, dict-like or functions transformations to apply to that axis' values.

Use either ``mapper`` and ``axis`` to specify the axis to target with ``mapper``, or ``index`` and/or ``columns``.

.. versionchanged:: 0.24.0

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis to rename.

copy : bool, default True

Also copy underlying data.

inplace : bool, default False

Modifies the object directly, instead of creating a new Series or DataFrame.

## Returns

-----

Series, DataFrame, or None

The same type as the caller or None if `inplace` is True.

## See Also

-----

Series.rename : Alter Series index labels or name.

DataFrame.rename : Alter DataFrame index labels or name.

Index.rename : Set new names on index.

## Notes

-----

Prior to version 0.21.0, ``rename\_axis`` could also be used to change the axis *labels* by passing a mapping or scalar. This behavior is deprecated and will be removed in a future version. Use ``rename`` instead.

``DataFrame.rename\_axis`` supports two calling conventions

\* ``rename\_axis(index=index\_mapper, columns=columns\_mapper, ...)``

\* ``rename\_axis mapper, axis={'index', 'columns'}, ...)``

The first calling convention will only modify the names of the index and/or the names of the Index object that is the columns. In this case, the parameter ``copy`` is ignored.

The second calling convention will modify the names of the the corresponding index if mapper is a list or a scalar. However, if mapper is dict-like or a function, it will use the deprecated behavior of modifying the axis *labels*.

We *highly* recommend using keyword arguments to clarify your intent.

## Examples

-----

**Series**

```
>>> s = pd.Series(["dog", "cat", "monkey"])
```

```
>>> s
```

```
0 dog
```

```
1 cat
```

```
2 monkey
```

```
dtype: object
```

```
>>> s.rename_axis("animal")
```

```
animal
```

```
0 dog
```

```
1 cat
```

```
2 monkey
```

```
dtype: object
```

**DataFrame**

```
>>> df = pd.DataFrame({"num_legs": [4, 4, 2],
```

```
... "num_arms": [0, 0, 2]},
```

```
... ["dog", "cat", "monkey"])
```

```
>>> df
```

```
 num_legs num_arms
```

```
dog 4 0
```

```
cat 4 0
```

```
monkey 2 2
```

```
>>> df = df.rename_axis("animal")
```

```
>>> df
```

```
 num_legs num_arms
```

```
animal
```

```
dog 4 0
```

```
cat 4 0
```

```
monkey 2 2
```

```
>>> df = df.rename_axis("limbs", axis="columns")
```

```
>>> df
```



	limbs	num_legs	num_arms
animal			
dog		4	0
cat		4	0
monkey		2	2

**\*\*MultiIndex\*\***

```
>>> df.index = pd.MultiIndex.from_product(['mammal'],
... ['dog', 'cat', 'monkey']],
... names=['type', 'name'])
```

```
>>> df
limbs num_legs num_arms
type name
mammal dog 4 0
 cat 4 0
 monkey 2 2
```

```
>>> df.rename_axis(index={'type': 'class'})
```

```
limbs num_legs num_arms
class name
mammal dog 4 0
 cat 4 0
 monkey 2 2
```

```
>>> df.rename_axis(columns=str.upper)
```

```
LIMBS num_legs num_arms
type name
mammal dog 4 0
 cat 4 0
 monkey 2 2
```

```
resample(self, rule, how=None, axis=0, fill_method=None, closed=None, label=None, convention='start', kind=None, loffset=None, limit=None, base=0, on=None, level=None)
Resample time-series data.
```

Convenience method for frequency conversion and resampling of time series. Object must have a datetime-like index (``DatetimeIndex``, ``PeriodIndex``, or ``TimedeltaIndex``), or pass datetime-like values to the ``on`` or ``level`` keyword.

Parameters

-----

rule : str

The offset string or object representing target conversion.

how : str

Method for down/re-sampling, default to 'mean' for downsampling.

.. deprecated:: 0.18.0

The new syntax is ```.resample(...).mean()``, or ``.resample(...).apply(<func>)```

axis : {0 or 'index', 1 or 'columns'}, default 0

Which axis to use for up- or down-sampling. For ``Series`` this will default to 0, i.e. along the rows. Must be ``DatetimeIndex``, ``TimedeltaIndex`` or ``PeriodIndex``.

fill\_method : str, default None

Filling method for upsampling.

.. deprecated:: 0.18.0

The new syntax is ```.resample(...).<func>()``, e.g. ``.resample(...).pad()```

closed : {'right', 'left'}, default None

Which side of bin interval is closed. The default is 'left' for all frequency offsets except for 'M', 'A', 'Q', 'BM', 'BA', 'BQ', and 'W' which all have a default of 'right'.

label : {'right', 'left'}, default None

Which bin edge label to label bucket with. The default is 'left' for all frequency offsets except for 'M', 'A', 'Q', 'BM', 'BA', 'BQ', and 'W' which all have a default of 'right'.

convention : {'start', 'end', 's', 'e'}, default 'start'

For ``PeriodIndex`` only, controls whether to use the start or end of ``rule``.

```

kind : {'timestamp', 'period'}, optional, default None
 Pass 'timestamp' to convert the resulting index to a
 `DateTimeIndex` or 'period' to convert it to a `PeriodIndex`.
 By default the input representation is retained.
loffset : timedelta, default None
 Adjust the resampled time labels.
limit : int, default None
 Maximum size gap when reindexing with `fill_method`.

 .. deprecated:: 0.18.0
base : int, default 0
 For frequencies that evenly subdivide 1 day, the "origin" of the
 aggregated intervals. For example, for '5min' frequency, base could
 range from 0 through 4. Defaults to 0.
on : str, optional
 For a DataFrame, column to use instead of index for resampling.
 Column must be datetime-like.

 .. versionadded:: 0.19.0

level : str or int, optional
 For a MultiIndex, level (name or number) to use for
 resampling. `level` must be datetime-like.

 .. versionadded:: 0.19.0

```

#### Returns

-----

Resampler object

#### See Also

-----

```

groupby : Group by mapping, function, label, or list of labels.
Series.resample : Resample a Series.
DataFrame.resample: Resample a DataFrame.

```

#### Notes

-----

See the `user guide

<<http://pandas.pydata.org/pandas-docs/stable/timeseries.html#resampling>>`\_ for more.

To learn more about the offset strings, please see `this link

<<http://pandas.pydata.org/pandas-docs/stable/timeseries.html#offset-aliases>>`\_\_.

#### Examples

-----

Start by creating a series with 9 one minute timestamps.

```

>>> index = pd.date_range('1/1/2000', periods=9, freq='T')
>>> series = pd.Series(range(9), index=index)
>>> series
2000-01-01 00:00:00 0
2000-01-01 00:01:00 1
2000-01-01 00:02:00 2
2000-01-01 00:03:00 3
2000-01-01 00:04:00 4
2000-01-01 00:05:00 5
2000-01-01 00:06:00 6
2000-01-01 00:07:00 7
2000-01-01 00:08:00 8
Freq: T, dtype: int64

```

Downsample the series into 3 minute bins and sum the values of the timestamps falling into a bin.

```

>>> series.resample('3T').sum()
2000-01-01 00:00:00 3
2000-01-01 00:03:00 12
2000-01-01 00:06:00 21
Freq: 3T, dtype: int64

```

Downsample the series into 3 minute bins as above, but label each bin using the right edge instead of the left. Please note that the value in the bucket used as the label is not included in the bucket, which it labels. For example, in the original series the bucket ``2000-01-01 00:03:00`` contains the value 3, but the summed value in the resampled bucket with the label ``2000-01-01 00:03:00`` does not include 3 (if it did, the summed value would be 6, not 3). To include this value close the right side of the bin interval as illustrated in the example below this one.

```
>>> series.resample('3T', label='right').sum()
2000-01-01 00:03:00 3
2000-01-01 00:06:00 12
2000-01-01 00:09:00 21
Freq: 3T, dtype: int64
```

Downsample the series into 3 minute bins as above, but close the right side of the bin interval.

```
>>> series.resample('3T', label='right', closed='right').sum()
2000-01-01 00:00:00 0
2000-01-01 00:03:00 6
2000-01-01 00:06:00 15
2000-01-01 00:09:00 15
Freq: 3T, dtype: int64
```

Upsample the series into 30 second bins.

```
>>> series.resample('30S').asfreq()[0:5] # Select first 5 rows
2000-01-01 00:00:00 0.0
2000-01-01 00:00:30 NaN
2000-01-01 00:01:00 1.0
2000-01-01 00:01:30 NaN
2000-01-01 00:02:00 2.0
Freq: 30S, dtype: float64
```

Upsample the series into 30 second bins and fill the ``NaN`` values using the ``pad`` method.

```
>>> series.resample('30S').pad()[0:5]
2000-01-01 00:00:00 0
2000-01-01 00:00:30 0
2000-01-01 00:01:00 1
2000-01-01 00:01:30 1
2000-01-01 00:02:00 2
Freq: 30S, dtype: int64
```

Upsample the series into 30 second bins and fill the ``NaN`` values using the ``bfill`` method.

```
>>> series.resample('30S').bfill()[0:5]
2000-01-01 00:00:00 0
2000-01-01 00:00:30 1
2000-01-01 00:01:00 1
2000-01-01 00:01:30 2
2000-01-01 00:02:00 2
Freq: 30S, dtype: int64
```

Pass a custom function via ``apply``

```
>>> def custom_resampler(array_like):
... return np.sum(array_like) + 5
...
>>> series.resample('3T').apply(custom_resampler)
2000-01-01 00:00:00 8
2000-01-01 00:03:00 17
2000-01-01 00:06:00 26
Freq: 3T, dtype: int64
```

For a Series with a PeriodIndex, the keyword ``convention`` can be used to control whether to use the start or end of ``rule``.



For a DataFrame with MultiIndex, the keyword `level` can be used to specify on which level the resampling needs to take place.

```
>>> days = pd.date_range('1/1/2000', periods=4, freq='D')
>>> d2 = dict({'price': [10, 11, 9, 13, 14, 18, 17, 19],
... 'volume': [50, 60, 40, 100, 50, 100, 40, 50]})
>>> df2 = pd.DataFrame(d2,
... index=pd.MultiIndex.from_product([days,
... ['morning',
... 'afternoon']])
>>> df2
```

		price	volume
2000-01-01	morning	10	50
	afternoon	11	60
2000-01-02	morning	9	40
	afternoon	13	100
2000-01-03	morning	14	50
	afternoon	18	100
2000-01-04	morning	17	40
	afternoon	19	50

```
>>> df2.resample('D', level=0).sum()
 price volume
2000-01-01 21 110
2000-01-02 22 140
2000-01-03 32 150
2000-01-04 36 90
```

sample(self, n=None, frac=None, replace=False, weights=None, random\_state=None, axis=None)

Return a random sample of items from an axis of object.

You can use `random\_state` for reproducibility.

Parameters

-----

n : int, optional

Number of items from axis to return. Cannot be used with `frac`.

Default = 1 if `frac` = None.

frac : float, optional

Fraction of axis items to return. Cannot be used with `n`.

replace : bool, default False

Sample with or without replacement.

weights : str or ndarray-like, optional

Default 'None' results in equal probability weighting.

If passed a Series, will align with target object on index. Index

values in weights not found in sampled object will be ignored and

index values in sampled object not in weights will be assigned

weights of zero.

If called on a DataFrame, will accept the name of a column

when axis = 0.

Unless weights are a Series, weights must be same length as axis

being sampled.

If weights do not sum to 1, they will be normalized to sum to 1.

Missing values in the weights column will be treated as zero.

Infinite values not allowed.

random\_state : int or numpy.random.RandomState, optional

Seed for the random number generator (if int), or numpy RandomState object.

axis : int or string, optional

Axis to sample. Accepts axis number or name. Default is stat axis for given data type (0 for Series and DataFrames, 1 for Panels).

Returns

-----

Series or DataFrame

A new object of same type as caller containing `n` items randomly sampled from the caller object.

See Also

-----

numpy.random.choice: Generates a random sample from a given 1-D numpy

array.

## Examples

-----

```
>>> df = pd.DataFrame({'num_legs': [2, 4, 8, 0],
... 'num_wings': [2, 0, 0, 0],
... 'num_specimen_seen': [10, 2, 1, 8]},
... index=['falcon', 'dog', 'spider', 'fish'])
```

```
>>> df
 num_legs num_wings num_specimen_seen
falcon 2 2 10
dog 4 0 2
spider 8 0 1
fish 0 0 8
```

Extract 3 random elements from the ``Series`` ``df['num\_legs']``:  
Note that we use ``random\_state`` to ensure the reproducibility of  
the examples.

```
>>> df['num_legs'].sample(n=3, random_state=1)
fish 0
spider 8
falcon 2
Name: num_legs, dtype: int64
```

A random 50% sample of the ``DataFrame`` with replacement:

```
>>> df.sample(frac=0.5, replace=True, random_state=1)
 num_legs num_wings num_specimen_seen
dog 4 0 2
fish 0 0 8
```

Using a DataFrame column as weights. Rows with larger value in the  
``num\_specimen\_seen`` column are more likely to be sampled.

```
>>> df.sample(n=2, weights='num_specimen_seen', random_state=1)
 num_legs num_wings num_specimen_seen
falcon 2 2 10
fish 0 0 8
```

`select(self, crit, axis=0)`

Return data corresponding to axis labels matching criteria.

.. deprecated:: 0.21.0

Use `df.loc[df.index.map(crit)]` to select via labels

## Parameters

-----

`crit` : function

To be called on each index (label). Should return True or False

`axis` : int

## Returns

-----

`selection` : same type as caller

`set_axis(self, labels, axis=0, inplace=None)`

Assign desired index to given axis.

Indexes for column or row labels can be changed by assigning  
a list-like or Index.

.. versionchanged:: 0.21.0

The signature is now ``labels`` and ``axis``, consistent with  
the rest of pandas API. Previously, the ``axis`` and ``labels``  
arguments were respectively the first and second positional  
arguments.

## Parameters

-----

`labels` : list-like, Index

The values for the new index.

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis to update. The value 0 identifies the rows, and 1 identifies the columns.

inplace : boolean, default None

Whether to return a new %(klass)s instance.

.. warning::

``inplace=None`` currently falls back to to True, but in a future version, will default to False. Use inplace=True explicitly rather than relying on the default.

Returns

-----

renamed : %(klass)s or None

An object of same type as caller if inplace=False, None otherwise.

See Also

-----

DataFrame.rename\_axis : Alter the name of the index or columns.

Examples

-----

**\*\*Series\*\***

```
>>> s = pd.Series([1, 2, 3])
```

```
>>> s
```

```
0 1
```

```
1 2
```

```
2 3
```

```
dtype: int64
```

```
>>> s.set_axis(['a', 'b', 'c'], axis=0, inplace=False)
```

```
a 1
```

```
b 2
```

```
c 3
```

```
dtype: int64
```

The original object is not modified.

```
>>> s
```

```
0 1
```

```
1 2
```

```
2 3
```

```
dtype: int64
```

**\*\*DataFrame\*\***

```
>>> df = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
```

Change the row labels.

```
>>> df.set_axis(['a', 'b', 'c'], axis='index', inplace=False)
```

```
 A B
```

```
a 1 4
```

```
b 2 5
```

```
c 3 6
```

Change the column labels.

```
>>> df.set_axis(['I', 'II'], axis='columns', inplace=False)
```

```
 I II
```

```
0 1 4
```

```
1 2 5
```

```
2 3 6
```

Now, update the labels inplace.

```
>>> df.set_axis(['i', 'ii'], axis='columns', inplace=True)
```

```
>>> df
 i ii
0 1 4
1 2 5
2 3 6
```

```
slice_shift(self, periods=1, axis=0)
```

Equivalent to `shift` without copying data. The shifted data will not include the dropped periods and the shifted axis will be smaller than the original.

Parameters

-----

periods : int

Number of periods to move, can be positive or negative

Returns

-----

shifted : same type as caller

Notes

-----

While the `slice\_shift` is faster than `shift`, you may pay for it later during alignment.

```
squeeze(self, axis=None)
```

Squeeze 1 dimensional axis objects into scalars.

Series or DataFrames with a single element are squeezed to a scalar. DataFrames with a single column or a single row are squeezed to a Series. Otherwise the object is unchanged.

This method is most useful when you don't know if your object is a Series or DataFrame, but you do know it has just a single column. In that case you can safely call `squeeze` to ensure you have a Series.

Parameters

-----

axis : {0 or 'index', 1 or 'columns', None}, default None

A specific axis to squeeze. By default, all length-1 axes are squeezed.

.. versionadded:: 0.20.0

Returns

-----

DataFrame, Series, or scalar

The projection after squeezing `axis` or all the axes.

See Also

-----

Series.iloc : Integer-location based indexing for selecting scalars.

DataFrame.iloc : Integer-location based indexing for selecting Series.

Series.to\_frame : Inverse of DataFrame.squeeze for a single-column DataFrame.

Examples

-----

```
>>> primes = pd.Series([2, 3, 5, 7])
```

Slicing might produce a Series with a single value:

```
>>> even_primes = primes[primes % 2 == 0]
```

```
>>> even_primes
```

```
0 2
```

```
dtype: int64
```

```
>>> even_primes.squeeze()
```

```
2
```

Squeezing objects with more than one value in every axis does nothing:



```
>>> odd_primes = primes[primes % 2 == 1]
>>> odd_primes
1 3
2 5
3 7
dtype: int64
```

```
>>> odd_primes.squeeze()
1 3
2 5
3 7
dtype: int64
```

Squeezing is even more effective when used with DataFrames.

```
>>> df = pd.DataFrame([[1, 2], [3, 4]], columns=['a', 'b'])
>>> df
 a b
0 1 2
1 3 4
```

Slicing a single column will produce a DataFrame with the columns having only one value:

```
>>> df_a = df[['a']]
>>> df_a
 a
0 1
1 3
```

So the columns can be squeezed down, resulting in a Series:

```
>>> df_a.squeeze('columns')
0 1
1 3
Name: a, dtype: int64
```

Slicing a single row from a single column will produce a single scalar DataFrame:

```
>>> df_0a = df.loc[df.index < 1, ['a']]
>>> df_0a
 a
0 1
```

Squeezing the rows produces a single scalar Series:

```
>>> df_0a.squeeze('rows')
a 1
Name: 0, dtype: int64
```

Squeezing all axes will project directly into a scalar:

```
>>> df_0a.squeeze()
1
```

```
swapaxes(self, axis1, axis2, copy=True)
 Interchange axes and swap values axes appropriately.
```

```
Returns

y : same as input
```

```
tail(self, n=5)
 Return the last `n` rows.
```

This function returns last `n` rows from the object based on position. It is useful for quickly verifying data, for example, after sorting or appending rows.

Parameters

```

n : int, default 5
 Number of rows to select.
```

#### Returns

```

type of caller
 The last `n` rows of the caller object.
```

#### See Also

```

DataFrame.head : The first `n` rows of the caller object.
```

#### Examples

```

>>> df = pd.DataFrame({'animal':['alligator', 'bee', 'falcon', 'lion',
... 'monkey', 'parrot', 'shark', 'whale', 'zebra']})
>>> df
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
5 parrot
6 shark
7 whale
8 zebra
```

Viewing the last 5 lines

```
>>> df.tail()
 animal
4 monkey
5 parrot
6 shark
7 whale
8 zebra
```

Viewing the last `n` lines (three in this case)

```
>>> df.tail(3)
 animal
6 shark
7 whale
8 zebra
```

```
take(self, indices, axis=0, convert=None, is_copy=True, **kwargs)
 Return the elements in the given *positional* indices along an axis.
```

This means that we are not indexing according to actual values in the index attribute of the object. We are indexing according to the actual position of the element in the object.

#### Parameters

```

indices : array-like
 An array of ints indicating which positions to take.
axis : {0 or 'index', 1 or 'columns', None}, default 0
 The axis on which to select elements. ``0`` means that we are
 selecting rows, ``1`` means that we are selecting columns.
convert : bool, default True
 Whether to convert negative indices into positive ones.
 For example, ``-1`` would map to the ``len(axis) - 1``.
 The conversions are similar to the behavior of indexing a
 regular Python list.
```

```
.. deprecated:: 0.21.0
 In the future, negative indices will always be converted.
```

```
is_copy : bool, default True
 Whether to return a copy of the original object or not.
```

**\*\*kwargs**  
For compatibility with :meth:`numpy.take`. Has no effect on the output.

**Returns**

-----

taken : same type as caller

An array-like containing the elements taken from the object.

**See Also**

-----

DataFrame.loc : Select a subset of a DataFrame by labels.

DataFrame.iloc : Select a subset of a DataFrame by positions.

numpy.take : Take elements from an array along an axis.

**Examples**

-----

```
>>> df = pd.DataFrame([('falcon', 'bird', 389.0),
... ('parrot', 'bird', 24.0),
... ('lion', 'mammal', 80.5),
... ('monkey', 'mammal', np.nan)],
... columns=['name', 'class', 'max_speed'],
... index=[0, 2, 3, 1])
>>> df
```

	name	class	max_speed
0	falcon	bird	389.0
2	parrot	bird	24.0
3	lion	mammal	80.5
1	monkey	mammal	NaN

Take elements at positions 0 and 3 along the axis 0 (default).

Note how the actual indices selected (0 and 1) do not correspond to our selected indices 0 and 3. That's because we are selecting the 0th and 3rd rows, not rows whose indices equal 0 and 3.

```
>>> df.take([0, 3])
 name class max_speed
0 falcon bird 389.0
1 monkey mammal NaN
```

Take elements at indices 1 and 2 along the axis 1 (column selection).

```
>>> df.take([1, 2], axis=1)
 class max_speed
0 bird 389.0
2 bird 24.0
3 mammal 80.5
1 mammal NaN
```

We may take elements using negative integers for positive indices, starting from the end of the object, just like with Python lists.

```
>>> df.take([-1, -2])
 name class max_speed
1 monkey mammal NaN
3 lion mammal 80.5
```

to\_clipboard(self, excel=True, sep=None, \*\*kwargs)  
Copy object to the system clipboard.

Write a text representation of object to the system clipboard.  
This can be pasted into Excel, for example.

**Parameters**

-----

excel : bool, default True

- True, use the provided separator, writing in a csv format for allowing easy pasting into excel.
- False, write a string representation of the object to the clipboard.

```
sep : str, default ``\t``
 Field delimiter.
**kwargs
 These parameters will be passed to DataFrame.to_csv.
```

See Also

-----  
DataFrame.to\_csv : Write a DataFrame to a comma-separated values  
 (csv) file.  
read\_clipboard : Read text from clipboard and pass to read\_table.

Notes

-----  
Requirements for your platform.

- Linux : `xclip`, or `xsel` (with `gtk` or `PyQt4` modules)
- Windows : none
- OS X : none

Examples

-----  
Copy the contents of a DataFrame to the clipboard.

```
>>> df = pd.DataFrame([[1, 2, 3], [4, 5, 6]], columns=['A', 'B', 'C'])
>>> df.to_clipboard(sep=',')
... # Wrote the following to the system clipboard:
... # ,A,B,C
... # 0,1,2,3
... # 1,4,5,6
```

We can omit the the index by passing the keyword `index` and setting it to false.

```
>>> df.to_clipboard(sep=',', index=False)
... # Wrote the following to the system clipboard:
... # A,B,C
... # 1,2,3
... # 4,5,6
```

to\_dense(self)

Return dense representation of NDFrame (as opposed to sparse).

to\_excel(self, excel\_writer, sheet\_name='Sheet1', na\_rep='', float\_format=None, columns=None, header=True, index=True, index\_label=None, startrow=0, startcol=0, engine=None, merge\_cells=True, encoding=None, inf\_rep='inf', verbose=True, freeze\_panes=None)

Write object to an Excel sheet.

To write a single object to an Excel .xlsx file it is only necessary to specify a target file name. To write to multiple sheets it is necessary to create an `ExcelWriter` object with a target file name, and specify a sheet in the file to write to.

Multiple sheets may be written to by specifying unique `sheet\_name`. With all data written to the file it is necessary to save the changes. Note that creating an `ExcelWriter` object with a file name that already exists will result in the contents of the existing file being erased.

Parameters

-----  
excel\_writer : str or ExcelWriter object  
 File path or existing ExcelWriter.  
sheet\_name : str, default 'Sheet1'  
 Name of sheet which will contain DataFrame.  
na\_rep : str, default ''  
 Missing data representation.  
float\_format : str, optional  
 Format string for floating point numbers. For example  
 ``float\_format="%.2f"`` will format 0.1234 to 0.12.  
columns : sequence or list of str, optional  
 Columns to write.  
header : bool or list of str, default True  
 Write out the column names. If a list of string is given it is

assumed to be aliases for the column names.

`index` : bool, default True  
Write row names (`index`).

`index_label` : str or sequence, optional  
Column label for index column(s) if desired. If not specified, and ``header`` and ``index`` are True, then the index names are used. A sequence should be given if the DataFrame uses MultiIndex.

`startrow` : int, default 0  
Upper left cell row to dump data frame.

`startcol` : int, default 0  
Upper left cell column to dump data frame.

`engine` : str, optional  
Write engine to use, 'openpyxl' or 'xlsxwriter'. You can also set this via the options ``io.excel.xlsx.writer``, ``io.excel.xls.writer``, and ``io.excel.xlsm.writer``.

`merge_cells` : bool, default True  
Write MultiIndex and Hierarchical Rows as merged cells.

`encoding` : str, optional  
Encoding of the resulting excel file. Only necessary for xlwt, other writers support unicode natively.

`inf_rep` : str, default 'inf'  
Representation for infinity (there is no native representation for infinity in Excel).

`verbose` : bool, default True  
Display more information in the error logs.

`freeze_panes` : tuple of int (length 2), optional  
Specifies the one-based bottommost row and rightmost column that is to be frozen.

.. versionadded:: 0.20.0.

#### See Also

-----

`to_csv` : Write DataFrame to a comma-separated values (csv) file.

`ExcelWriter` : Class for writing DataFrame objects into excel sheets.

`read_excel` : Read an Excel file into a pandas DataFrame.

`read_csv` : Read a comma-separated values (csv) file into DataFrame.

#### Notes

-----

For compatibility with `:meth:`~DataFrame.to_csv``, `to_excel` serializes lists and dicts to strings before writing.

Once a workbook has been saved it is not possible write further data without rewriting the whole workbook.

#### Examples

-----

Create, write to and save a workbook:

```
>>> df1 = pd.DataFrame([['a', 'b'], ['c', 'd']],
... index=['row 1', 'row 2'],
... columns=['col 1', 'col 2'])
>>> df1.to_excel("output.xlsx") # doctest: +SKIP
```

To specify the sheet name:

```
>>> df1.to_excel("output.xlsx",
... sheet_name='Sheet_name_1') # doctest: +SKIP
```

If you wish to write to more than one sheet in the workbook, it is necessary to specify an ExcelWriter object:

```
>>> df2 = df1.copy()
>>> with pd.ExcelWriter('output.xlsx') as writer: # doctest: +SKIP
... df1.to_excel(writer, sheet_name='Sheet_name_1')
... df2.to_excel(writer, sheet_name='Sheet_name_2')
```

To set the library that is used to write the Excel file, you can pass the ``engine`` keyword (the default engine is automatically chosen depending on the file extension):

```
>>> df1.to_excel('output1.xlsx', engine='xlsxwriter') # doctest: +SKIP
```

```
to_hdf(self, path_or_buf, key, **kwargs)
```

Write the contained data to an HDF5 file using HDFStore.

Hierarchical Data Format (HDF) is self-describing, allowing an application to interpret the structure and contents of a file with no outside information. One HDF file can hold a mix of related objects which can be accessed as a group or as individual objects.

In order to add another DataFrame or Series to an existing HDF file please use append mode and a different a key.

For more information see the :ref:`user guide <io.hdf5>`.

#### Parameters

-----

`path_or_buf` : str or pandas.HDFStore

File path or HDFStore object.

`key` : str

Identifier for the group in the store.

`mode` : {'a', 'w', 'r+'}, default 'a'

Mode to open file:

- 'w': write, a new file is created (an existing file with the same name would be deleted).
- 'a': append, an existing file is opened for reading and writing, and if the file does not exist it is created.
- 'r+': similar to 'a', but the file must already exist.

`format` : {'fixed', 'table'}, default 'fixed'

Possible values:

- 'fixed': Fixed format. Fast writing/reading. Not-appendable, nor searchable.
- 'table': Table format. Write as a PyTables Table structure which may perform worse but allow more flexible operations like searching / selecting subsets of the data.

`append` : bool, default False

For Table formats, append the input data to the existing.

`data_columns` : list of columns or True, optional

List of columns to create as indexed data columns for on-disk queries, or True to use all columns. By default only the axes of the object are indexed. See :ref:`io.hdf5-query-data-columns`. Applicable only to format='table'.

`complevel` : {0-9}, optional

Specifies a compression level for data.

A value of 0 disables compression.

`complib` : {'zlib', 'lzo', 'bzip2', 'blosc'}, default 'zlib'

Specifies the compression library to be used.

As of v0.20.2 these additional compressors for Blosc are supported (default if no compressor specified: 'blosc:blosclz'):

{'blosc:blosclz', 'blosc:lz4', 'blosc:lz4hc', 'blosc:snappy', 'blosc:zlib', 'blosc:zstd'}.

Specifying a compression library which is not available issues a ValueError.

`fletcher32` : bool, default False

If applying compression use the fletcher32 checksum.

`dropna` : bool, default False

If true, ALL nan rows will not be written to store.

`errors` : str, default 'strict'

Specifies how encoding and decoding errors are to be handled.

See the errors argument for :func:`open` for a full list of options.

#### See Also

-----

`DataFrame.read_hdf` : Read from HDF file.

`DataFrame.to_parquet` : Write a DataFrame to the binary parquet format.

`DataFrame.to_sql` : Write to a sql table.

`DataFrame.to_feather` : Write out feather-format for DataFrames.

`DataFrame.to_csv` : Write out to a csv file.

## Examples

-----

```
>>> df = pd.DataFrame({'A': [1, 2, 3], 'B': [4, 5, 6]},
... index=['a', 'b', 'c'])
>>> df.to_hdf('data.h5', key='df', mode='w')
```

We can add another object to the same file:

```
>>> s = pd.Series([1, 2, 3, 4])
>>> s.to_hdf('data.h5', key='s')
```

Reading from HDF file:

```
>>> pd.read_hdf('data.h5', 'df')
A B
a 1 4
b 2 5
c 3 6
>>> pd.read_hdf('data.h5', 's')
0 1
1 2
2 3
3 4
dtype: int64
```

Deleting file with data:

```
>>> import os
>>> os.remove('data.h5')
```

```
to_json(self, path_or_buf=None, orient=None, date_format=None, double_precision=10,
force_ascii=True, date_unit='ms', default_handler=None, lines=False, compression='infer',
index=True)
```

Convert the object to a JSON string.

Note NaN's and None will be converted to null and datetime objects will be converted to UNIX timestamps.

## Parameters

-----

**path\_or\_buf** : string or file handle, optional  
File path or object. If not specified, the result is returned as a string.

**orient** : string  
Indication of expected JSON string format.

### \* Series

- default is 'index'
- allowed values are: {'split', 'records', 'index', 'table'}

### \* DataFrame

- default is 'columns'
- allowed values are:  
{ 'split', 'records', 'index', 'columns', 'values', 'table' }

### \* The format of the JSON string

- 'split' : dict like {'index' -> [index],  
                      'columns' -> [columns], 'data' -> [values]}
- 'records' : list like  
            [{column -> value}, ... , {column -> value}]
- 'index' : dict like {index -> {column -> value}}
- 'columns' : dict like {column -> {index -> value}}
- 'values' : just the values array
- 'table' : dict like {'schema': {schema}, 'data': {data}}  
describing the data, and the data component is  
like ``orient='records'``.

.. versionchanged:: 0.20.0

```

date_format : {None, 'epoch', 'iso'}
 Type of date conversion. 'epoch' = epoch milliseconds,
 'iso' = ISO8601. The default depends on the `orient`. For
 ``orient='table'``, the default is 'iso'. For all other orients,
 the default is 'epoch'.
double_precision : int, default 10
 The number of decimal places to use when encoding
 floating point values.
force_ascii : bool, default True
 Force encoded string to be ASCII.
date_unit : string, default 'ms' (milliseconds)
 The time unit to encode to, governs timestamp and ISO8601
 precision. One of 's', 'ms', 'us', 'ns' for second, millisecond,
 microsecond, and nanosecond respectively.
default_handler : callable, default None
 Handler to call if object cannot otherwise be converted to a
 suitable format for JSON. Should receive a single argument which is
 the object to convert and return a serialisable object.
lines : bool, default False
 If 'orient' is 'records' write out line delimited json format. Will
 throw ValueError if incorrect 'orient' since others are not list
 like.

 .. versionadded:: 0.19.0

```

```

compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}

```

A string representing the compression to use in the output file, only used when the first argument is a filename. By default, the compression is inferred from the filename.

```

 .. versionadded:: 0.21.0
 .. versionchanged:: 0.24.0
 'infer' option added and set to default
index : bool, default True
 Whether to include the index values in the JSON string. Not
 including the index (``index=False``) is only supported when
 orient is 'split' or 'table'.

 .. versionadded:: 0.23.0

```

See Also

-----

read\_json

Examples

-----

```

>>> df = pd.DataFrame([['a', 'b'], ['c', 'd']],
... index=['row 1', 'row 2'],
... columns=['col 1', 'col 2'])
>>> df.to_json(orient='split')
'{"columns":["col 1","col 2"],
 "index":["row 1","row 2"],
 "data":[["a","b"],["c","d"]}]'

```

Encoding/decoding a Dataframe using ``'records'`` formatted JSON.  
Note that index labels are not preserved with this encoding.

```

>>> df.to_json(orient='records')
'[{"col 1":"a","col 2":"b"}, {"col 1":"c","col 2":"d"}]'

```

Encoding/decoding a Dataframe using ``'index'`` formatted JSON:

```

>>> df.to_json(orient='index')
'{"row 1":{"col 1":"a","col 2":"b"},"row 2":{"col 1":"c","col 2":"d"}}'

```

Encoding/decoding a Dataframe using ``'columns'`` formatted JSON:

```

>>> df.to_json(orient='columns')
'{"col 1":{"row 1":"a","row 2":"c"},"col 2":{"row 1":"b","row 2":"d"}}'

```



Encoding/decoding a Dataframe using ``'values'`` formatted JSON:

```
>>> df.to_json(orient='values')
'[[{"a","b"}, {"c","d"}]'
```

Encoding with Table Schema

```
>>> df.to_json(orient='table')
'{"schema": {"fields": [{"name": "index", "type": "string"},
 {"name": "col 1", "type": "string"},
 {"name": "col 2", "type": "string"}],
 "primaryKey": "index",
 "pandas_version": "0.20.0"},
 "data": [{"index": "row 1", "col 1": "a", "col 2": "b"},
 {"index": "row 2", "col 1": "c", "col 2": "d"}]}'
```

`to_latex(self, buf=None, columns=None, col_space=None, header=True, index=True, na_rep='NaN', formatters=None, float_format=None, sparsify=None, index_names=True, bold_rows=False, column_format=None, longtable=None, escape=None, encoding=None, decimal='.', multicolumn=None, multicolumn_format=None, multirow=None)`

Render an object to a LaTeX tabular environment table.

Render an object to a tabular environment table. You can splice this into a LaTeX document. Requires `\usepackage{booktabs}`.

.. versionchanged:: 0.20.2  
Added to Series

Parameters

-----

`buf` : file descriptor or None

Buffer to write to. If None, the output is returned as a string.

`columns` : list of label, optional

The subset of columns to write. Writes all columns by default.

`col_space` : int, optional

The minimum width of each column.

`header` : bool or list of str, default True

Write out the column names. If a list of strings is given, it is assumed to be aliases for the column names.

`index` : bool, default True

Write row names (index).

`na_rep` : str, default 'NaN'

Missing data representation.

`formatters` : list of functions or dict of {str: function}, optional

Formatter functions to apply to columns' elements by position or name. The result of each function must be a unicode string.

List must be of length equal to the number of columns.

`float_format` : str, optional

Format string for floating point numbers.

`sparsify` : bool, optional

Set to False for a DataFrame with a hierarchical index to print every multiindex key at each row. By default, the value will be read from the config module.

`index_names` : bool, default True

Prints the names of the indexes.

`bold_rows` : bool, default False

Make the row labels bold in the output.

`column_format` : str, optional

The columns format as specified in `LaTeX table format` <<https://en.wikibooks.org/wiki/LaTeX/Tables>>`\_\_ e.g. 'rcl' for 3 columns. By default, 'l' will be used for all columns except columns of numbers, which default to 'r'.

`longtable` : bool, optional

By default, the value will be read from the pandas config module. Use a longtable environment instead of tabular. Requires adding a `\usepackage{longtable}` to your LaTeX preamble.

`escape` : bool, optional

By default, the value will be read from the pandas config module. When set to False prevents from escaping latex special characters in column names.

`encoding` : str, optional

A string representing the encoding to use in the output file, defaults to 'ascii' on Python 2 and 'utf-8' on Python 3.

decimal : str, default '.'  
Character recognized as decimal separator, e.g. ',' in Europe.  
.. versionadded:: 0.18.0

multicolumn : bool, default True  
Use \multicolumn to enhance MultiIndex columns.  
The default will be read from the config module.  
.. versionadded:: 0.20.0

multicolumn\_format : str, default 'l'  
The alignment for multicolumns, similar to 'column\_format'  
The default will be read from the config module.  
.. versionadded:: 0.20.0

multirow : bool, default False  
Use \multirow to enhance MultiIndex rows. Requires adding a \usepackage{multirow} to your LaTeX preamble. Will print centered labels (instead of top-aligned) across the contained rows, separating groups via clines. The default will be read from the pandas config module.  
.. versionadded:: 0.20.0

#### Returns

-----  
str or None  
If buf is None, returns the resulting LaTeX format as a string. Otherwise returns None.

#### See Also

-----  
DataFrame.to\_string : Render a DataFrame to a console-friendly tabular output.  
DataFrame.to\_html : Render a DataFrame as an HTML table.

#### Examples

-----  
>>> df = pd.DataFrame({'name': ['Raphael', 'Donatello'],  
... 'mask': ['red', 'purple'],  
... 'weapon': ['sai', 'bo staff']})  
>>> df.to\_latex(index=False) # doctest: +NORMALIZE\_WHITESPACE  
'\\begin{tabular}{lll}\\n\\toprule\\n name & mask & weapon  
\\\\\\n\\midrule\\n Raphael & red & sai \\\\\\n Donatello &  
purple & bo staff \\\\\\n\\bottomrule\\n\\end{tabular}\\n'

to\_msgpack(self, path\_or\_buf=None, encoding='utf-8', \*\*kwargs)  
Serialize object to input file path using msgpack format.

THIS IS AN EXPERIMENTAL LIBRARY and the storage format may not be stable until a future release.

#### Parameters

-----  
path : string File path, buffer-like, or None  
if None, return generated string  
append : bool whether to append to an existing msgpack (default is False)  
compress : type of compressor (zlib or blosc), default to None (no compression)

to\_pickle(self, path, compression='infer', protocol=4)  
Pickle (serialize) object to file.

#### Parameters

-----  
path : str  
File path where the pickled object will be stored.  
compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}, default 'infer'

A string representing the compression to use in the output file. By default, infers from the file extension in specified path.

.. versionadded:: 0.20.0  
protocol : int

Int which indicates which protocol should be used by the pickler, default HIGHEST\_PROTOCOL (see [1]\_ paragraph 12.1.2). The possible values for this parameter depend on the version of Python. For Python 2.x, possible values are 0, 1, 2. For Python >= 3.0, 3 is a valid value. For Python >= 3.4, 4 is a valid value. A negative value for the protocol parameter is equivalent to setting its value to HIGHEST\_PROTOCOL.

.. [1] <https://docs.python.org/3/library/pickle.html>  
.. versionadded:: 0.21.0

See Also

-----  
read\_pickle : Load pickled pandas object (or any object) from file.  
DataFrame.to\_hdf : Write DataFrame to an HDF5 file.  
DataFrame.to\_sql : Write DataFrame to a SQL database.  
DataFrame.to\_parquet : Write a DataFrame to the binary parquet format.

Examples

-----  
>>> original\_df = pd.DataFrame({"foo": range(5), "bar": range(5, 10)})  
>>> original\_df  
 foo bar  
0 0 5  
1 1 6  
2 2 7  
3 3 8  
4 4 9  
>>> original\_df.to\_pickle("./dummy.pkl")  
  
>>> unpickled\_df = pd.read\_pickle("./dummy.pkl")  
>>> unpickled\_df  
 foo bar  
0 0 5  
1 1 6  
2 2 7  
3 3 8  
4 4 9  
  
>>> import os  
>>> os.remove("./dummy.pkl")

to\_sql(self, name, con, schema=None, if\_exists='fail', index=True, index\_label=None, chunksize=None, dtype=None, method=None)

Write records stored in a DataFrame to a SQL database.

Databases supported by SQLAlchemy [1]\_ are supported. Tables can be newly created, appended to, or overwritten.

Parameters

-----  
name : string

Name of SQL table.

con : sqlalchemy.engine.Engine or sqlite3.Connection

Using SQLAlchemy makes it possible to use any DB supported by that library. Legacy support is provided for sqlite3.Connection objects.

schema : string, optional

Specify the schema (if database flavor supports this). If None, use default schema.

if\_exists : {'fail', 'replace', 'append'}, default 'fail'

How to behave if the table already exists.

\* fail: Raise a ValueError.

\* replace: Drop the table before inserting new values.

\* append: Insert new values to the existing table.

index : bool, default True

Write DataFrame index as a column. Uses `index\_label` as the column name in the table.

index\_label : string or sequence, default None

Column label for index column(s). If None is given (default) and `index` is True, then the index names are used.

A sequence should be given if the DataFrame uses MultiIndex.

chunksizе : int, optional  
 Rows will be written in batches of this size at a time. By default, all rows will be written at once.

dtype : dict, optional  
 Specifying the datatype for columns. The keys should be the column names and the values should be the SQLAlchemy types or strings for the sqlite3 legacy mode.

method : {None, 'multi', callable}, default None  
 Controls the SQL insertion clause used:

- \* None : Uses standard SQL ``INSERT`` clause (one per row).
- \* 'multi': Pass multiple values in a single ``INSERT`` clause.
- \* callable with signature ``(pd\_table, conn, keys, data\_iter)``.

Details and a sample callable implementation can be found in the section :ref:`insert method <io.sql.method>`.

.. versionadded:: 0.24.0

#### Raises

-----

#### ValueError

When the table already exists and `if\_exists` is 'fail' (the default).

#### See Also

-----

read\_sql : Read a DataFrame from a table.

#### Notes

-----

Timezone aware datetime columns will be written as ``Timestamp with timezone`` type with SQLAlchemy if supported by the database. Otherwise, the datetimes will be stored as timezone unaware timestamps local to the original timezone.

.. versionadded:: 0.24.0

#### References

-----

- .. [1] <http://docs.sqlalchemy.org>
- .. [2] <https://www.python.org/dev/peps/pep-0249/>

#### Examples

-----

Create an in-memory SQLite database.

```
>>> from sqlalchemy import create_engine
>>> engine = create_engine('sqlite://', echo=False)
```

Create a table from scratch with 3 rows.

```
>>> df = pd.DataFrame({'name' : ['User 1', 'User 2', 'User 3']})
>>> df
```

```
 name
0 User 1
1 User 2
2 User 3
```

```
>>> df.to_sql('users', con=engine)
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 1'), (1, 'User 2'), (2, 'User 3')]
```

```
>>> df1 = pd.DataFrame({'name' : ['User 4', 'User 5']})
>>> df1.to_sql('users', con=engine, if_exists='append')
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 1'), (1, 'User 2'), (2, 'User 3'),
 (0, 'User 4'), (1, 'User 5')]
```

Overwrite the table with just ``df1``.

```
>>> df1.to_sql('users', con=engine, if_exists='replace',
... index_label='id')
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 4'), (1, 'User 5')]
```

Specify the dtype (especially useful for integers with missing values). Notice that while pandas is forced to store the data as floating point, the database supports nullable integers. When fetching the data with Python, we get back integer scalars.

```
>>> df = pd.DataFrame({"A": [1, None, 2]})
>>> df
 A
0 1.0
1 NaN
2 2.0

>>> from sqlalchemy.types import Integer
>>> df.to_sql('integers', con=engine, index=False,
... dtype={"A": Integer()})

>>> engine.execute("SELECT * FROM integers").fetchall()
[(1,), (None,), (2,)]
```

`to_xarray(self)`

Return an xarray object from the pandas object.

Returns

-----

xarray.DataArray or xarray.Dataset

Data in the pandas structure converted to Dataset if the object is a DataFrame, or a DataArray if the object is a Series.

See Also

-----

DataFrame.to\_hdf : Write DataFrame to an HDF5 file.

DataFrame.to\_parquet : Write a DataFrame to the binary parquet format.

Notes

-----

See the `xarray docs <<http://xarray.pydata.org/en/stable/>>`\_\_

Examples

-----

```
>>> df = pd.DataFrame([('falcon', 'bird', 389.0, 2),
... ('parrot', 'bird', 24.0, 2),
... ('lion', 'mammal', 80.5, 4),
... ('monkey', 'mammal', np.nan, 4)],
... columns=['name', 'class', 'max_speed',
... 'num_legs'])
>>> df
 name class max_speed num_legs
0 falcon bird 389.0 2
1 parrot bird 24.0 2
2 lion mammal 80.5 4
3 monkey mammal NaN 4
```

```
>>> df.to_xarray()
<xarray.Dataset>
Dimensions: (index: 4)
Coordinates:
 * index (index) int64 0 1 2 3
Data variables:
 name (index) object 'falcon' 'parrot' 'lion' 'monkey'
 class (index) object 'bird' 'bird' 'mammal' 'mammal'
 max_speed (index) float64 389.0 24.0 80.5 nan
 num_legs (index) int64 2 2 4 4
```

```
>>> df['max_speed'].to_xarray()
<xarray.DataArray 'max_speed' (index: 4)>
array([389. , 24. , 80.5, nan])
```

```

Coordinates:
 * index (index) int64 0 1 2 3

>>> dates = pd.to_datetime(['2018-01-01', '2018-01-01',
... '2018-01-02', '2018-01-02'])
>>> df_multiindex = pd.DataFrame({'date': dates,
... 'animal': ['falcon', 'parrot', 'falcon',
... 'parrot'],
... 'speed': [350, 18, 361, 15]}).set_index(['date',
... 'animal'])
>>> df_multiindex

```

```

 speed
date animal
2018-01-01 falcon 350
 parrot 18
2018-01-02 falcon 361
 parrot 15

```

```

>>> df_multiindex.to_xarray()
<xarray.Dataset>
Dimensions: (animal: 2, date: 2)
Coordinates:
 * date (date) datetime64[ns] 2018-01-01 2018-01-02
 * animal (animal) object 'falcon' 'parrot'
Data variables:
 speed (date, animal) int64 350 18 361 15

```

```

truncate(self, before=None, after=None, axis=None, copy=True)
 Truncate a Series or DataFrame before and after some index value.

```

This is a useful shorthand for boolean indexing based on index values above or below certain thresholds.

#### Parameters

```

before : date, string, int
 Truncate all rows before this index value.
after : date, string, int
 Truncate all rows after this index value.
axis : {0 or 'index', 1 or 'columns'}, optional
 Axis to truncate. Truncates the index (rows) by default.
copy : boolean, default is True,
 Return a copy of the truncated section.

```

#### Returns

```

type of caller
 The truncated Series or DataFrame.

```

#### See Also

```

DataFrame.loc : Select a subset of a DataFrame by label.
DataFrame.iloc : Select a subset of a DataFrame by position.

```

#### Notes

```

If the index being truncated contains only datetime values,
`before` and `after` may be specified as strings instead of
Timestamps.

```

#### Examples

```

>>> df = pd.DataFrame({'A': ['a', 'b', 'c', 'd', 'e'],
... 'B': ['f', 'g', 'h', 'i', 'j'],
... 'C': ['k', 'l', 'm', 'n', 'o']},
... index=[1, 2, 3, 4, 5])
>>> df
 A B C
1 a f k
2 b g l
3 c h m
4 d i n

```

```
5 e j o
```

```
>>> df.truncate(before=2, after=4)
```

```
 A B C
2 b g l
3 c h m
4 d i n
```

The columns of a DataFrame can be truncated.

```
>>> df.truncate(before="A", after="B", axis="columns")
```

```
 A B
1 a f
2 b g
3 c h
4 d i
5 e j
```

For Series, only rows can be truncated.

```
>>> df['A'].truncate(before=2, after=4)
```

```
2 b
3 c
4 d
Name: A, dtype: object
```

The index values in ``truncate`` can be datetimes or string dates.

```
>>> dates = pd.date_range('2016-01-01', '2016-02-01', freq='s')
```

```
>>> df = pd.DataFrame(index=dates, data={'A': 1})
```

```
>>> df.tail()
```

```
 A
2016-01-31 23:59:56 1
2016-01-31 23:59:57 1
2016-01-31 23:59:58 1
2016-01-31 23:59:59 1
2016-02-01 00:00:00 1
```

```
>>> df.truncate(before=pd.Timestamp('2016-01-05'),
... after=pd.Timestamp('2016-01-10')).tail()
```

```
 A
2016-01-09 23:59:56 1
2016-01-09 23:59:57 1
2016-01-09 23:59:58 1
2016-01-09 23:59:59 1
2016-01-10 00:00:00 1
```

Because the index is a DatetimeIndex containing only dates, we can specify `before` and `after` as strings. They will be coerced to Timestamps before truncation.

```
>>> df.truncate('2016-01-05', '2016-01-10').tail()
```

```
 A
2016-01-09 23:59:56 1
2016-01-09 23:59:57 1
2016-01-09 23:59:58 1
2016-01-09 23:59:59 1
2016-01-10 00:00:00 1
```

Note that ``truncate`` assumes a 0 value for any unspecified time component (midnight). This differs from partial string slicing, which returns any partially matching dates.

```
>>> df.loc['2016-01-05':'2016-01-10', :].tail()
```

```
 A
2016-01-10 23:59:55 1
2016-01-10 23:59:56 1
2016-01-10 23:59:57 1
2016-01-10 23:59:58 1
2016-01-10 23:59:59 1
```

```

| tshift(self, periods=1, freq=None, axis=0)
| Shift the time index, using the index's frequency if available.
|
| Parameters
| -----
| periods : int
| Number of periods to move, can be positive or negative
| freq : DateOffset, timedelta, or time rule string, default None
| Increment to use from the tseries module or time rule (e.g. 'EOM')
| axis : int or basestring
| Corresponds to the axis that contains the Index
|
| Returns
| -----
| shifted : NDFrame
|
| Notes
| -----
| If freq is not specified then tries to use the freq or inferred_freq
| attributes of the index. If neither of those attributes exist, a
| ValueError is thrown
|
| tz_convert(self, tz, axis=0, level=None, copy=True)
| Convert tz-aware axis to target time zone.
|
| Parameters
| -----
| tz : string or pytz.timezone object
| axis : the axis to convert
| level : int, str, default None
| If axis is a MultiIndex, convert a specific level. Otherwise
| must be None
| copy : boolean, default True
| Also make a copy of the underlying data
|
| Returns
| -----
|
| Raises
| -----
| TypeError
| If the axis is tz-naive.
|
| tz_localize(self, tz, axis=0, level=None, copy=True, ambiguous='raise', nonexistent='
raise')
| Localize tz-naive index of a Series or DataFrame to target time zone.
|
| This operation localizes the Index. To localize the values in a
| timezone-naive Series, use :meth:`Series.dt.tz_localize`.
|
| Parameters
| -----
| tz : string or pytz.timezone object
| axis : the axis to localize
| level : int, str, default None
| If axis is a MultiIndex, localize a specific level. Otherwise
| must be None
| copy : boolean, default True
| Also make a copy of the underlying data
| ambiguous : 'infer', bool-ndarray, 'NaT', default 'raise'
| When clocks moved backward due to DST, ambiguous times may arise.
| For example in Central European Time (UTC+01), when going from
| 03:00 DST to 02:00 non-DST, 02:30:00 local time occurs both at
| 00:30:00 UTC and at 01:30:00 UTC. In such a situation, the
| `ambiguous` parameter dictates how ambiguous times should be
| handled.
|
| - 'infer' will attempt to infer fall dst-transition hours based on
| order
| - bool-ndarray where True signifies a DST time, False designates
| a non-DST time (note that this flag is only applicable for
| ambiguous times)

```



- 'NaT' will return NaT where there are ambiguous times
- 'raise' will raise an AmbiguousTimeError if there are ambiguous times

nonexistent : str, default 'raise'

A nonexistent time does not exist in a particular timezone where clocks moved forward due to DST. Valid valuse are:

- 'shift\_forward' will shift the nonexistent time forward to the closest existing time
- 'shift\_backward' will shift the nonexistent time backward to the closest existing time
- 'NaT' will return NaT where there are nonexistent times
- timedelta objects will shift nonexistent times by the timedelta
- 'raise' will raise an NonExistentTimeError if there are nonexistent times

.. versionadded:: 0.24.0

#### Returns

-----

Series or DataFrame

Same type as the input.

#### Raises

-----

TypeError

If the TimeSeries is tz-aware and tz is not None.

#### Examples

-----

Localize local times:

```
>>> s = pd.Series([1],
... index=pd.DatetimeIndex(['2018-09-15 01:30:00']))
>>> s.tz_localize('CET')
2018-09-15 01:30:00+02:00 1
dtype: int64
```

Be careful with DST changes. When there is sequential data, pandas can infer the DST time:

```
>>> s = pd.Series(range(7), index=pd.DatetimeIndex([
... '2018-10-28 01:30:00',
... '2018-10-28 02:00:00',
... '2018-10-28 02:30:00',
... '2018-10-28 02:00:00',
... '2018-10-28 02:30:00',
... '2018-10-28 03:00:00',
... '2018-10-28 03:30:00']))
>>> s.tz_localize('CET', ambiguous='infer')
2018-10-28 01:30:00+02:00 0
2018-10-28 02:00:00+02:00 1
2018-10-28 02:30:00+02:00 2
2018-10-28 02:00:00+01:00 3
2018-10-28 02:30:00+01:00 4
2018-10-28 03:00:00+01:00 5
2018-10-28 03:30:00+01:00 6
dtype: int64
```

In some cases, inferring the DST is impossible. In such cases, you can pass an ndarray to the ambiguous parameter to set the DST explicitly

```
>>> s = pd.Series(range(3), index=pd.DatetimeIndex([
... '2018-10-28 01:20:00',
... '2018-10-28 02:36:00',
... '2018-10-28 03:46:00']))
>>> s.tz_localize('CET', ambiguous=np.array([True, True, False]))
2018-10-28 01:20:00+02:00 0
2018-10-28 02:36:00+02:00 1
2018-10-28 03:46:00+01:00 2
dtype: int64
```

If the DST transition causes nonexistent times, you can shift these dates forward or backwards with a `timedelta` object or `'shift_forward'` or `'shift_backwards'`.

```
>>> s = pd.Series(range(2), index=pd.DatetimeIndex([
... '2015-03-29 02:30:00',
... '2015-03-29 03:30:00']))
>>> s.tz_localize('Europe/Warsaw', nonexistent='shift_forward')
2015-03-29 03:00:00+02:00 0
2015-03-29 03:30:00+02:00 1
dtype: int64
>>> s.tz_localize('Europe/Warsaw', nonexistent='shift_backward')
2015-03-29 01:59:59.999999999+01:00 0
2015-03-29 03:30:00+02:00 1
dtype: int64
>>> s.tz_localize('Europe/Warsaw', nonexistent=pd.Timedelta('1H'))
2015-03-29 03:30:00+02:00 0
2015-03-29 03:30:00+02:00 1
dtype: int64
```

`where(self, cond, other=nan, inplace=False, axis=None, level=None, errors='raise', try_cast=False, raise_on_error=None)`

Replace values where the condition is False.

Parameters

`cond` : boolean NDFrame, array-like, or callable

Where `'cond'` is True, keep the original value. Where False, replace with corresponding value from `'other'`. If `'cond'` is callable, it is computed on the NDFrame and should return boolean NDFrame or array. The callable must not change input NDFrame (though pandas doesn't check it).

.. versionadded:: 0.18.1

A callable can be used as `cond`.

`other` : scalar, NDFrame, or callable

Entries where `'cond'` is False are replaced with corresponding value from `'other'`. If `other` is callable, it is computed on the NDFrame and should return scalar or NDFrame. The callable must not change input NDFrame (though pandas doesn't check it).

.. versionadded:: 0.18.1

A callable can be used as `other`.

`inplace` : boolean, default False

Whether to perform the operation in place on the data.

`axis` : int, default None

Alignment axis if needed.

`level` : int, default None

Alignment level if needed.

`errors` : str, {'raise', 'ignore'}, default 'raise'

Note that currently this parameter won't affect the results and will always coerce to a suitable dtype.

- 'raise' : allow exceptions to be raised.

- 'ignore' : suppress exceptions. On error return original object.

`try_cast` : boolean, default False

Try to cast the result back to the input type (if possible).

`raise_on_error` : boolean, default True

Whether to raise on invalid data types (e.g. trying to where on strings).

.. deprecated:: 0.21.0

Use `'errors'`.

Returns

`wh` : same type as caller

See Also

-----  
:func:`DataFrame.mask` : Return an object of same shape as  
self.

Notes

-----  
The where method is an application of the if-then idiom. For each element in the calling DataFrame, if ``cond`` is ``True`` the element is used; otherwise the corresponding element from the DataFrame ``other`` is used.

The signature for :func:`DataFrame.where` differs from :func:`numpy.where`. Roughly ``df1.where(m, df2)`` is equivalent to ``np.where(m, df1, df2)``.

For further details and examples see the ``where`` documentation in :ref:`indexing <indexing.where\_mask>`.

Examples

-----  
>>> s = pd.Series(range(5))  
>>> s.where(s > 0)  
0 NaN  
1 1.0  
2 2.0  
3 3.0  
4 4.0  
dtype: float64  
  
>>> s.mask(s > 0)  
0 0.0  
1 NaN  
2 NaN  
3 NaN  
4 NaN  
dtype: float64  
  
>>> s.where(s > 1, 10)  
0 10  
1 10  
2 2  
3 3  
4 4  
dtype: int64  
  
>>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A', 'B'])  
>>> m = df % 3 == 0  
>>> df.where(m, -df)  
 A B  
0 0 -1  
1 -2 3  
2 -4 -5  
3 6 -7  
4 -8 9  
>>> df.where(m, -df) == np.where(m, df, -df)  
 A B  
0 True True  
1 True True  
2 True True  
3 True True  
4 True True  
>>> df.where(m, -df) == df.mask(~m, -df)  
 A B  
0 True True  
1 True True  
2 True True  
3 True True  
4 True True

xs(self, key, axis=0, level=None, drop\_level=True)

Return cross-section from the Series/DataFrame.

This method takes a `key` argument to select data at a particular level of a MultiIndex.

#### Parameters

-----  
key : label or tuple of label

Label contained in the index, or partially in a MultiIndex.

axis : {0 or 'index', 1 or 'columns'}, default 0

Axis to retrieve cross-section on.

level : object, defaults to first n levels (n=1 or len(key))

In case of a key partially contained in a MultiIndex, indicate which levels are used. Levels can be referred by label or position.

drop\_level : bool, default True

If False, returns object with same levels as self.

#### Returns

-----  
Series or DataFrame

Cross-section from the original Series or DataFrame corresponding to the selected index levels.

#### See Also

-----  
DataFrame.loc : Access a group of rows and columns by label(s) or a boolean array.

DataFrame.iloc : Purely integer-location based indexing for selection by position.

#### Notes

-----  
`xs` can not be used to set values.

MultiIndex Slicers is a generic way to get/set values on any level or levels.

It is a superset of `xs` functionality, see :ref:`MultiIndex Slicers <advanced.mi\_slicers>`.

#### Examples

-----  
>>> d = {'num\_legs': [4, 4, 2, 2],  
... 'num\_wings': [0, 0, 2, 2],  
... 'class': ['mammal', 'mammal', 'mammal', 'bird'],  
... 'animal': ['cat', 'dog', 'bat', 'penguin'],  
... 'locomotion': ['walks', 'walks', 'flies', 'walks']}  
>>> df = pd.DataFrame(data=d)  
>>> df = df.set\_index(['class', 'animal', 'locomotion'])  
>>> df

			num_legs	num_wings
class	animal	locomotion		
mammal	cat	walks	4	0
	dog	walks	4	0
	bat	flies	2	2
bird	penguin	walks	2	2

Get values at specified index

```
>>> df.xs('mammal')
 num_legs num_wings
animal locomotion
cat walks 4 0
dog walks 4 0
bat flies 2 2
```

Get values at several indexes

```
>>> df.xs(('mammal', 'dog'))
 num_legs num_wings
locomotion
walks 4 0
```

Get values at specified index and level

```
>>> df.xs('cat', level=1)
 num_legs num_wings
class locomotion
mammal walks 4 0
```

Get values at several indexes and levels

```
>>> df.xs(('bird', 'walks'),
... level=[0, 'locomotion'])
 num_legs num_wings
animal
penguin 2 2
```

Get values at specified column and axis

```
>>> df.xs('num_wings', axis=1)
class animal locomotion
mammal cat walks 0
 dog walks 0
 bat flies 2
bird penguin walks 2
Name: num_wings, dtype: int64
```

---

Data descriptors inherited from pandas.core.generic.NDFrame:

at

Access a single value for a row/column label pair.

Similar to ``loc``, in that both provide label-based lookups. Use ``at`` if you only need to get or set a single value in a DataFrame or Series.

Raises

-----

KeyError

When label does not exist in DataFrame

See Also

-----

DataFrame.iat : Access a single value for a row/column pair by integer position.

DataFrame.loc : Access a group of rows and columns by label(s).

Series.at : Access a single value using a label.

Examples

-----

```
>>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],
... index=[4, 5, 6], columns=['A', 'B', 'C'])
>>> df
 A B C
4 0 2 3
5 0 4 1
6 10 20 30
```

Get value at specified row/column pair

```
>>> df.at[4, 'B']
2
```

Set value at specified row/column pair

```
>>> df.at[4, 'B'] = 10
>>> df.at[4, 'B']
10
```

Get value within a Series

```
>>> df.loc[5].at['B']
4
```

## blocks

Internal property, property synonym for `as_blocks()`.

.. deprecated:: 0.21.0

## iat

Access a single value for a row/column pair by integer position.

Similar to ```iloc```, in that both provide integer-based lookups. Use ```iat``` if you only need to get or set a single value in a DataFrame or Series.

Raises

-----

IndexError

When integer position is out of bounds

See Also

-----

DataFrame.at : Access a single value for a row/column label pair.

DataFrame.loc : Access a group of rows and columns by label(s).

DataFrame.iloc : Access a group of rows and columns by integer position(s).

Examples

-----

```
>>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],
... columns=['A', 'B', 'C'])
```

```
>>> df
 A B C
0 0 2 3
1 0 4 1
2 10 20 30
```

Get value at specified row/column pair

```
>>> df.iat[1, 2]
1
```

Set value at specified row/column pair

```
>>> df.iat[1, 2] = 10
>>> df.iat[1, 2]
10
```

Get value within a series

```
>>> df.loc[0].iat[1]
2
```

## iloc

Purely integer-location based indexing for selection by position.

```iloc[]``` is primarily integer position based (from ```0``` to ```length-1``` of the axis), but may also be used with a boolean array.

Allowed inputs are:

- An integer, e.g. ```5```.
- A list or array of integers, e.g. ```[4, 3, 0]```.
- A slice object with ints, e.g. ```1:7```.
- A boolean array.
- A ```callable``` function with one argument (the calling Series, DataFrame or Panel) and that returns valid output for indexing (one of the above). This is useful in method chains, when you don't have a reference to the calling object, but would like to base your selection on some value.

```iloc``` will raise ```IndexError``` if a requested indexer is out-of-bounds, except `*slice*` indexers which allow out-of-bounds indexing (this conforms with python/numpy `*slice*` semantics).

See more at ref: `Selection by Position <indexing.integer>`.

See Also

-----

DataFrame.iat : Fast integer location scalar accessor.

DataFrame.loc : Purely label-location based indexer for selection by label.

Series.iloc : Purely integer-location based indexing for  
selection by position.

Examples

-----

```
>>> mydict = [{'a': 1, 'b': 2, 'c': 3, 'd': 4},
... {'a': 100, 'b': 200, 'c': 300, 'd': 400},
... {'a': 1000, 'b': 2000, 'c': 3000, 'd': 4000 }]
>>> df = pd.DataFrame(mydict)
```

```
>>> df
```

	a	b	c	d
0	1	2	3	4
1	100	200	300	400
2	1000	2000	3000	4000

**\*\*Indexing just the rows\*\***

With a scalar integer.

```
>>> type(df.iloc[0])
<class 'pandas.core.series.Series'>
>>> df.iloc[0]
a 1
b 2
c 3
d 4
Name: 0, dtype: int64
```

With a list of integers.

```
>>> df.iloc[[0]]
 a b c d
0 1 2 3 4
>>> type(df.iloc[[0]])
<class 'pandas.core.frame.DataFrame'>
```

```
>>> df.iloc[[0, 1]]
 a b c d
0 1 2 3 4
1 100 200 300 400
```

With a `slice` object.

```
>>> df.iloc[:3]
 a b c d
0 1 2 3 4
1 100 200 300 400
2 1000 2000 3000 4000
```

With a boolean mask the same length as the index.

```
>>> df.iloc[[True, False, True]]
 a b c d
0 1 2 3 4
2 1000 2000 3000 4000
```

With a callable, useful in method chains. The `x` passed to the ``lambda`` is the DataFrame being sliced. This selects the rows whose index label even.

```
>>> df.iloc[lambda x: x.index % 2 == 0]
 a b c d
0 1 2 3 4
2 1000 2000 3000 4000
```

**\*\*Indexing both axes\*\***

You can mix the indexer types for the index and columns. Use ``:``` to select the entire axis.

With scalar integers.

```
>>> df.iloc[0, 1]
2
```

With lists of integers.

```
>>> df.iloc[[0, 2], [1, 3]]
 b d
0 2 4
2 2000 4000
```

With `slice` objects.

```
>>> df.iloc[1:3, 0:3]
 a b c
1 100 200 300
2 1000 2000 3000
```

With a boolean array whose length matches the columns.

```
>>> df.iloc[:, [True, False, True, False]]
 a c
0 1 3
1 100 300
2 1000 3000
```

With a callable function that expects the Series or DataFrame.

```
>>> df.iloc[:, lambda df: [0, 2]]
 a c
0 1 3
1 100 300
2 1000 3000
```

`is_copy`

Return the copy.

`ix`

A primarily label-location based indexer, with integer position fallback.

Warning: Starting in 0.20.0, the `.ix` indexer is deprecated, in favor of the more strict `.iloc` and `.loc` indexers.

```.ix[]``` supports mixed integer and label based access. It is primarily label based, but will fall back to integer positional access unless the corresponding axis is of integer type.

```.ix``` is the most general indexer and will support any of the inputs in ```.loc``` and ```.iloc```. ```.ix``` also supports floating point label schemes. ```.ix``` is exceptionally useful when dealing with mixed positional and label based hierarchical indexes.

However, when an axis is integer based, ONLY label based access and not positional access is supported. Thus, in such cases, it's usually better to be explicit and use ```.iloc``` or ```.loc```.

See more at :ref:`Advanced Indexing <advanced>`.

`loc`

Access a group of rows and columns by label(s) or a boolean array.

```.loc[]``` is primarily label based, but may also be used with a boolean array.

Allowed inputs are:

- A single label, e.g. ``5`` or ``'a'``, (note that ``5`` is interpreted as a *label* of the index, and **never** as an integer position along the index).
- A list or array of labels, e.g. ``['a', 'b', 'c']``.
- A slice object with labels, e.g. ``'a':'f'``.

.. warning:: Note that contrary to usual python slices, **both** the start and the stop are included

- A boolean array of the same length as the axis being sliced, e.g. ``[True, False, True]``.
- A ``callable`` function with one argument (the calling Series, DataFrame or Panel) and that returns valid output for indexing (one of the above)

See more at :ref:`Selection by Label <indexing.label>`

Raises

KeyError:

when any items are not found

See Also

DataFrame.at : Access a single value for a row/column label pair.

DataFrame.iloc : Access group of rows and columns by integer position(s).

DataFrame.xs : Returns a cross-section (row(s) or column(s)) from the Series/DataFrame.

Series.loc : Access group of values using labels.

Examples

Getting values

```
>>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
...                    index=['cobra', 'viper', 'sidewinder'],
...                    columns=['max_speed', 'shield'])
>>> df
```

	max_speed	shield
cobra	1	2
viper	4	5
sidewinder	7	8

Single label. Note this returns the row as a Series.

```
>>> df.loc['viper']
max_speed    4
shield       5
Name: viper, dtype: int64
```

List of labels. Note using ``[['']]`` returns a DataFrame.

```
>>> df.loc[['viper', 'sidewinder']]
          max_speed  shield
viper           4       5
sidewinder       7       8
```

Single label for row and column

```
>>> df.loc['cobra', 'shield']
2
```

Slice with labels for row and single label for column. As mentioned above, note that both the start and stop of the slice are included.

```
>>> df.loc['cobra':'viper', 'max_speed']
cobra    1
viper    4
Name: max_speed, dtype: int64
```

Boolean list with the same length as the row axis

```
>>> df.loc[[False, False, True]]
           max_speed  shield
sidewinder           7       8
```

Conditional that returns a boolean Series

```
>>> df.loc[df['shield'] > 6]
           max_speed  shield
sidewinder           7       8
```

Conditional that returns a boolean Series with column labels specified

```
>>> df.loc[df['shield'] > 6, ['max_speed']]
           max_speed
sidewinder           7
```

Callable that returns a boolean Series

```
>>> df.loc[lambda df: df['shield'] == 8]
           max_speed  shield
sidewinder           7       8
```

****Setting values****

Set value for all items matching the list of labels

```
>>> df.loc[['viper', 'sidewinder'], ['shield']] = 50
>>> df
           max_speed  shield
cobra              1        2
viper              4       50
sidewinder          7       50
```

Set value for an entire row

```
>>> df.loc['cobra'] = 10
>>> df
           max_speed  shield
cobra             10       10
viper              4       50
sidewinder          7       50
```

Set value for an entire column

```
>>> df.loc[:, 'max_speed'] = 30
>>> df
           max_speed  shield
cobra              30       10
viper              30       50
sidewinder          30       50
```

Set value for rows matching callable condition

```
>>> df.loc[df['shield'] > 35] = 0
>>> df
           max_speed  shield
cobra             30       10
viper              0        0
sidewinder          0        0
```

****Getting values on a DataFrame with an index that has integer labels****

Another example using integers for the index

```
>>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
...                    index=[7, 8, 9], columns=['max_speed', 'shield'])
>>> df
           max_speed  shield
7              1        2
8              4        5
9              7        8
```

Slice with integer labels for rows. As mentioned above, note that both the start and stop of the slice are included.

```
>>> df.loc[7:9]
   max_speed  shield
7          1        2
8          4        5
9          7        8
```

****Getting values with a MultiIndex****

A number of examples using a DataFrame with a MultiIndex

```
>>> tuples = [
...     ('cobra', 'mark i'), ('cobra', 'mark ii'),
...     ('sidewinder', 'mark i'), ('sidewinder', 'mark ii'),
...     ('viper', 'mark ii'), ('viper', 'mark iii')
... ]
>>> index = pd.MultiIndex.from_tuples(tuples)
>>> values = [[12, 2], [0, 4], [10, 20],
...           [1, 4], [7, 1], [16, 36]]
>>> df = pd.DataFrame(values, columns=['max_speed', 'shield'], index=index)
>>> df
```

		max_speed	shield
cobra	mark i	12	2
	mark ii	0	4
sidewinder	mark i	10	20
	mark ii	1	4
viper	mark ii	7	1
	mark iii	16	36

Single label. Note this returns a DataFrame with a single index.

```
>>> df.loc['cobra']
   max_speed  shield
mark i      12        2
mark ii      0        4
```

Single index tuple. Note this returns a Series.

```
>>> df.loc[('cobra', 'mark ii')]
max_speed    0
shield        4
Name: (cobra, mark ii), dtype: int64
```

Single label for row and column. Similar to passing in a tuple, this returns a Series.

```
>>> df.loc['cobra', 'mark i']
max_speed    12
shield        2
Name: (cobra, mark i), dtype: int64
```

Single tuple. Note using ``[[]]`` returns a DataFrame.

```
>>> df.loc[['cobra', 'mark ii']]
   max_speed  shield
cobra mark ii      0        4
```

Single tuple for the index with a single label for the column

```
>>> df.loc[('cobra', 'mark i'), 'shield']
2
```

Slice from index tuple to single label

```
>>> df.loc[('cobra', 'mark i'):'viper']
   max_speed  shield
cobra mark i      12        2
      mark ii      0        4
sidewinder mark i      10       20
      mark ii      1        4
```

viper	mark ii	7	1
	mark iii	16	36

Slice from index tuple to index tuple

```
>>> df.loc[('cobra', 'mark i'):(('viper', 'mark ii'))]
               max_speed  shield
cobra      mark i         12      2
           mark ii          0      4
sidewinder mark i         10     20
           mark ii          1      4
viper      mark ii          7      1
```

Data and other attributes inherited from pandas.core.generic.NDFrame:

timetuple = None

Methods inherited from pandas.core.base.PandasObject:

__sizeof__(self)
Generates the total memory usage for an object that returns
either a value or Series of values

Methods inherited from pandas.core.base.StringMixin:

__bytes__(self)
Return a string representation for a particular object.

Invoked by bytes(obj) in py3 only.
Yields a bytestring in both py2/py3.

__repr__(self)
Return a string representation for a particular object.

Yields Bytestring in Py2, Unicode String in py3.

__str__(self)
Return a string representation for a particular Object

Invoked by str(df) in both py2/py3.
Yields Bytestring in Py2, Unicode String in py3.

Methods inherited from pandas.core.accessor.DirNamesMixin:

__dir__(self)
Provide method name lookup and completion
Only provide 'public' methods

Index and Data Lists

We can create a Series from Python lists (also from NumPy arrays)

In [14]:

```
myindex = ['USA', 'Canada', 'Mexico']
```

In [15]:

```
mydata = [1776, 1867, 1821]
```

In [16]:

```
myser = pd.Series(data=mydata)
```

In [17]:

```
myser
```

```
Out[17]:
```

```
0    1776
1    1867
2    1821
dtype: int64
```

```
In [18]:
```

```
pd.Series(data=mydata,index=myindex)
```

```
Out[18]:
```

```
USA        1776
Canada     1867
Mexico     1821
dtype: int64
```

```
In [23]:
```

```
ran_data = np.random.randint(0,100,4)
```

```
In [24]:
```

```
ran_data
```

```
Out[24]:
```

```
array([39, 35, 37, 23])
```

```
In [26]:
```

```
names = ['Andrew','Bobo','Claire','David']
```

```
In [27]:
```

```
ages = pd.Series(ran_data,names)
```

```
In [28]:
```

```
ages
```

```
Out[28]:
```

```
Andrew    39
Bobo       35
Claire     37
David      23
dtype: int32
```

From a Dictionary

```
In [29]:
```

```
ages = {'Sammy':5,'Frank':10,'Spike':7}
```

```
In [30]:
```

```
ages
```

```
Out[30]:
```

```
{'Frank': 10, 'Sammy': 5, 'Spike': 7}
```

```
In [31]:
```

```
pd.Series(ages)
```

```
Out[31]:
```

```
Out[31]:  
Sammy      5  
Frank     10  
Spike      7  
dtype: int64
```

Key Ideas of a Series

Named Index

In [32]:

```
# Imaginary Sales Data for 1st and 2nd Quarters for Global Company  
q1 = {'Japan': 80, 'China': 450, 'India': 200, 'USA': 250}  
q2 = {'Brazil': 100, 'China': 500, 'India': 210, 'USA': 260}
```

In [33]:

```
# Convert into Pandas Series  
sales_Q1 = pd.Series(q1)  
sales_Q2 = pd.Series(q2)
```

In [34]:

```
sales_Q1
```

Out[34]:

```
Japan      80  
China    450  
India     200  
USA       250  
dtype: int64
```

In [35]:

```
# Call values based on Named Index  
sales_Q1['Japan']
```

Out[35]:

```
80
```

In [36]:

```
# Integer Based Location information also retained!  
sales_Q1[0]
```

Out[36]:

```
80
```

Be careful with potential errors!

In [37]:

```
# Wrong Name  
# sales_Q1['France']
```

In [38]:

```
# Accidental Extra Space  
# sales_Q1['USA ']
```

In [39]:

```
# Capitalization Mistake
```

```
# sales_Q1['usa']
```

Operations

```
In [40]:
```

```
# Grab just the index keys
sales_Q1.keys()
```

```
Out[40]:
```

```
Index(['Japan', 'China', 'India', 'USA'], dtype='object')
```

```
In [41]:
```

```
# Can Perform Operations Broadcasted across entire Series
sales_Q1 * 2
```

```
Out[41]:
```

```
Japan      160
China      900
India       400
USA         500
dtype: int64
```

```
In [42]:
```

```
sales_Q2 / 100
```

```
Out[42]:
```

```
Brazil      1.0
China        5.0
India        2.1
USA          2.6
dtype: float64
```

Between Series

```
In [43]:
```

```
# Notice how Pandas informs you of mismatch with NaN
sales_Q1 + sales_Q2
```

```
Out[43]:
```

```
Brazil      NaN
China       950.0
India       410.0
Japan        NaN
USA         510.0
dtype: float64
```

```
In [44]:
```

```
# You can fill these with any value you want
sales_Q1.add(sales_Q2, fill_value=0)
```

```
Out[44]:
```

```
Brazil      100.0
China       950.0
India       410.0
Japan        80.0
USA         510.0
dtype: float64
```

That is all we need to know about Series, up next, DataFrames!

